

CMU, Pittsburgh, 9 November 2018

How to Solve a Linear System

Goran Flegar

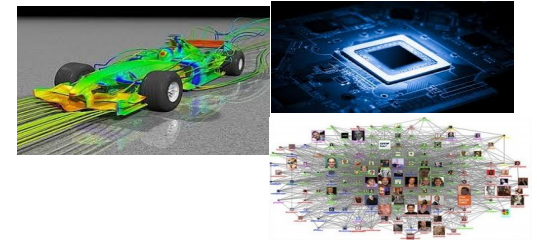
Hartwig Anzt, Yen-Chen Chen, Terry Cojean, Jack Dongarra, Nick Higham, Pratik Nayak, Enrique S. Quintana-Ortí, Mike Tsai



Scan me
for slides!

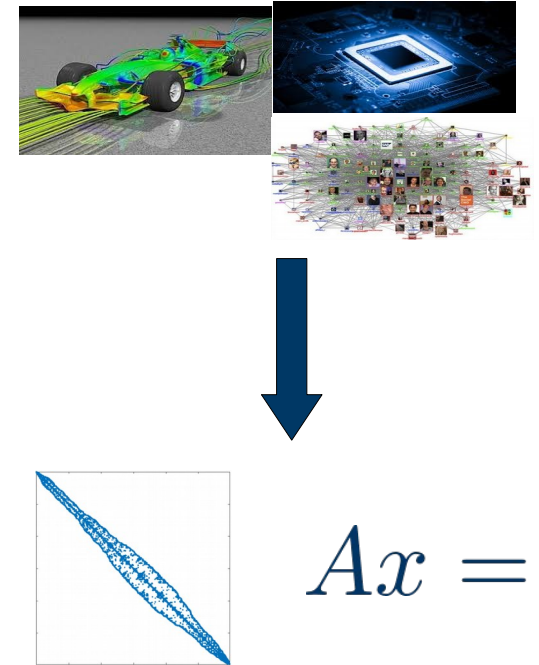
Sources of linear systems

- Real-world problem transformed into a linear system via:
 - PDE discretizations, graph representations
 - Large number of unknowns (1M+, full matrix 8TB)
 - Most matrix elements are zero



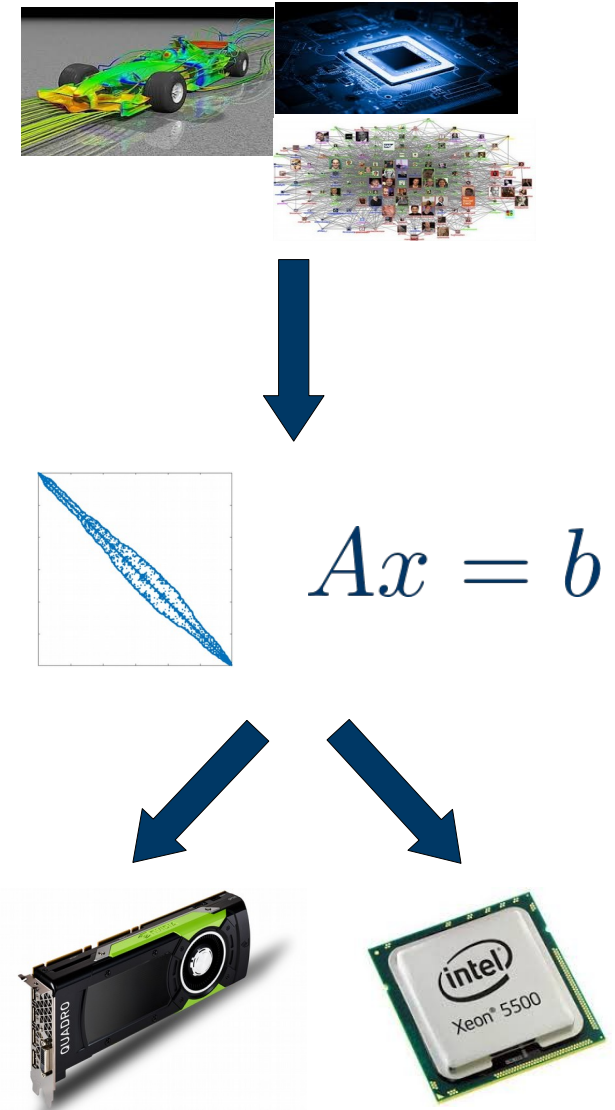
Sources of linear systems

- Real-world problem transformed into a linear system via:
 - PDE discretizations, graph representations
 - Large number of unknowns (1M+, full matrix 8TB)
 - Most matrix elements are zero
- Possible approach: iterative methods
 - Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver
 - Preconditioners
 - reduce the number of iterations



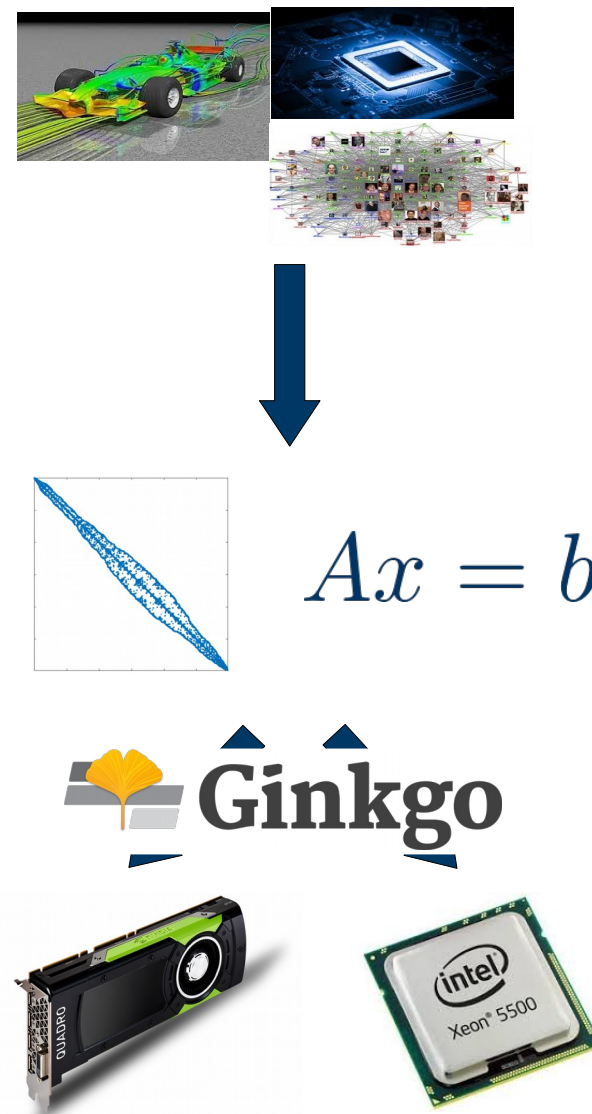
Sources of linear systems

- Real-world problem transformed into a linear system via:
 - PDE discretizations, graph representations
 - Large number of unknowns (1M+, full matrix 8TB)
 - Most matrix elements are zero
- Possible approach: iterative methods
 - Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver
 - Preconditioners
 - reduce the number of iterations
 - Special hardware (e.g. GPUs)
 - Probably not a good idea to implement everything from scratch...



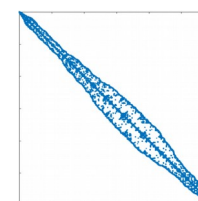
Sources of linear systems

- Real-world problem transformed into a linear system via:
 - PDE discretizations, graph representations
 - Large number of unknowns (1M+, full matrix 8TB)
 - Most matrix elements are zero
- Possible approach: iterative methods
 - Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver
 - Preconditioners
 - reduce the number of iterations
 - Special hardware (e.g. GPUs)
 - Probably not a good idea to implement everything from scratch...
 - Use a library instead:



Sources of linear systems

- Real-world problem transformed into a linear system via:
 - PDE discretizations, graph representations
 - Large number of unknowns (1M+, full matrix 8TB)
 - Most matrix elements are zero
- Possible approach: iterative methods
 - Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver
 - Preconditioners
 - reduce the number of iterations
 - Special hardware (e.g. GPUs)
 - Probably not a good idea to implement everything from scratch...



$$Ax = b$$



- Use a library instead:



The Ginkgo library

- Linear operator algebra library:
 - Matrix formats, preconditioners, (Krylov) solvers

Joint effort: Innovative Computing Lab at University of Tennessee, Knoxville; Karlsruhe Institute of Technology; University Jaume I



The Ginkgo library

- Linear operator algebra library:
 - Matrix formats, preconditioners, (Krylov) solvers
- Goals:
 - Backends for various devices (executors)
 - Easy composability – even when using building blocks from a different category (e.g. using a matrix or a solver as a preconditioner)
 - Easy extensibility
 - Support for matrix-free methods

Joint effort: Innovative Computing Lab at University of Tennessee, Knoxville; Karlsruhe Institute of Technology; Universitat Jaume I



Core idea

Matrix-vector product

$$x = Ab$$

Core idea

Matrix-vector product

$$x = Ab$$

Preconditioner

$$x = M^{-1}b$$

$$M \approx A$$

Core idea

Matrix-vector product

$$x = Ab$$

Preconditioner

$$x = M^{-1}b$$

$$\begin{aligned} M &\approx A \\ M^{-1} &= \Pi(A) \end{aligned}$$

Core idea

Matrix-vector product

$$x = Ab$$

Preconditioner

$$x = M^{-1}b$$

Solver

$$Ax = b$$

$$\begin{aligned} M &\approx A \\ M^{-1} &= \Pi(A) \end{aligned}$$

Core idea

Matrix-vector product

$$x = Ab$$

Preconditioner

$$x = M^{-1}b$$

$$M \approx A$$
$$M^{-1} = \Pi(A)$$

Solver

$$Ax = b$$
$$x = Sb$$

$$S = A^{-1}$$

Core idea

Matrix-vector product

$$x = Ab$$

Preconditioner

$$x = M^{-1}b$$

$$M \approx A$$

$$M^{-1} = \Pi(A)$$

Solver

$$Ax = b$$

$$x = Sb$$

$$S = A^{-1}$$

$$S = \Sigma(A)$$

Core idea

Matrix-vector product

$$x = Ab$$

Preconditioner

$$x = M^{-1}b$$

$$M \approx A$$
$$M^{-1} = \Pi(A)$$

Solver

$$Ax = b$$

$$x = Sb$$

$$S = A^{-1}$$
$$S = \Sigma(A)$$

All of them can be expressed as:

- applications of linear operators* (LinOp)

$$L : \mathbb{F}^m \rightarrow \mathbb{F}^n$$

* can be realized as a (non-linear) approximation

Core idea

Matrix-vector product

$$x = Ab$$

Preconditioner

$$x = M^{-1}b$$

$$M \approx A$$

$$M^{-1} = \Pi(A)$$

Solver

$$Ax = b$$

$$x = Sb$$

$$S = A^{-1}$$

$$S = \Sigma(A)$$

All of them can be expressed as:

- applications of linear operators* (LinOp)
- (non-linear) transformations applied to linear operators (Factory)

$$\begin{aligned} L &: \mathbb{F}^m \rightarrow \mathbb{F}^n \\ \Phi &: \mathbb{L}^{mn}(\mathbb{F}) \rightarrow \mathbb{L}^{kl}(\mathbb{F}) \end{aligned}$$

* can be realized as a (non-linear) approximation

Example

- Solving a symmetric positive definite linear system using the conjugate gradient (CG) solver, preconditioned with a block-Jacobi (BJ) preconditioner, with a system matrix stored in CSR format

Example

- Solving a symmetric positive definite linear system using the conjugate gradient (CG) solver, preconditioned with a block-Jacobi (BJ) preconditioner, with a system matrix stored in CSR format

1. Read the system matrix operator, right hand side and initial guess from file, standard input, or use existing data from the application

$$A_{CSR} \quad b_D \quad x_D$$

Example

- Solving a symmetric positive definite linear system using the conjugate gradient (CG) solver, preconditioned with a block-Jacobi (BJ) preconditioner, with a system matrix stored in CSR format

1. Read the system matrix operator, right hand side and initial guess from file, standard input, or use existing data from the application

$A_{CSR} \ b_D \ x_D$

2. Build the block-Jacobi preconditioner factory (no computation, binding parameters)

BJ

Example

- Solving a symmetric positive definite linear system using the conjugate gradient (CG) solver, preconditioned with a block-Jacobi (BJ) preconditioner, with a system matrix stored in CSR format

1. Read the system matrix operator, right hand side and initial guess from file, standard input, or use existing data from the application

$A_{\text{CSR}} \quad b_D \quad x_D$

2. Build the block-Jacobi preconditioner factory (no computation, binding parameters)

BJ

3. Use the preconditioner factory to build the preconditioned conjugate gradient factory (no computation, binding parameters)

CG_{BJ}

Example

- Solving a symmetric positive definite linear system using the conjugate gradient (CG) solver, preconditioned with a block-Jacobi (BJ) preconditioner, with a system matrix stored in CSR format

1. Read the system matrix operator, right hand side and initial guess from file, standard input, or use existing data from the application

$$A_{\text{CSR}} \quad b_D \quad x_D$$

2. Build the block-Jacobi preconditioner factory (no computation, binding parameters)

$$BJ$$

3. Use the preconditioner factory to build the preconditioned conjugate gradient factory (no computation, binding parameters)

$$CG_{BJ}$$

4. Generate the solver operator from the system matrix using the solver factory (preprocessing needed for the solver and preconditioner)

$$S_A = CG_{BJ}(A_{\text{CSR}})$$

Example

- Solving a symmetric positive definite linear system using the conjugate gradient (CG) solver, preconditioned with a block-Jacobi (BJ) preconditioner, with a system matrix stored in CSR format

1. Read the system matrix operator, right hand side and initial guess from file, standard input, or use existing data from the application

$$A_{\text{CSR}} \quad b_D \quad x_D$$

2. Build the block-Jacobi preconditioner factory (no computation, binding parameters)

$$BJ$$

3. Use the preconditioner factory to build the preconditioned conjugate gradient factory (no computation, binding parameters)

$$CG_{BJ}$$

4. Generate the solver operator from the system matrix using the solver factory (preprocessing needed for the solver and preconditioner)

$$S_A = CG_{BJ}(A_{\text{CSR}})$$

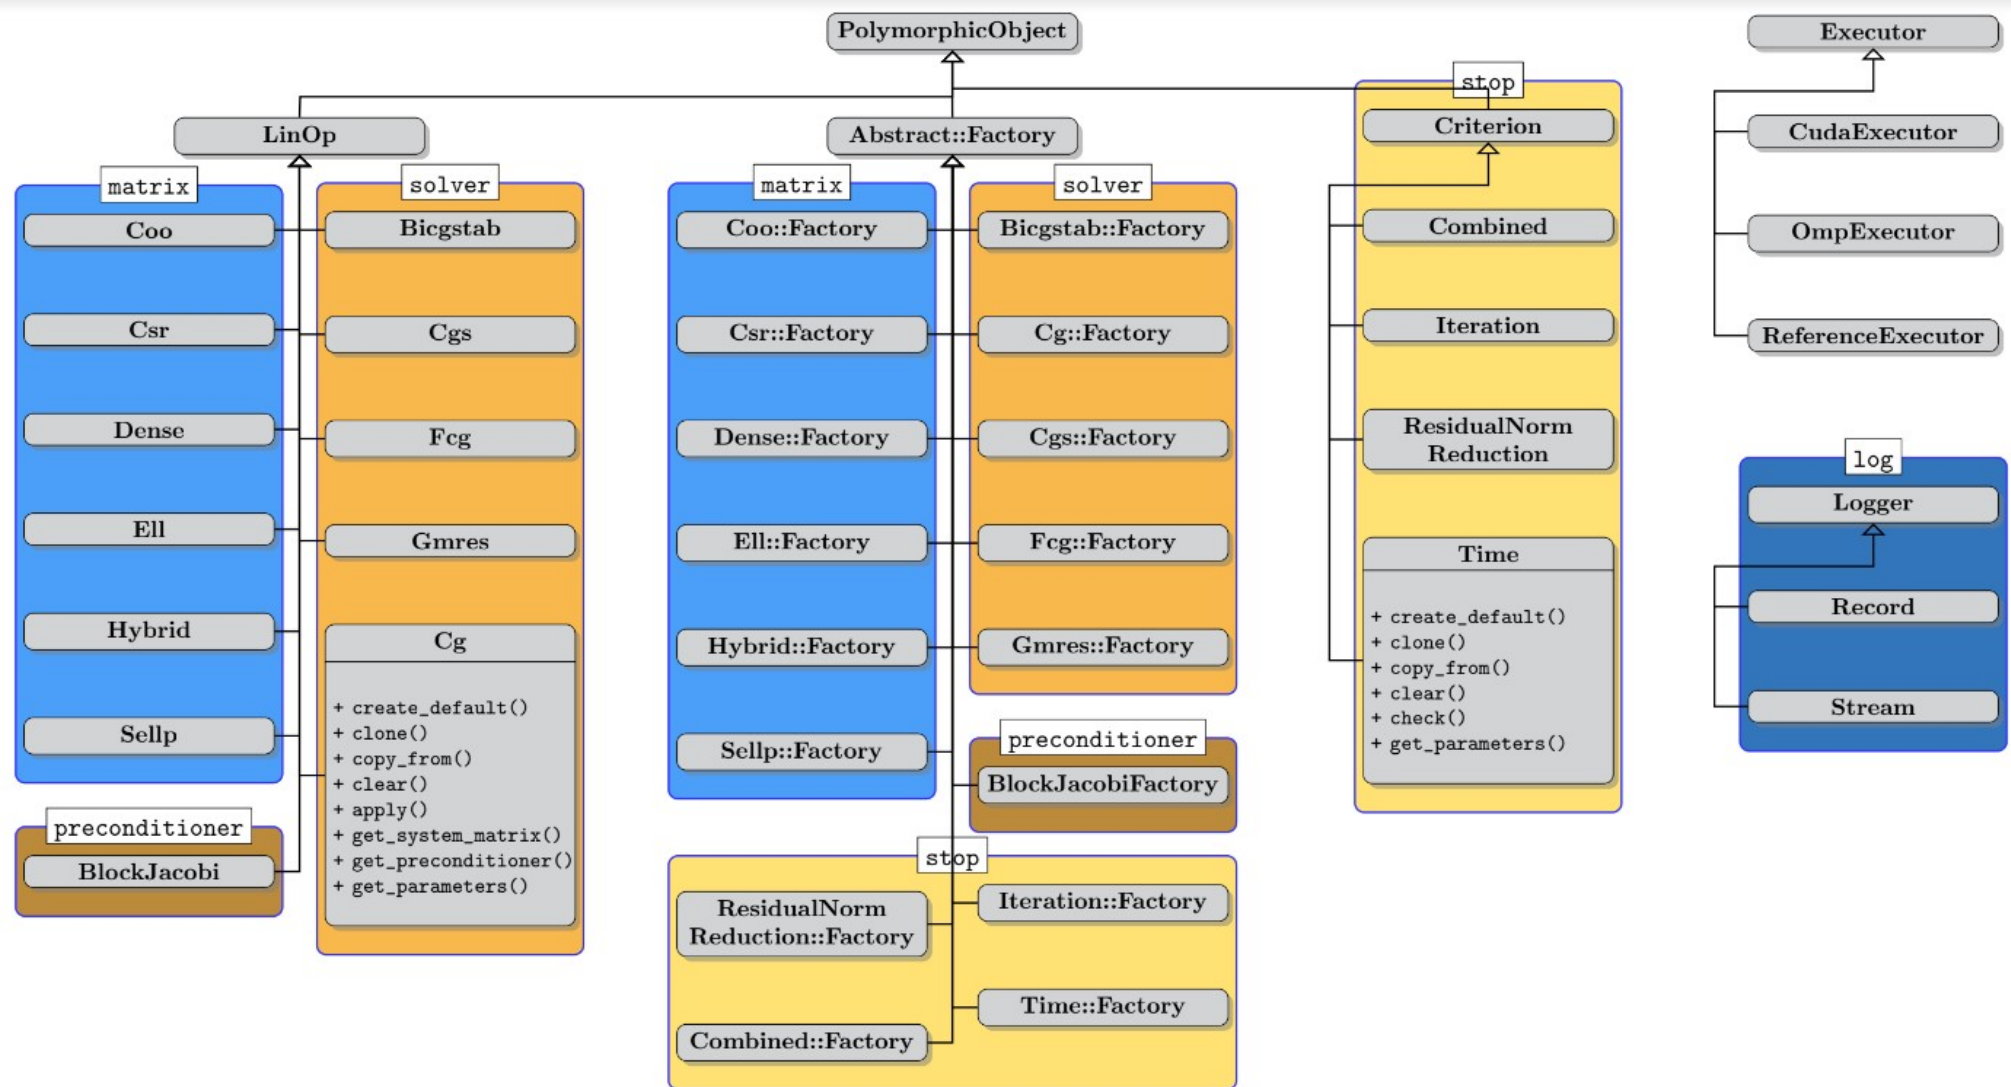
5. Solve the system (actual computation)

$$x_D = S_A b_D$$

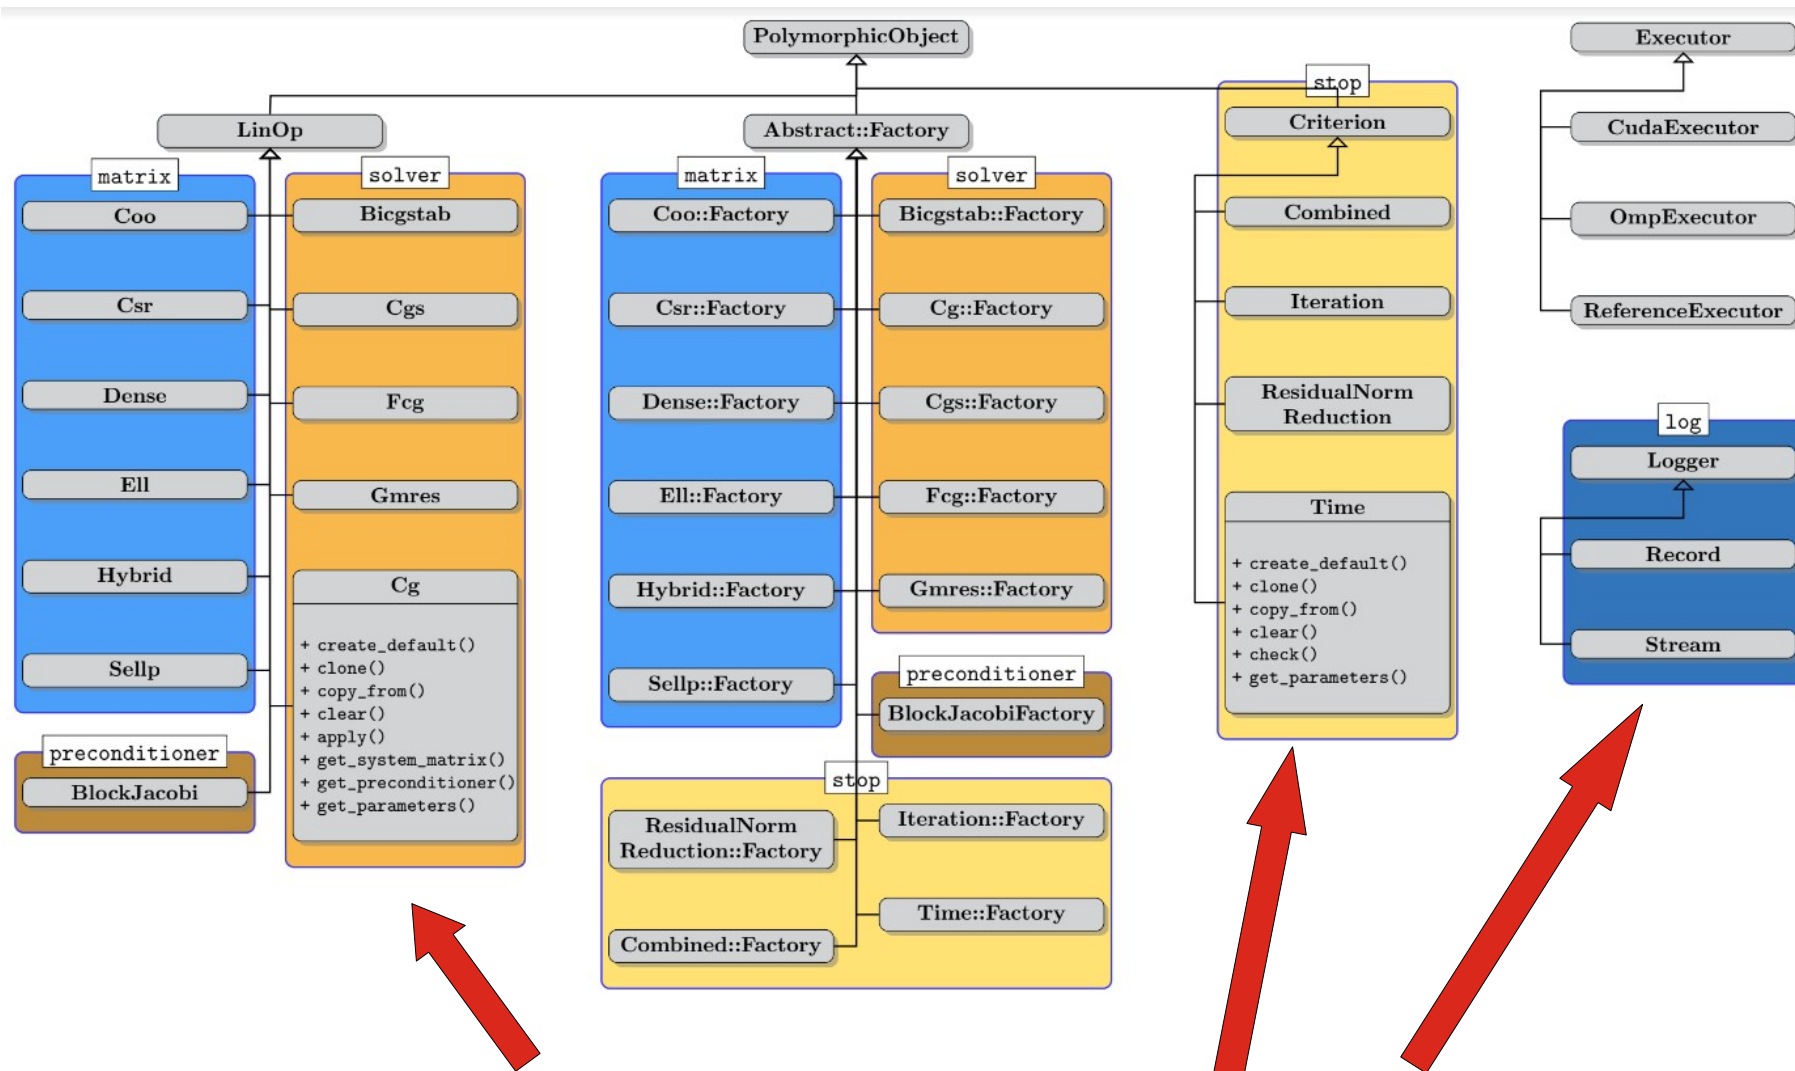
Example source code

```
int main() {  
    // Instantiate a CUDA executor  
    auto gpu = gko::CudaExecutor::create(0, gko::OmpExecutor::create());  
    // Read data  
    auto A = gko::read<gko::matrix::Csr<>>(std::cin, gpu);  
    auto b = gko::read<gko::matrix::Dense<>>(std::cin, gpu);  
    auto x = gko::read<gko::matrix::Dense<>>(std::cin, gpu);  
    // Create the solver  
    auto solver = gko::solver::Cg<>::build()  
        .with_preconditioner(  
            gko::preconditioner::Jacobi<>::build().with_max_block_size(32).on(gpu))  
        .with_criteria(  
            gko::stop::Iteration::build().with_max_iters(20u).on(gpu),  
            gko::stop::ResidualNormReduction<>::build()  
                .with_reduction_factor(1e-15).on(gpu))  
        .on(gpu);  
    // Solve system  
    solver->generate(give(A))->apply(lend(b), lend(x));  
    // Write result  
    write(std::cout, lend(x));  
}
```

Library features



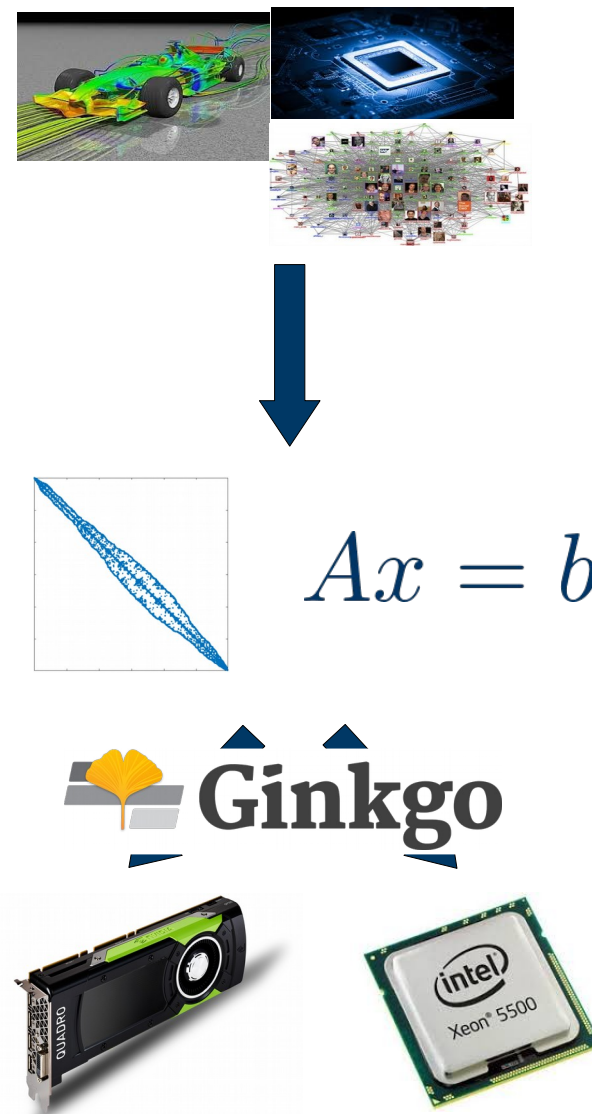
Library features: extensibility



users can provide new matrices, solvers, preconditioners, stopping criteria, loggers
Without recompiling the library!

Sources of linear systems

- Real-world problem transformed into a linear system via:
 - PDE discretizations, graph representations
 - Large number of unknowns (1M+, full matrix 8TB)
 - Most matrix elements are zero
- Possible approach: iterative methods
 - Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver
 - Preconditioners
 - reduce the number of iterations
 - Special hardware (e.g. GPUs)
 - Probably not a good idea to implement everything from scratch...
 - Use a library instead:



Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n}$$

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an
equivalent preconditioned system

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an
equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an
equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

~~$$M^{-1}A$$~~

**Do not compute the preconditioned
system matrix explicitly!**

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

~~$$M^{-1}A$$~~

$$y := (M^{-1}A)x$$

Do not compute the preconditioned system matrix explicitly!

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

~~$$M^{-1}A$$~~

Do not compute the preconditioned system matrix explicitly!

$$y := (M^{-1}A)x$$



$$z := Ax$$

$$y := M^{-1}z$$

Preconditioner application

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

~~$$M^{-1}A$$~~

Do not compute the preconditioned system matrix explicitly!

$$y := (M^{-1}A)x$$



Generate the preconditioner matrix, and store it in a form suitable for application

$$A \rightsquigarrow M^{-1}$$

Preconditioner setup



$$\begin{aligned} z &:= Ax \\ y &:= M^{-1}z \end{aligned}$$

Preconditioner application

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

~~$$M^{-1}A$$~~

Do not compute the preconditioned system matrix explicitly!

$$y := (M^{-1}A)x$$



Generate the preconditioner matrix, and store it in a form suitable for application

$$A \rightsquigarrow M^{-1}$$

Generation via factory



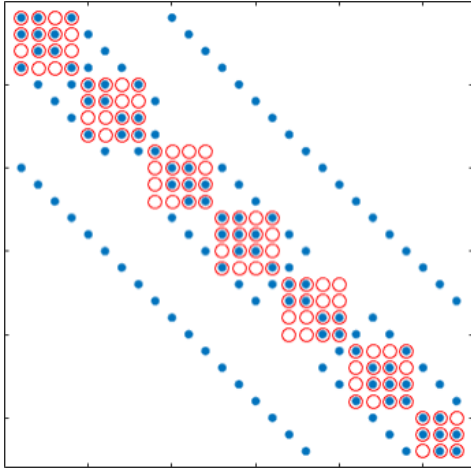
$$\begin{aligned} z &:= Ax \\ y &:= M^{-1}z \end{aligned}$$

Linear operator application



Ginkgo linear operator abstraction

Example: Block-Jacobi preconditioning



- Block-Jacobi preconditioning
 - Use only diagonal blocks for approximation

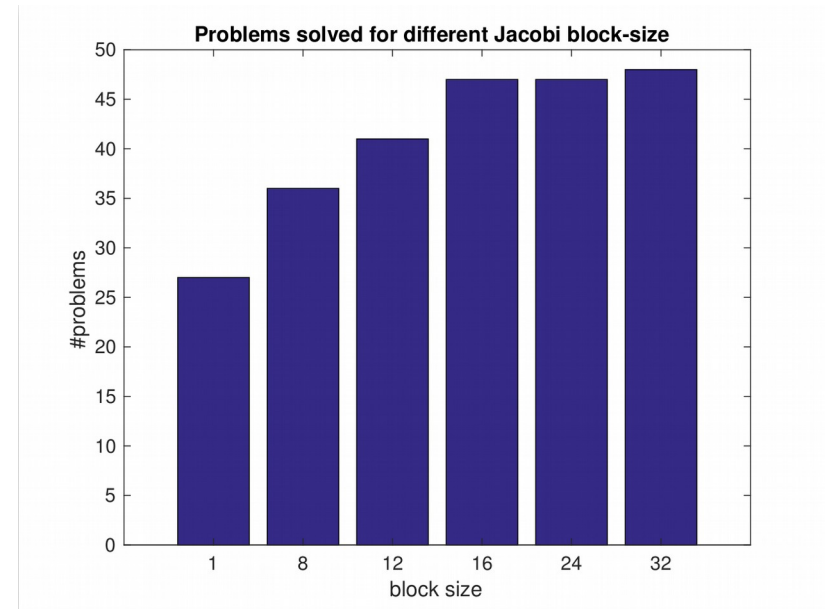
$$\text{diag}(A) = [D_1, \dots, D_k]$$

$$M := \text{diag}(D_1, \dots, D_k)$$

Anzt, Dongarra, Flegar, Quintana-Ortí, *Variable-size batched Gauss–Jordan elimination for block-Jacobi preconditioning on graphics processors*, ParCo

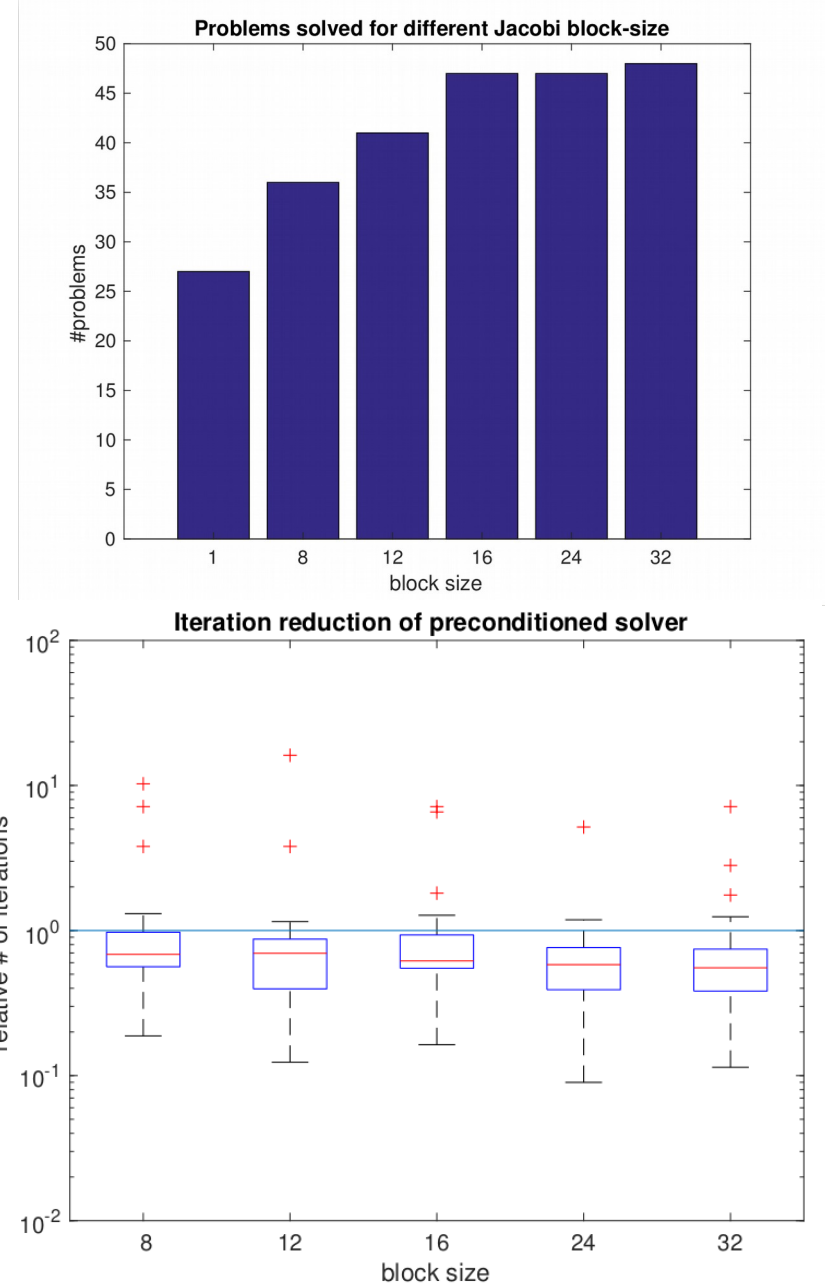
Benefits of block-Jacobi

- 56 matrices from SuiteSparse with inherent block structure
- MAGMA-sparse open source library
 - IDR solver
 - Supervariable agglomeration
 - Detects block structure of the matrix
- Improves the robustness of the solver



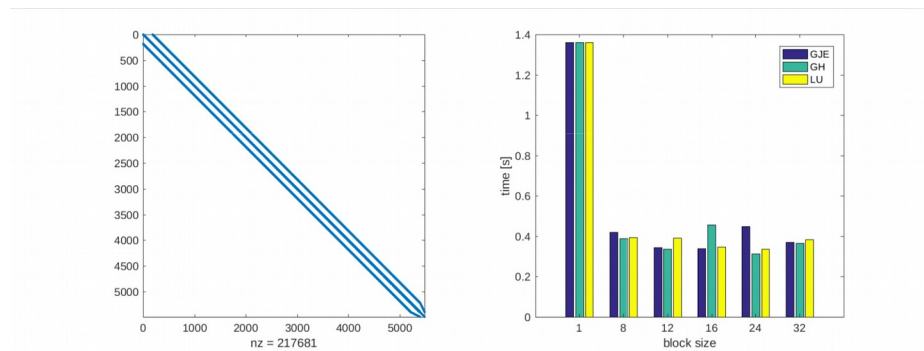
Benefits of block-Jacobi

- 56 matrices from SuiteSparse with inherent block structure
- MAGMA-sparse open source library
 - IDR solver
 - Supervariable agglomeration
 - Detects block structure of the matrix
- Improves the robustness of the solver
- Improves convergence of the solver

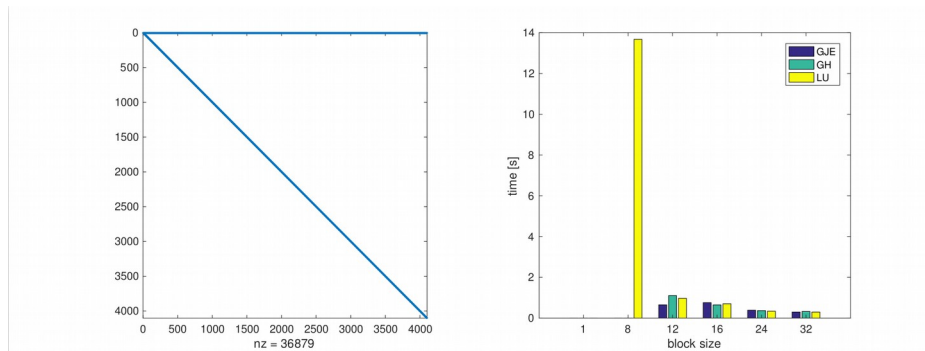


Complete solver runtime

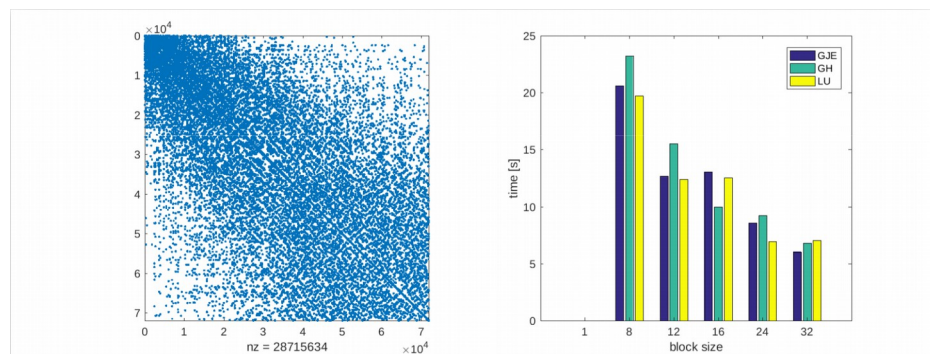
s2rmt3m1



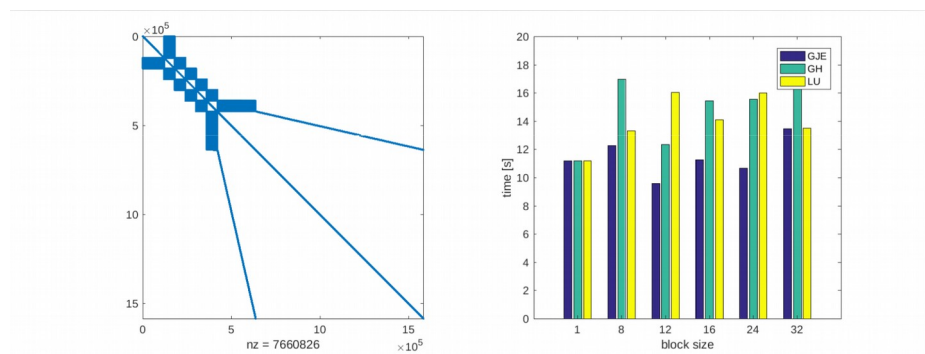
Chebyshev3



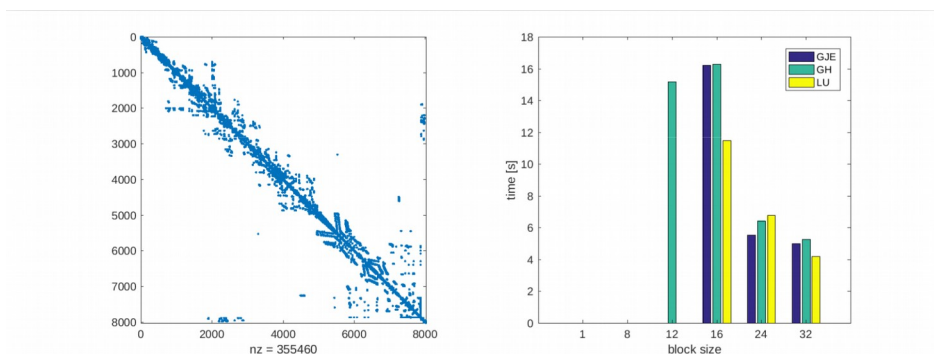
nd24k



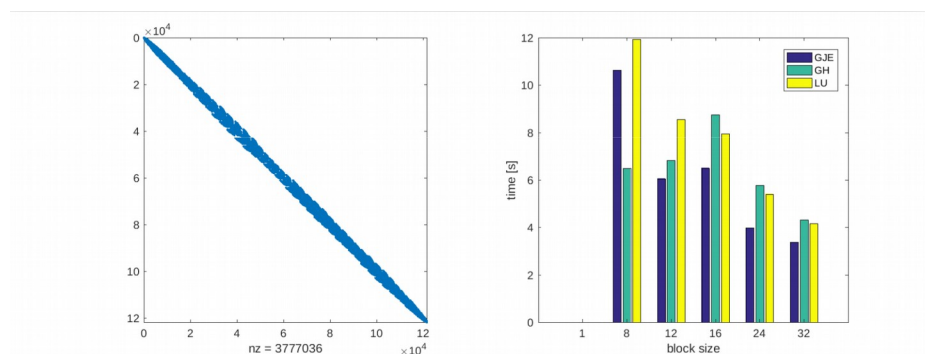
G3_circuit



bcsstk38



ship_003



Current Research: Adaptive precision block-Jacobi

Preconditioner is an **approximation** of the system matrix

- Applying a preconditioner inherently carries an **error**
- For block-Jacobi the relative error of z is usually around 0.01 – 0.1

$$z := M^{-1}y \approx A^{-1}y$$

Current Research: Adaptive precision block-Jacobi

Preconditioner is an **approximation** of the system matrix

- Applying a preconditioner inherently carries an **error**
- For block-Jacobi the relative error of z is usually around 0.01 – 0.1

$$z := M^{-1}y \approx A^{-1}y$$

Preconditioner application is **memory bounded**

- Most of the cost comes from reading the matrix from memory
- Idea: use **lower precision** to **store** the matrix

Current Research: Adaptive precision block-Jacobi

Preconditioner is an **approximation** of the system matrix

- Applying a preconditioner inherently carries an **error**
- For block-Jacobi the relative error of z is usually around 0.01 – 0.1

$$z := M^{-1}y \approx A^{-1}y$$

Preconditioner application is **memory bounded**

- Most of the cost comes from reading the matrix from memory
- Idea: use **lower precision** to **store** the matrix

Adaptive precision in inversion-based block-Jacobi:

- All **computation** is done in **double precision**
- Preconditioner matrix is **stored** in **lower precision**, with roundoff error “ u ”
- Error bound:

$$\frac{\|\delta z_i\|}{\|z_i\|} \lesssim (c_i \kappa(D_i) u_d + u) \kappa(D_i)$$

Current Research: Adaptive precision block-Jacobi

Preconditioner is an **approximation** of the system matrix

- Applying a preconditioner inherently carries an **error**
- For block-Jacobi the relative error of z is usually around 0.01 – 0.1

$$z := M^{-1}y \approx A^{-1}y$$

Preconditioner application is **memory bounded**

- Most of the cost comes from reading the matrix from memory
- Idea: use **lower precision** to **store** the matrix

Adaptive precision in inversion-based block-Jacobi:

- All **computation** is done in **double precision**
- Preconditioner matrix is **stored** in **lower precision**, with roundoff error “ u ”
- Error bound:

$$\frac{\|\delta z_i\|}{\|z_i\|} \lesssim (c_i \kappa(D_i) u_d + u) \kappa(D_i)$$

Assuming the preconditioner block is relatively well conditioned

- The error is determined by the product of u , and the condition number
- **Choose the precision for each block independently, such that at least 1 digit of the result is correct**

Experimental results

Determining the precision:

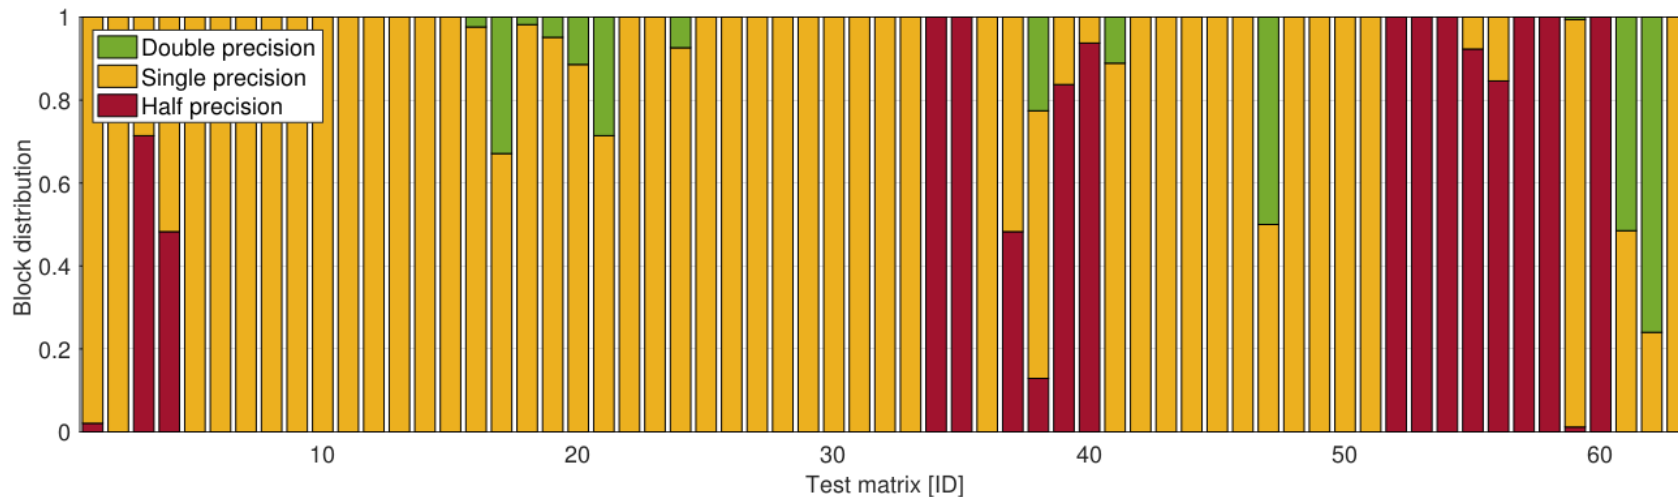
$$\begin{cases} \text{fp16} & \text{if } 0 < \kappa_1(D_i) \leq 10^2, \\ \text{fp32} & \text{if } 10^2 < \kappa_1(D_i) \leq 10^6, \text{ and} \\ \text{fp64} & \text{otherwise,} \end{cases}$$

Experimental results

Determining the precision:

$$\begin{cases} \text{fp16} & \text{if } 0 < \kappa_1(D_i) \leq 10^2, \\ \text{fp32} & \text{if } 10^2 < \kappa_1(D_i) \leq 10^6, \text{ and} \\ \text{fp64} & \text{otherwise,} \end{cases}$$

% of diagonal blocks stored in each precision *:

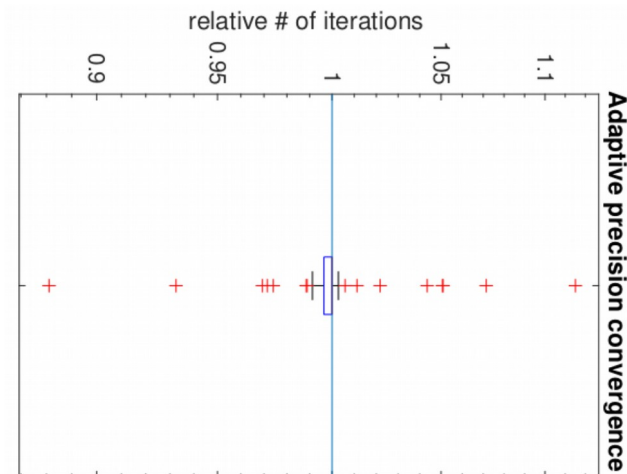


* Prototype implementation in MATLAB,
Results on 63 matrices from SuiteSparse

Experimental results

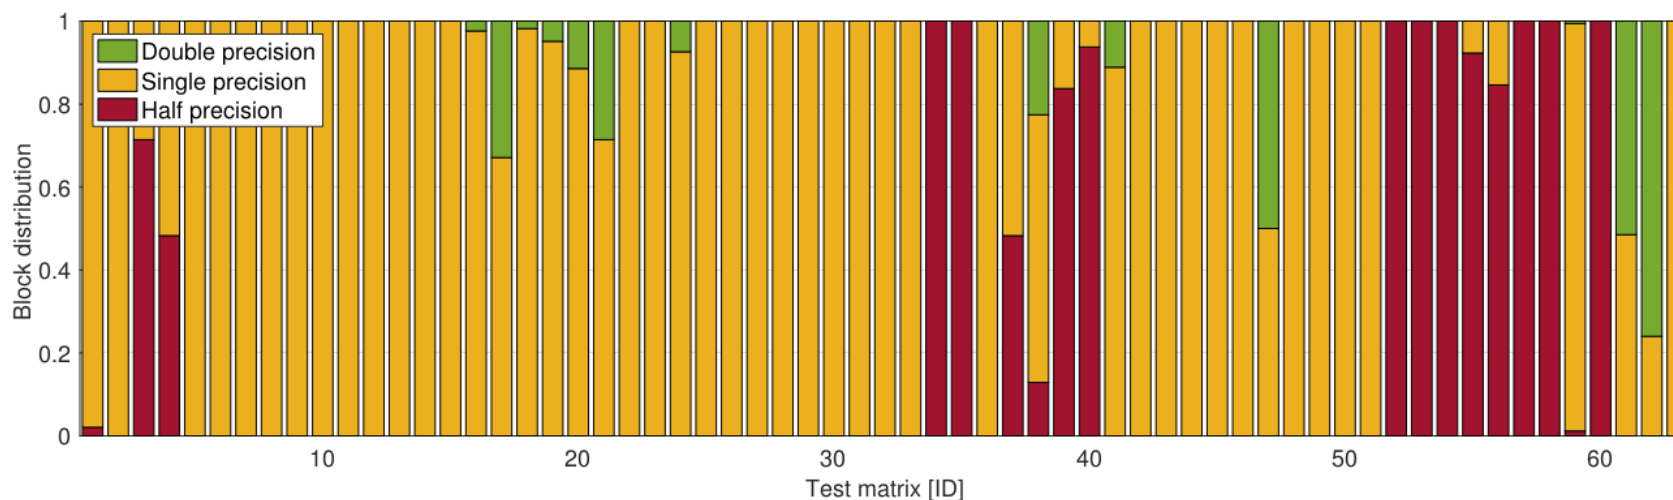
Determining the precision:

$$\begin{cases} \text{fp16} & \text{if } 0 < \kappa_1(D_i) \leq 10^2, \\ \text{fp32} & \text{if } 10^2 < \kappa_1(D_i) \leq 10^6, \text{ and} \\ \text{fp64} & \text{otherwise,} \end{cases}$$



Usually < 5% iteration increase

% of diagonal blocks stored in each precision *:



* Prototype implementation in MATLAB,
Results on 63 matrices from SuiteSparse

Energy efficiency of adaptive precision block-Jacobi

Memory access is an order of magnitude more expensive than computation! [Shalf 2013]

Energy efficiency of adaptive precision block-Jacobi

Memory access is an order of magnitude more expensive than computation! [Shalf 2013]

Energy model assumptions:

- Accessing 1 bit of data has a cost of 1 (energy unit)
- **Disregard** energy cost of **arithmetic** operations
- Total energy cost of each iteration:

$$\underbrace{14n \cdot \text{fp64}}_{\text{vector memory transfers}} + \underbrace{(2n + n_z) \cdot \text{fp64} + (n + n_z) \cdot \text{int32}}_{\text{CSR-SpMV memory transfers}} + \underbrace{2n \cdot \text{fp64} + \sum_{i=1}^N m_i^2 \cdot \text{fpxx}_i}_{\text{preconditioner memory transfers}}$$

Energy efficiency of adaptive precision block-Jacobi

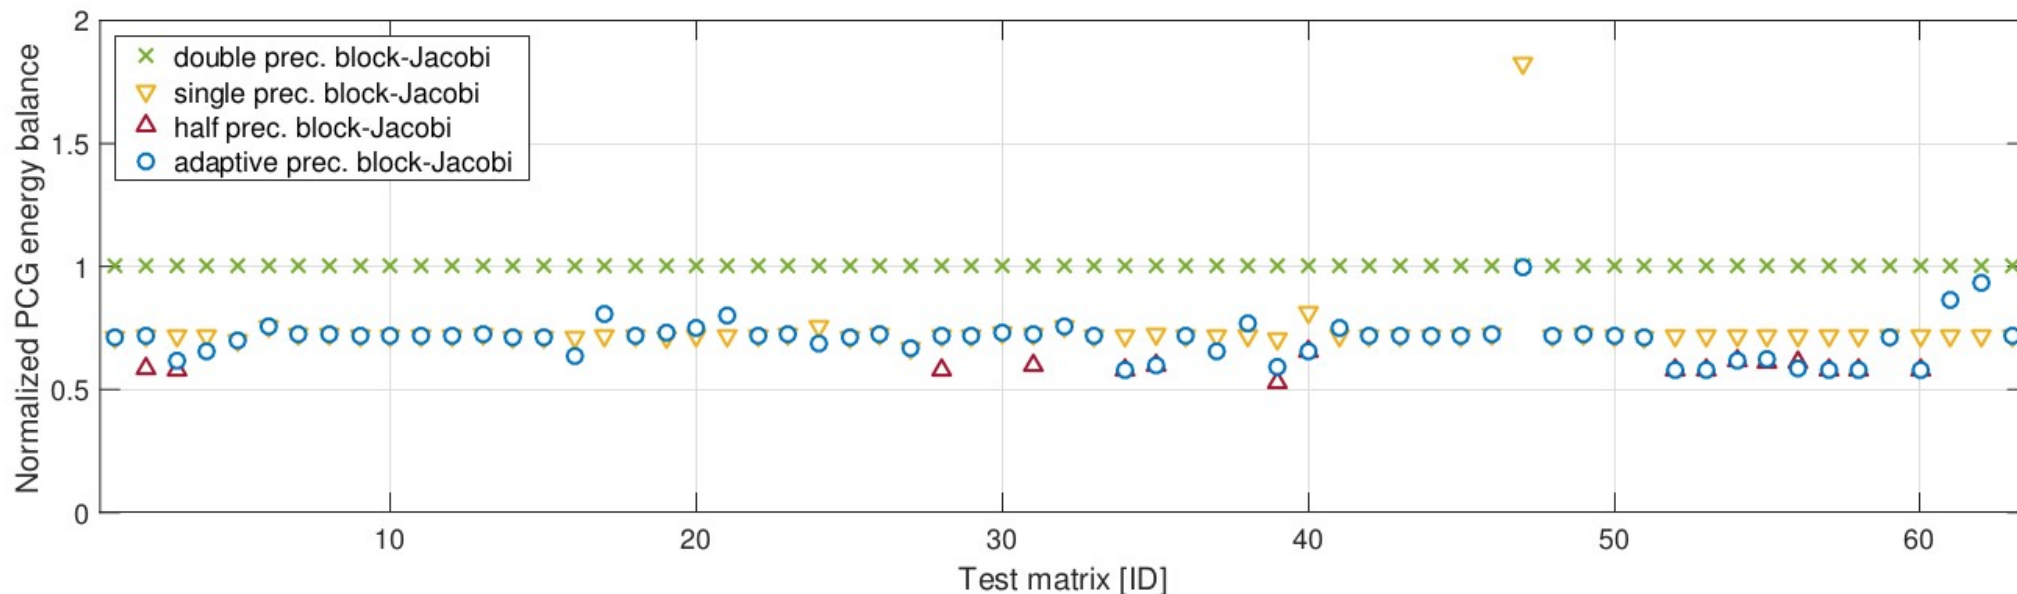
Memory access is an order of magnitude more expensive than computation! [Shalf 2013]

Energy model assumptions:

- Accessing 1 bit of data has a cost of 1 (energy unit)
- **Disregard** energy cost of **arithmetic** operations
- Total energy cost of each iteration:

$$\underbrace{14n \cdot \text{fp64}}_{\text{vector memory transfers}} + \underbrace{(2n + n_z) \cdot \text{fp64} + (n + n_z) \cdot \text{int32}}_{\text{CSR-SpMV memory transfers}} + \underbrace{2n \cdot \text{fp64} + \sum_{i=1}^N m_i^2 \cdot \text{fp} \times x_i}_{\text{preconditioner memory transfers}}$$

Predicted energy savings of adaptive precision block-Jacobi:



Overflow and underflow

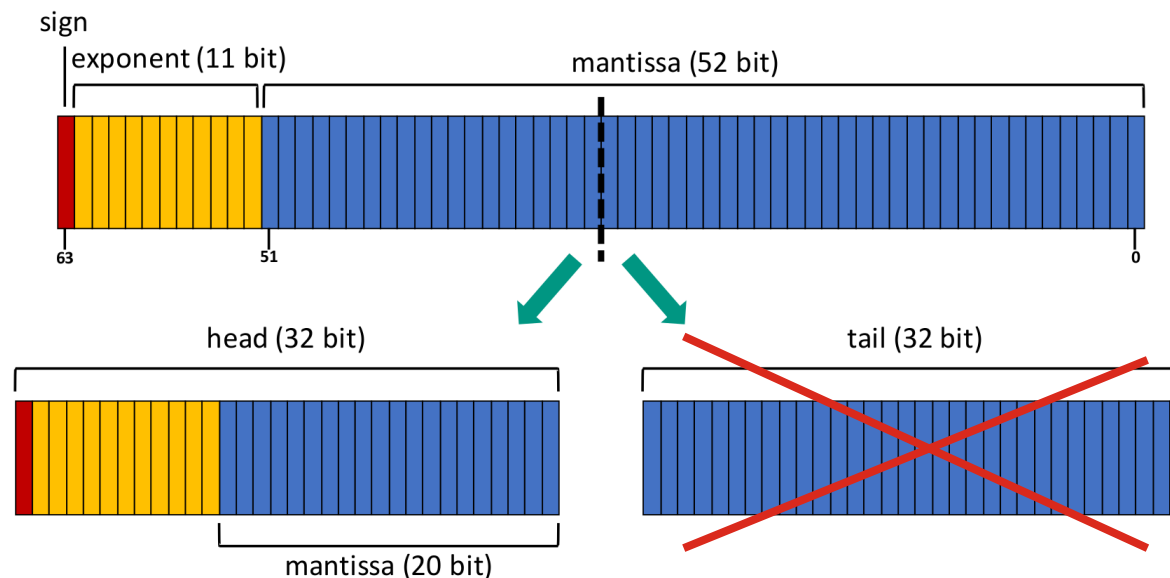
- Storing the block in lower precision can cause underflow or overflow in some values.
- This is not accounted for by the numerical analysis shown in the previous slides

Overflow and underflow

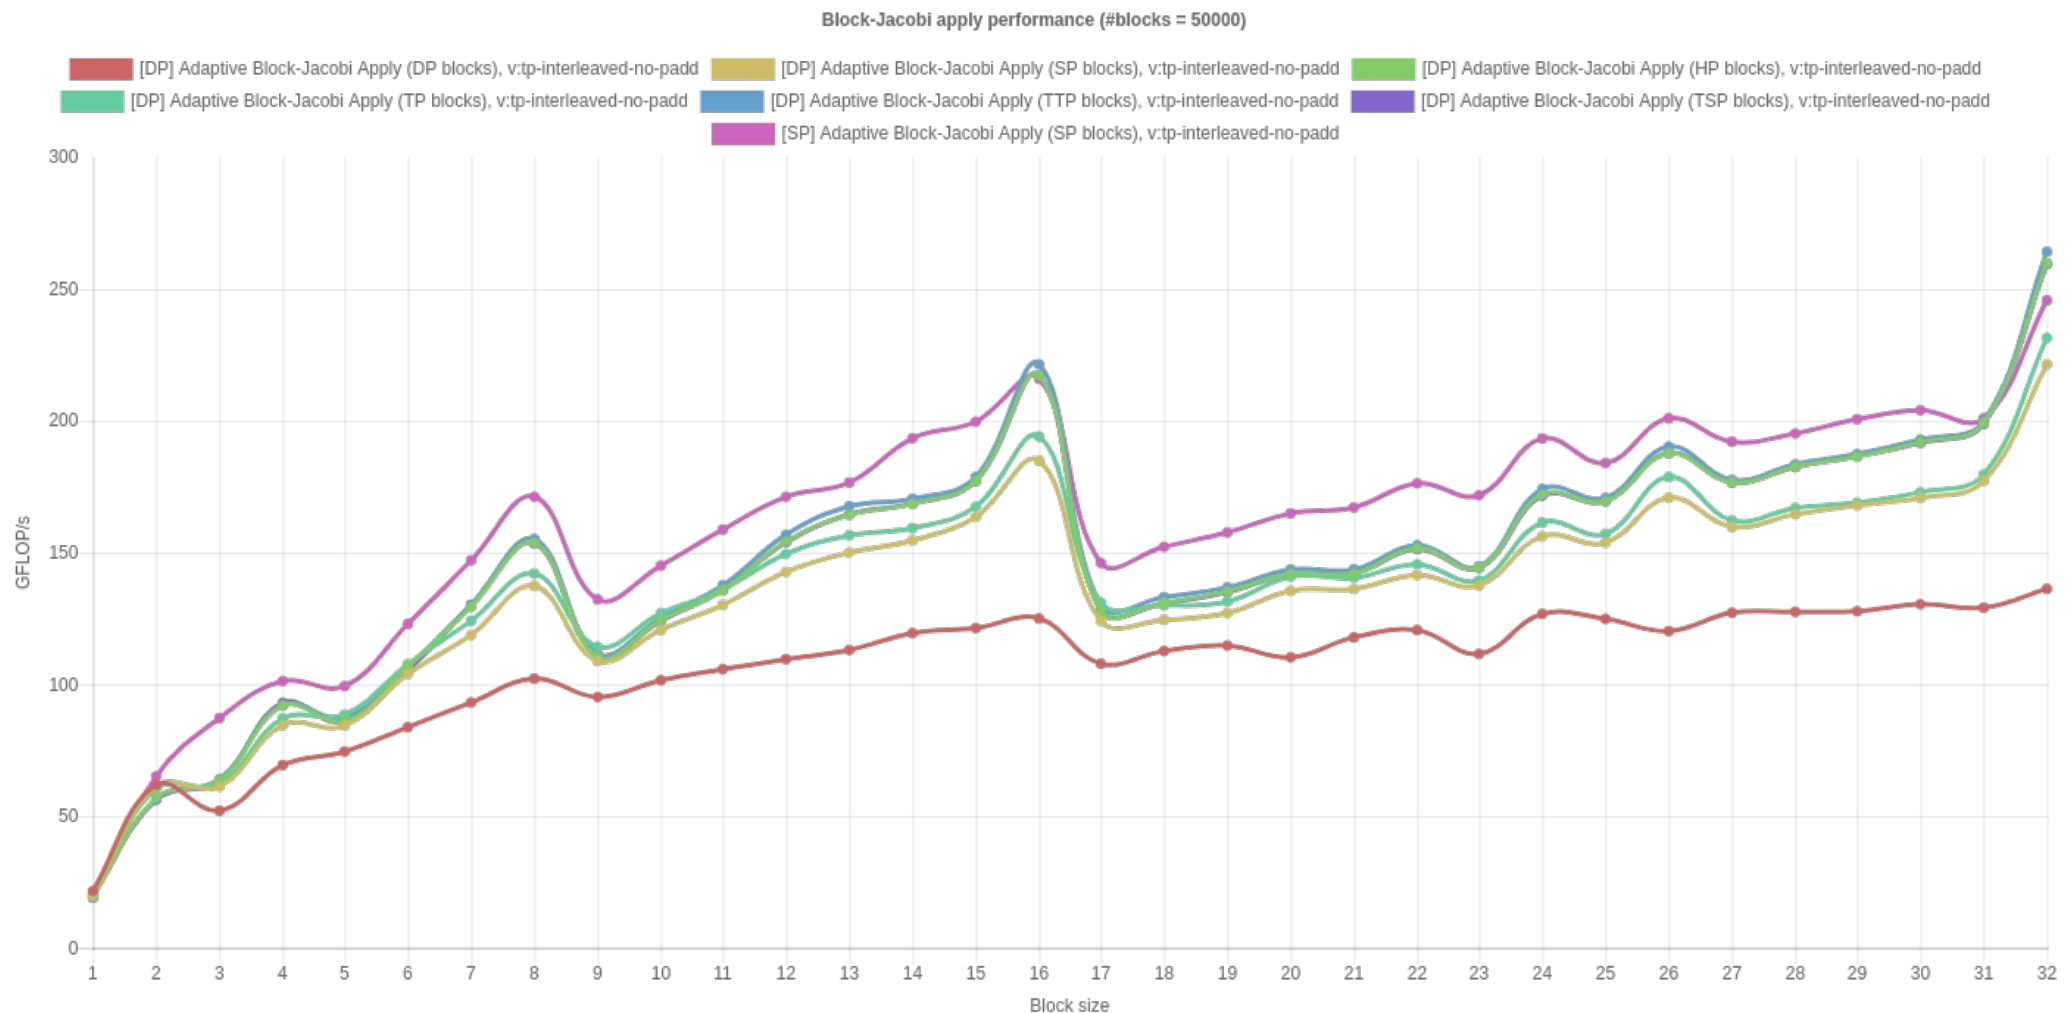
- Storing the block in lower precision can cause underflow or overflow in some values.
- This is not accounted for by the numerical analysis shown in the previous slides
- Two possible solutions:
 1. **Check the condition number** of the low precision block to verify there were no catastrophic overflows or underflows.

Overflow and underflow

- Storing the block in lower precision can cause underflow or overflow in some values.
- This is not accounted for by the numerical analysis shown in the previous slides
- Two possible solutions:
 - Check the condition number of the low precision block to verify there were no catastrophic overflows or underflows.
 - Use a custom storage format that preserves the number of exponent bits:

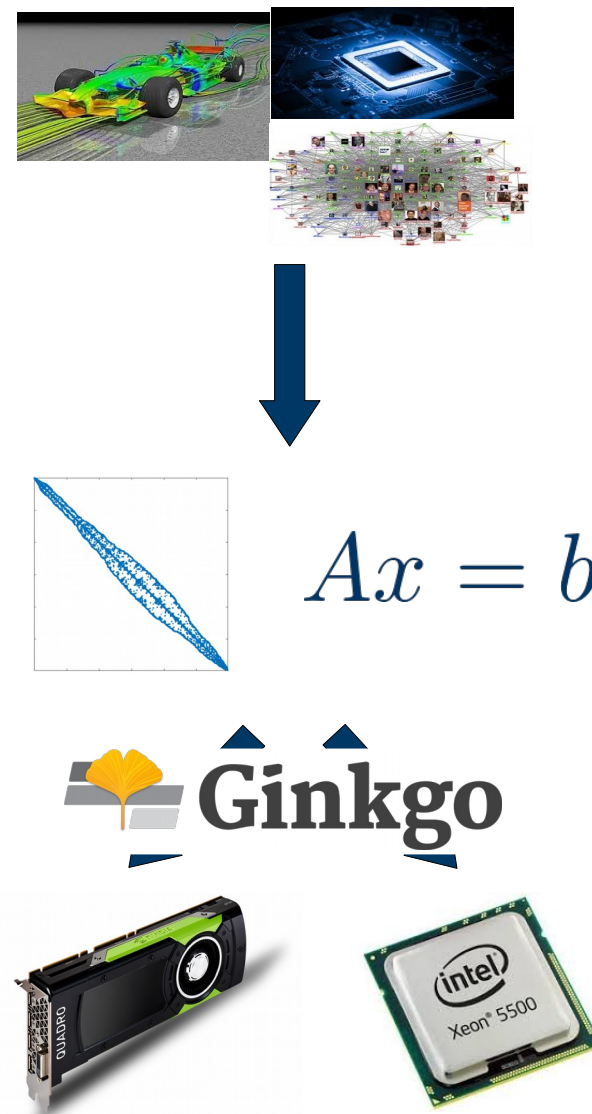


Initial GPU implementation



Sources of linear systems

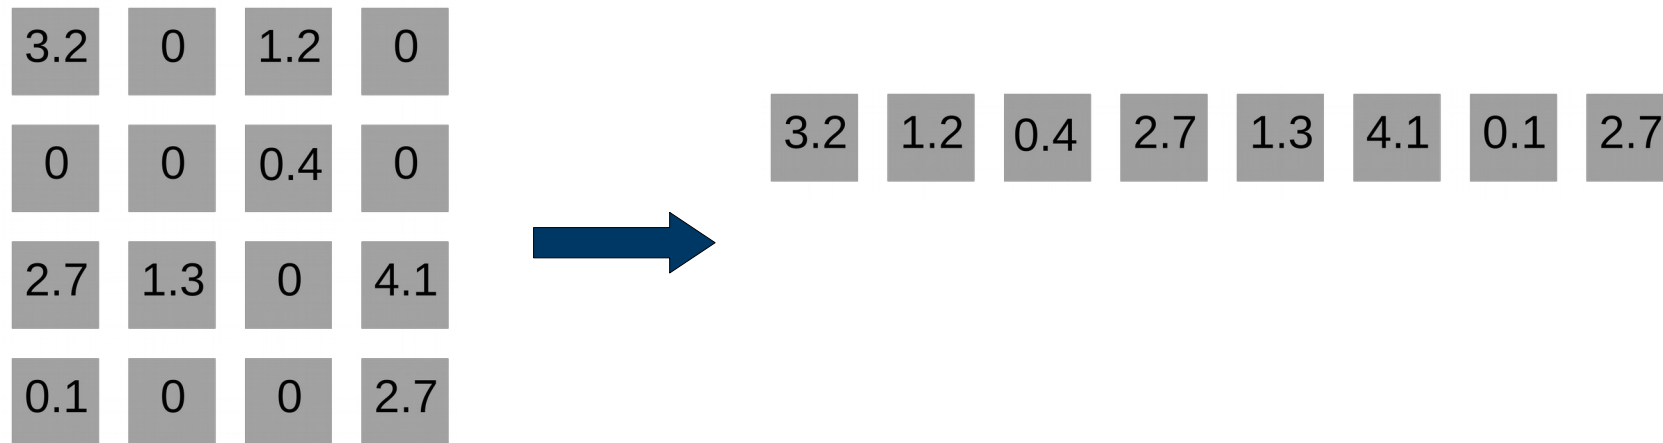
- Real-world problem transformed into a linear system via:
 - PDE discretizations, graph representations
 - Large number of unknowns (1M+, full matrix 8TB)
 - Most matrix elements are zero
- Possible approach: iterative methods
 - Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver
 - Preconditioners
 - reduce the number of iterations
 - Special hardware (e.g. GPUs)
 - Probably not a good idea to implement everything from scratch...
 - Use a library instead:



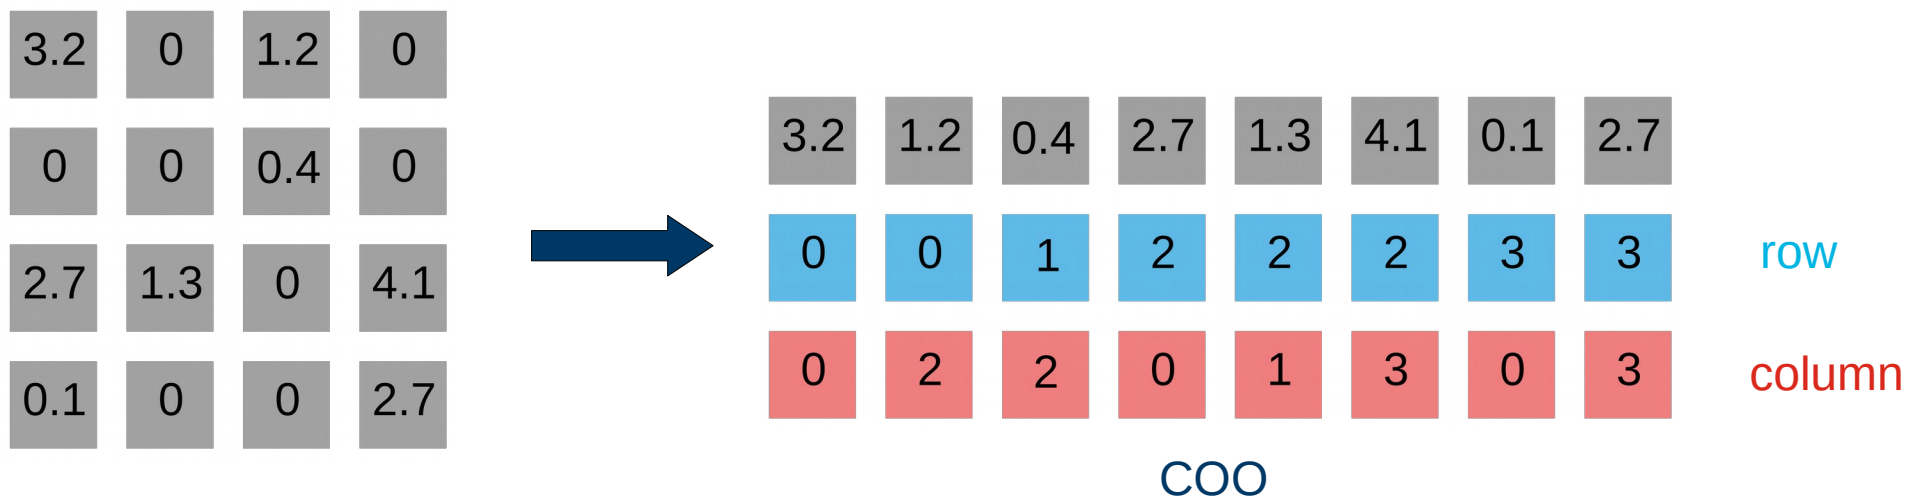
Sparse matrix formats

3.2	0	1.2	0
0	0	0.4	0
2.7	1.3	0	4.1
0.1	0	0	2.7

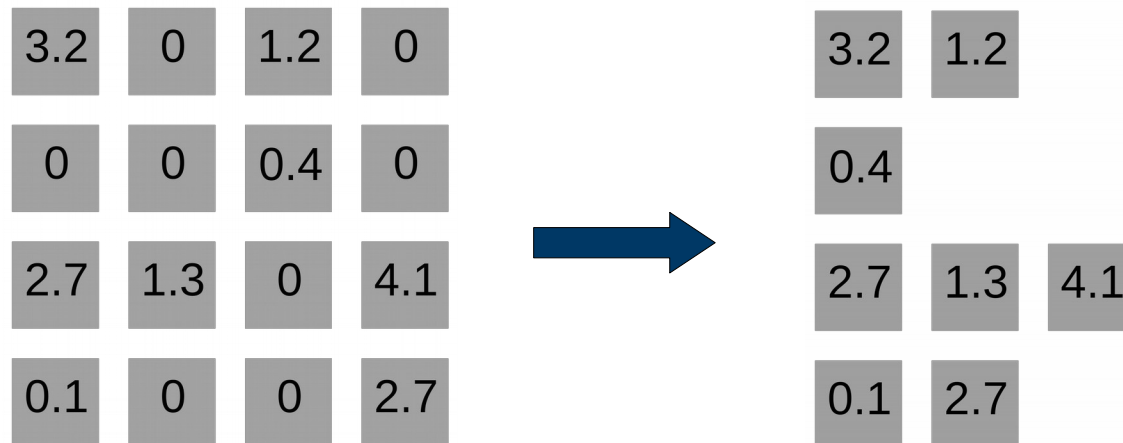
Sparse matrix formats



Sparse matrix formats



Sparse matrix formats



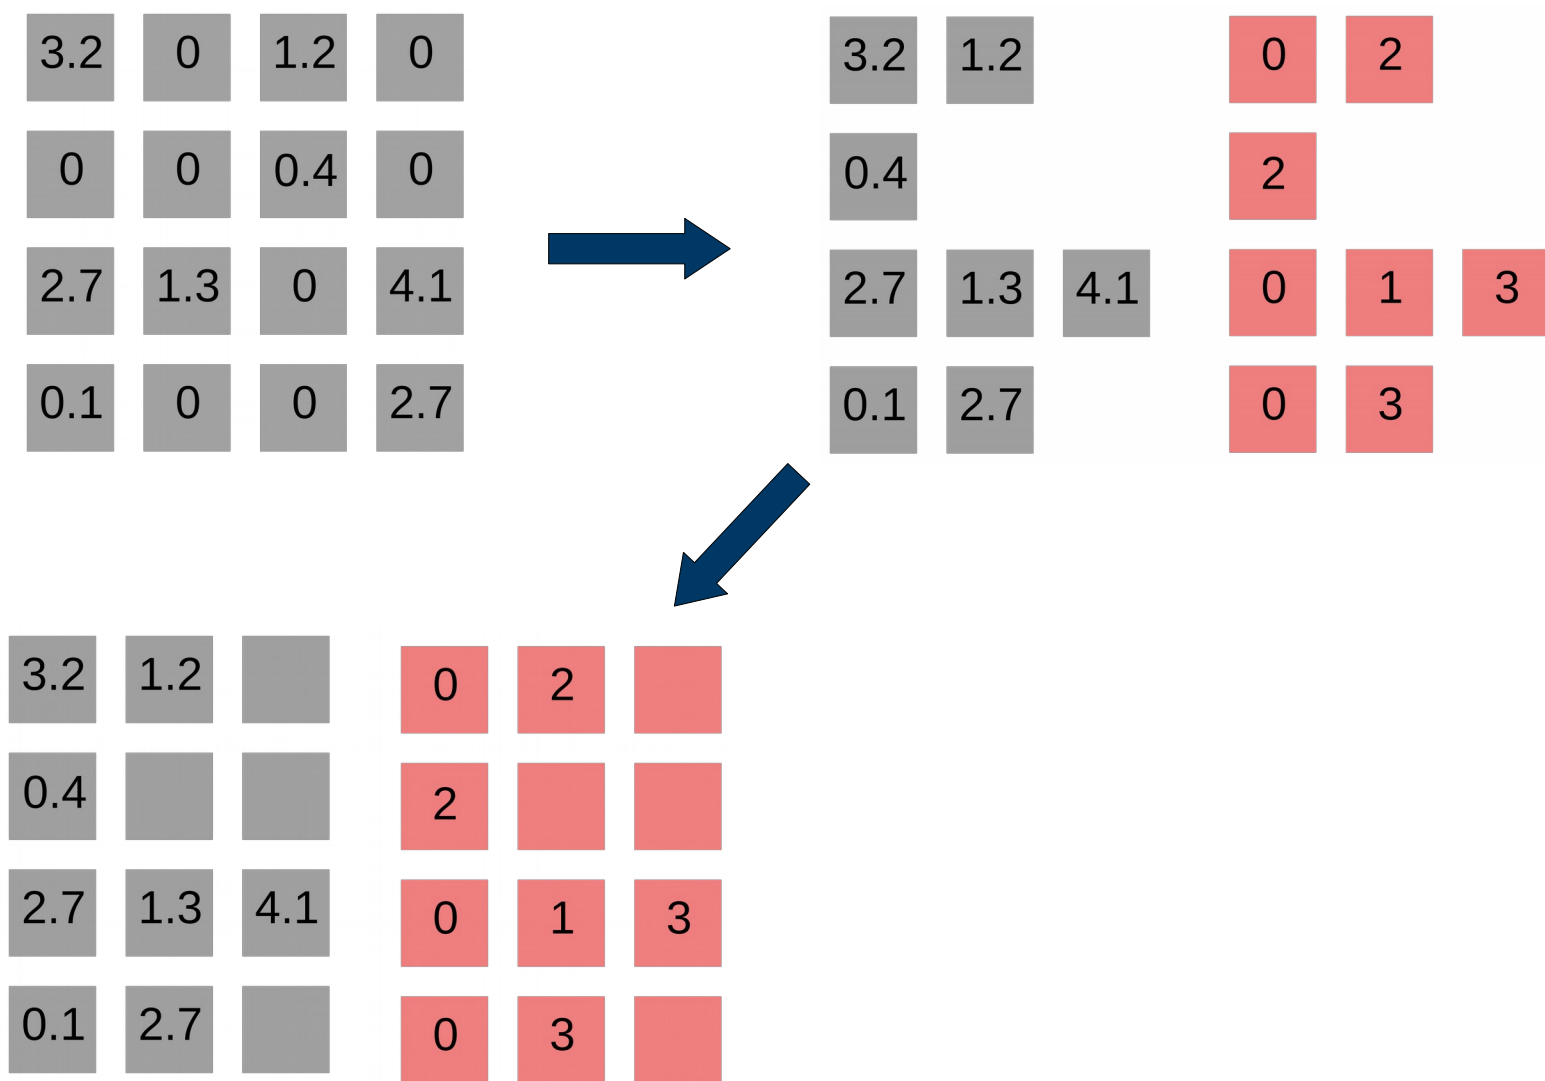
Sparse matrix formats

3.2	0	1.2	0
0	0	0.4	0
2.7	1.3	0	4.1
0.1	0	0	2.7



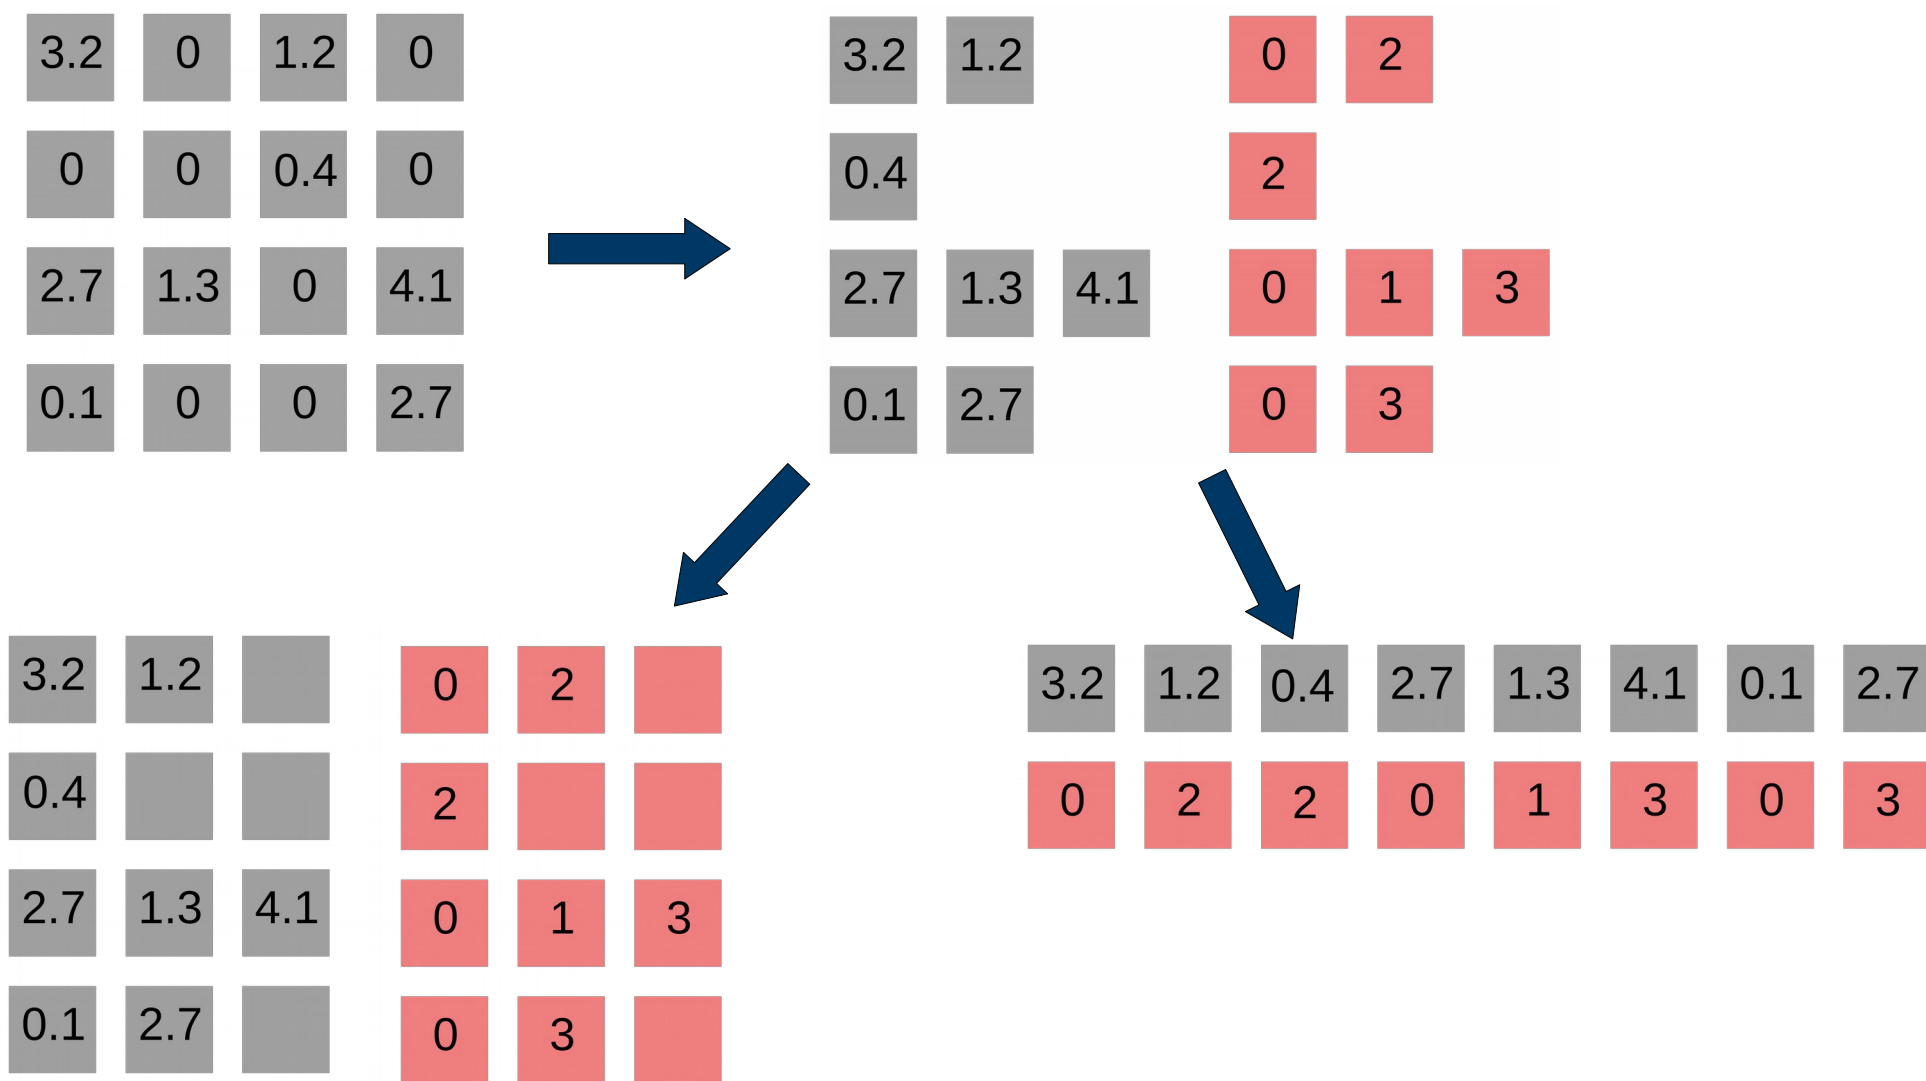
3.2	1.2			0	2	
0.4				2		
2.7	1.3	4.1		0	1	3
0.1	2.7			0	3	

Sparse matrix formats



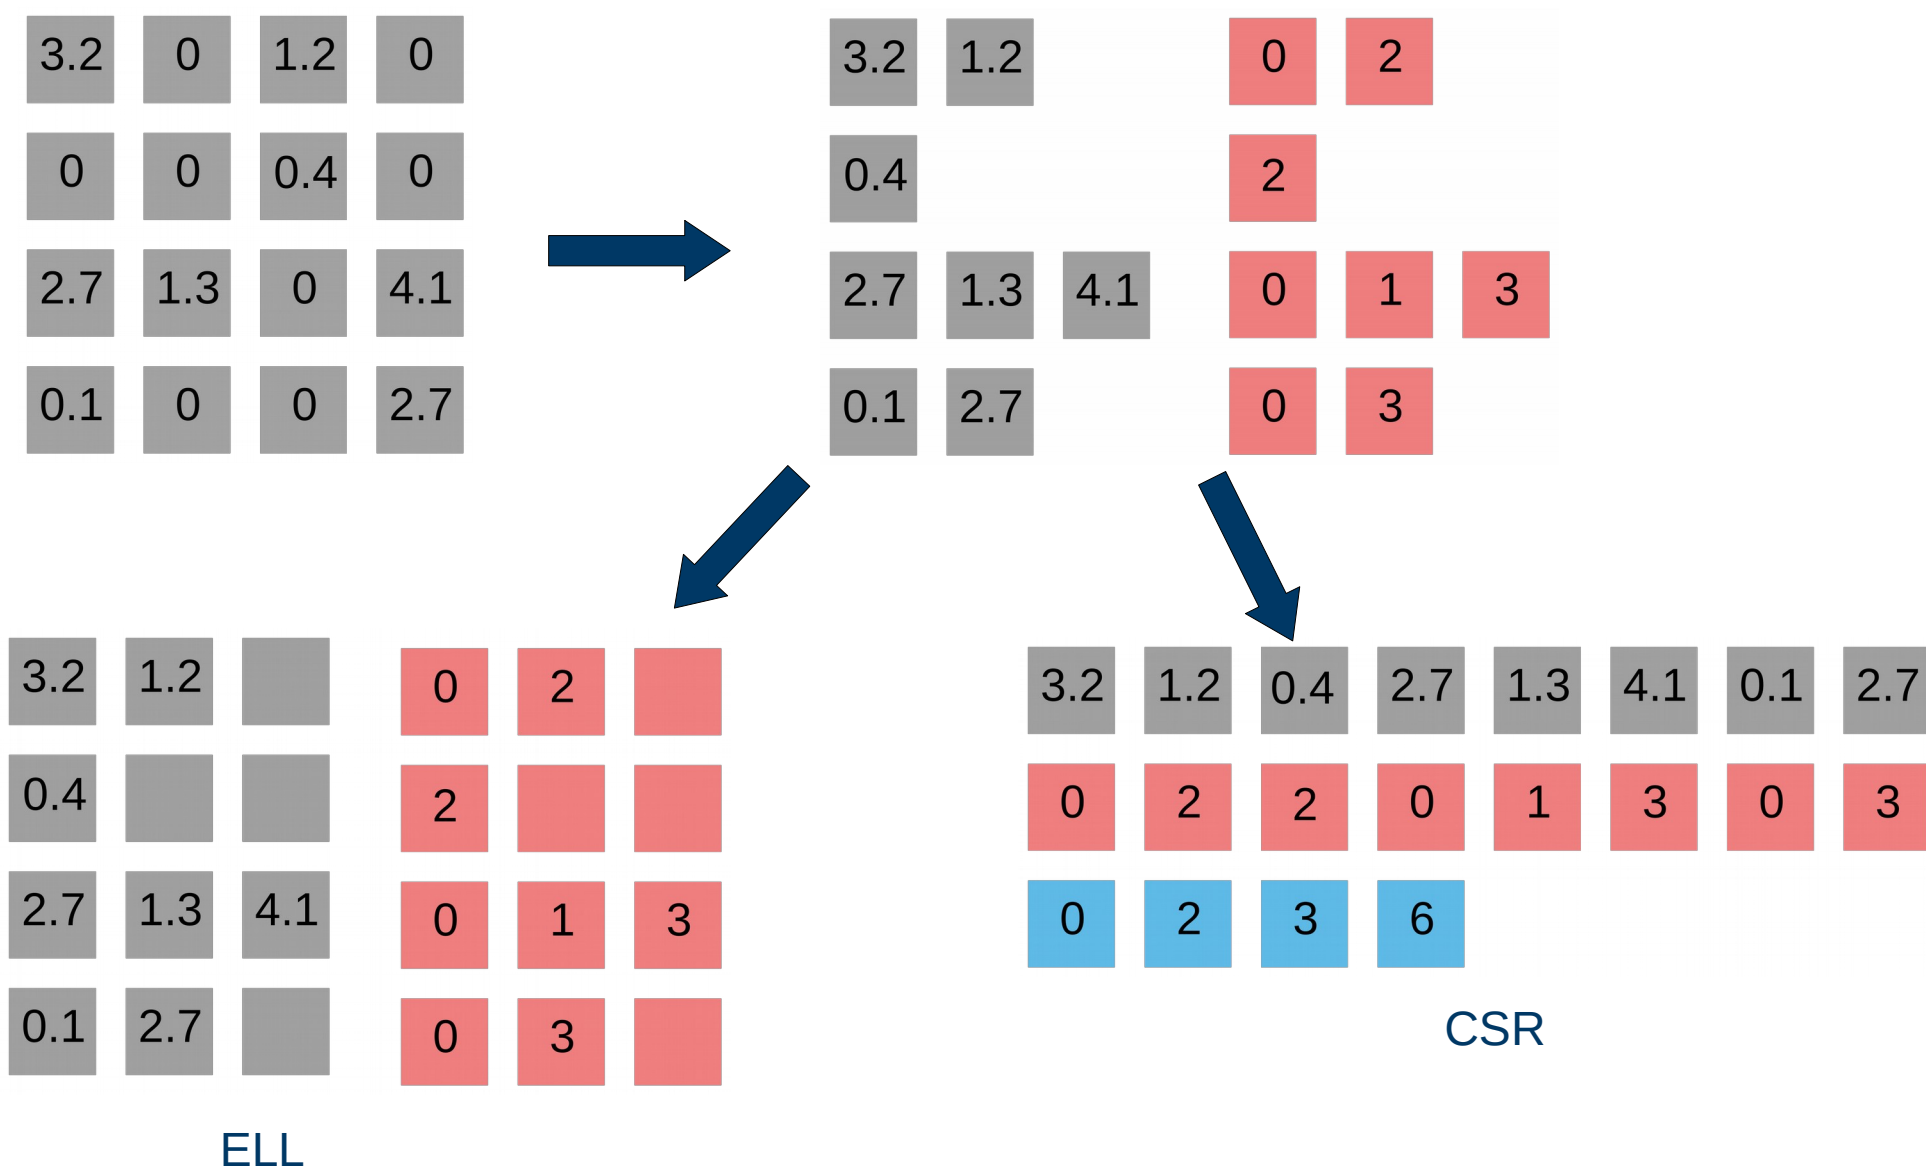
ELL

Sparse matrix formats



ELL

Sparse matrix formats



Sparse matrix formats

3.2	0	1.2	0
0	0	0.4	0
2.7	1.3	0	4.1
0.1	0	0	2.7



3.2
2.7
0.1

1.3

1.2

4.1

0.4

2.7

0

2

0

2

2

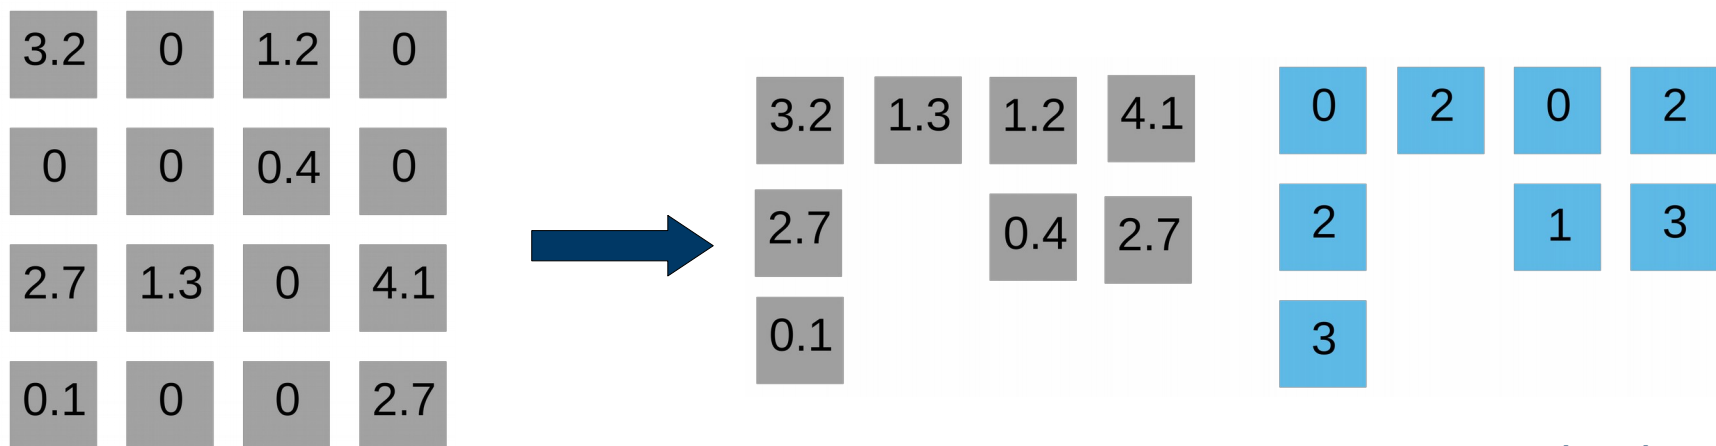
1

3

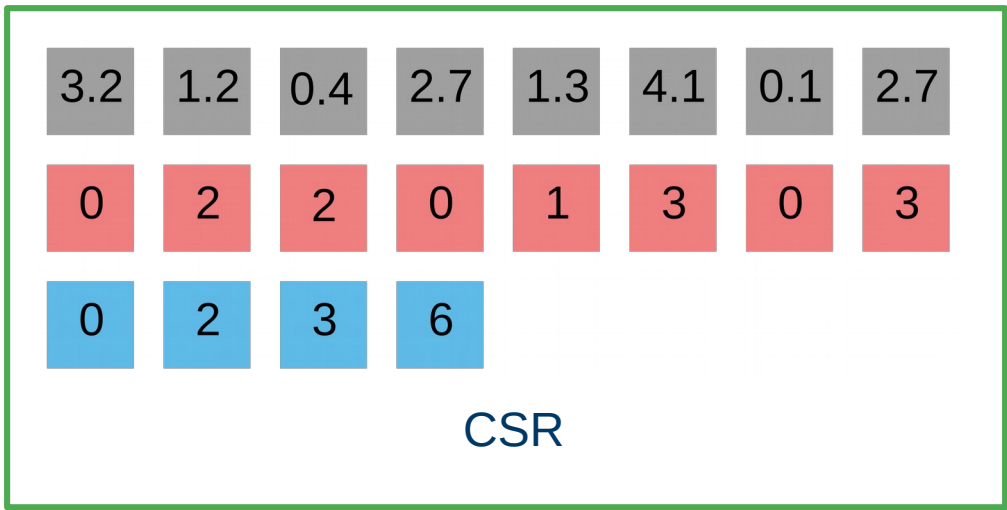
3

... leads to CSC

Sparse matrix formats



... leads to CSC

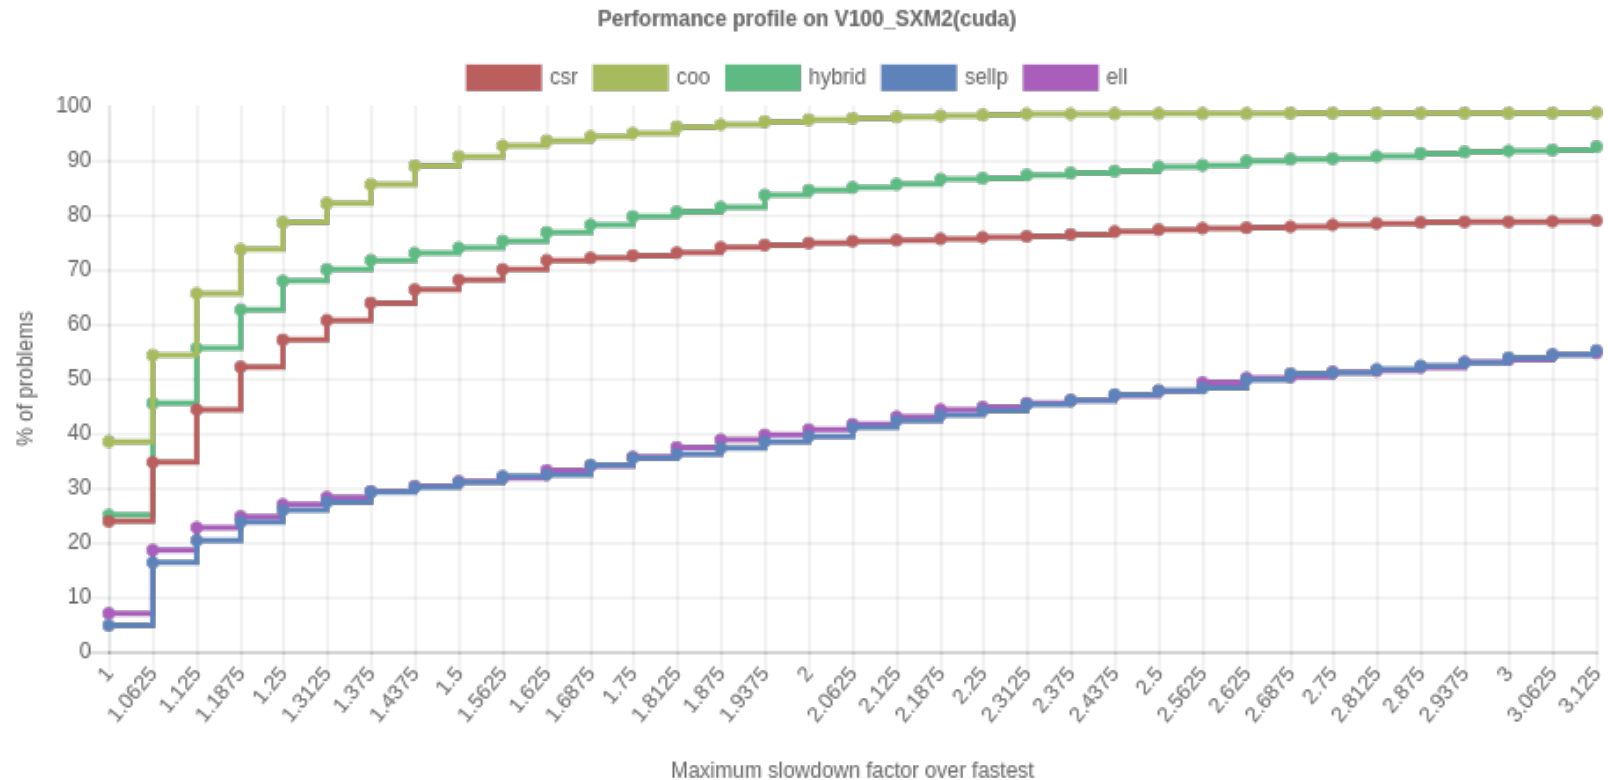


CSR

“Standard” approach

First things first

THERE IS NO “BEST” SPARSE MATRIX FORMAT / SpMV ALGORITHM



CSR SpMV

3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7
-----	-----	-----	-----	-----	-----	-----	-----

Values (val)

0	2	2	0	1	3	0	3
---	---	---	---	---	---	---	---

Column indexes (colidx)

0	2	3	6
---	---	---	---

Row pointers (rowptr)

CSR SpMV

3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7	Values (val)
0	2	2	0	1	3	0	3	Column indexes (colidx)
0	2	3	6					Row pointers (rowptr)

$$y := Ax$$

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

CSR SpMV

3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7	Values (val)
0	2	2	0	1	3	0	3	Column indexes (colidx)
0	2	3	6					Row pointers (rowptr)

$$y := Ax$$

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Bell & Garland '08

- parallelize outer loop

~ cuSPARSE SpMV

CSR SpMV

3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7	Values (val)
0	2	2	0	1	3	0	3	Column indexes (colidx)
0	2	3	6					Row pointers (rowptr)

$$y := Ax$$

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2   for (int i = 0; i < m; ++i) {  
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)  
4       y[i] += val[j] * x [ colidx[j] ];  
5   }  
6 }
```

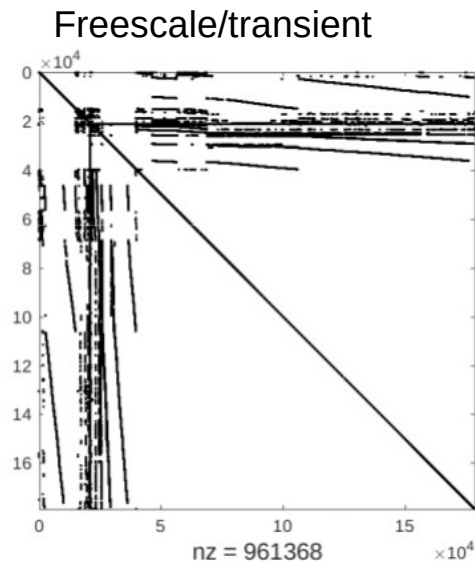
Bell & Garland '08

- parallelize outer loop

~ cuSPARSE SpMV

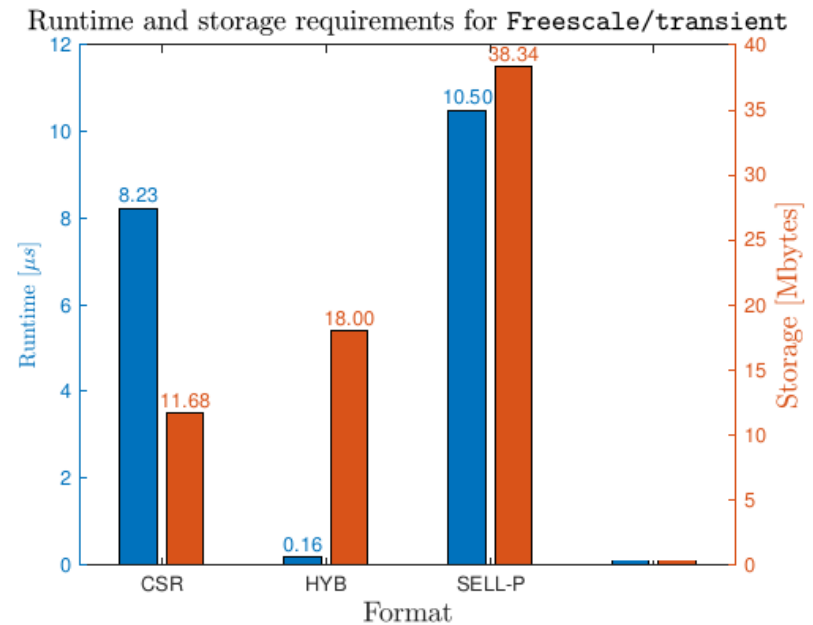
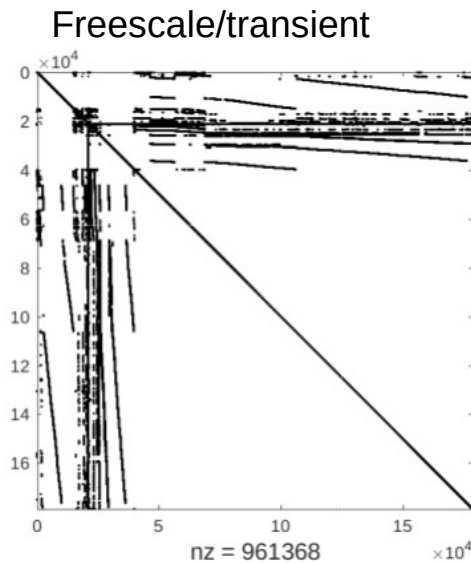
Load imbalance!

Example



Example

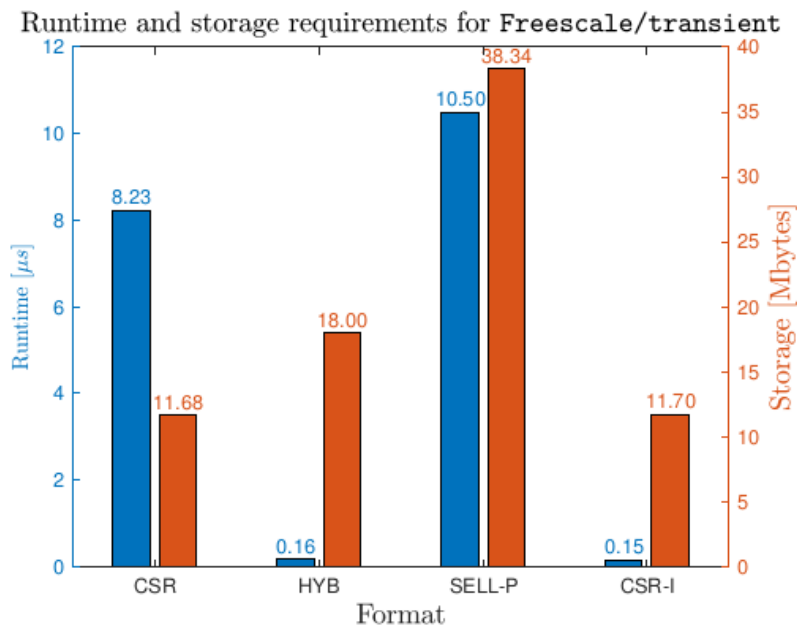
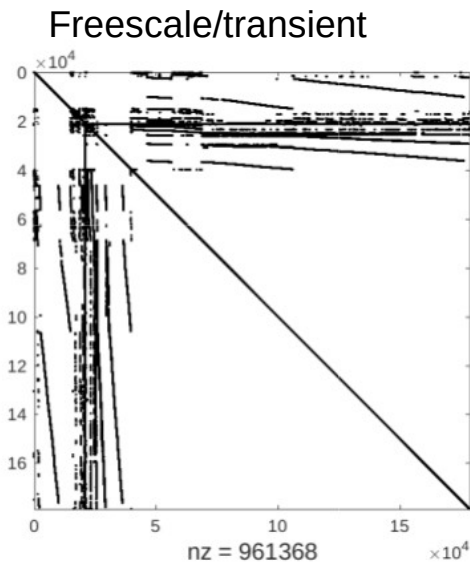
* GTX 1080



Can we do better than HYB using CSR?

Example

* GTX 1080



Can we do better than HYB using CSR?

55x speedup

YES!

Flegar, Anzt, *Overcoming Load Imbalance for Irregular Sparse Matrices*, IA3'17

Flegar, Quintana-Ortí, *Balanced CSR Sparse Matrix-Vector Product on Graphics Processors*, Euro-Par'17

How to do it

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2     for (int i = 0; i < m; ++i) {  
3         for (int j = rowptr[i]; j < rowptr[i+1]; ++j)  
4             y[i] += val[j] * x [ colidx[j] ];  
5     }  
6 }
```


How to do it

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2   for (int i = 0; i < m; ++i) {  
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)  
4       y[i] += val[j] * x[colidx[j]];  
5   }  
6 }
```



Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2   int row = -1, next_row = 0, nnz = rowptr[m];  
3   for (int i = 0; i < nnz; ++i) {  
4     while (i >= next_row) next_row = rowptr[++row+1];  
5     y[row] += val[i] * x[colidx[i]];  
6   }}
```

How to do it

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   int row = -1, next_row = 0, nnz = rowptr[m];
3   for (int i = 0; i < nnz; ++i) {
4     while (i >= next_row) next_row = rowptr[++row+1];
5     y[row] += val[i] * x[colidx[i]];
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
3   int row = -1, next_row = 0, nnz = rowptr[m];
4   for (int k = 0; k < T; ++k) {
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {
6       while (i >= next_row) next_row = rowptr[++row+1];
7       y[row] += val[i] * x[colidx[i]];
8     }}
```

How to do it

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   int row = -1, next_row = 0, nnz = rowptr[m];
3   for (int i = 0; i < nnz; ++i) {
4     while (i >= next_row) next_row = rowptr[++row+1];
5     y[row] += val[i] * x[colidx[i]];
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
3   int row = -1, next_row = 0, nnz = rowptr[m];
4   for (int k = 0; k < T; ++k) { Parallelize this!
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {
6       while (i >= next_row) next_row = rowptr[++row+1];
7       y[row] += val[i] * x[colidx[i]];
8     }}
```

How to do it

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2   for (int i = 0; i < m; ++i) {  
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)  
4       y[i] += val[j] * x[colidx[j]];  
5   }  
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2   int row = -1, next_row = 0, nnz = rowptr[m];  
3   for (int i = 0; i < nnz; ++i) {  
4     while (i >= next_row) next_row = rowptr[++row+1];  
5     y[row] += val[i] * x[colidx[i]];  
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;  
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
3   int row = -1, next_row = 0, nnz = rowptr[m];  
4   for (int k = 0; k < T; ++k) { Parallelize this!  
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {  
6       while (i >= next_row) next_row = rowptr[++row+1];  
7       y[row] += val[i] * x[colidx[i]];  
8     }}
```

Race conditions!

How to do it

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   int row = -1, next_row = 0, nnz = rowptr[m];
3   for (int i = 0; i < nnz; ++i) {
4     while (i >= next_row) next_row = rowptr[++row+1];
5     y[row] += val[i] * x[colidx[i]];
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
3   int row = -1, next_row = 0, nnz = rowptr[m];
4   for (int k = 0; k < T; ++k) {
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {
6       while (i >= next_row) next_row = rowptr[++row+1];
7       y[row] += val[i] * x[colidx[i]];
8     }}
```

Race conditions!

- Use atomics
- Accumulate partial result in the registers

How to do it

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   int row = -1, next_row = 0, nnz = rowptr[m];
3   for (int i = 0; i < nnz; ++i) {
4     while (i >= next_row) next_row = rowptr[++row+1];
5     y[row] += val[i] * x[colidx[i]];
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
3   int row = -1, next_row = 0, nnz = rowptr[m];
4   for (int k = 0; k < T; ++k) {
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {
6       while (i >= next_row) next_row = rowptr[++row+1];
7       y[row] += val[i] * x[colidx[i]];
8     }}
```

Parallelize this!

State between outer loop iterations!

Race conditions!

- Use atomics
- Accumulate partial result in the registers

How to do it

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   int row = -1, next_row = 0, nnz = rowptr[m];
3   for (int i = 0; i < nnz; ++i) {
4     while (i >= next_row) next_row = rowptr[++row+1];
5     y[row] += val[i] * x[colidx[i]];
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
3   int row = -1, next_row = 0, nnz = rowptr[m];
4   for (int k = 0; k < T; ++k) {
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {
6       while (i >= next_row) next_row = rowptr[++row+1];
7       y[row] += val[i] * x[colidx[i]];
8     }}
```

Race conditions!

- Use atomics
- Accumulate partial result in the registers

State between outer loop iterations!

Precompute starting value of “row” for each thread.

Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Spreading out threads causes strided memory access.

Assign one warp per chunk.

Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Spreading out threads causes strided memory access.

Assign one warp per chunk.

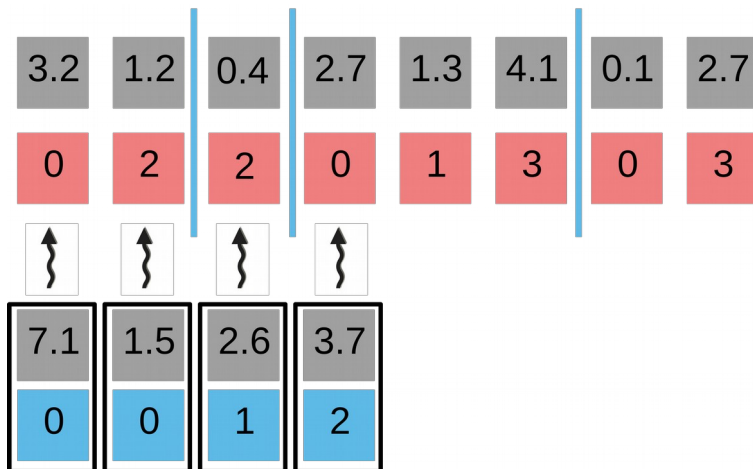
3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7
0	2	2	0	1	3	0	3

Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Spreading out threads causes strided memory access.

Assign one warp per chunk.

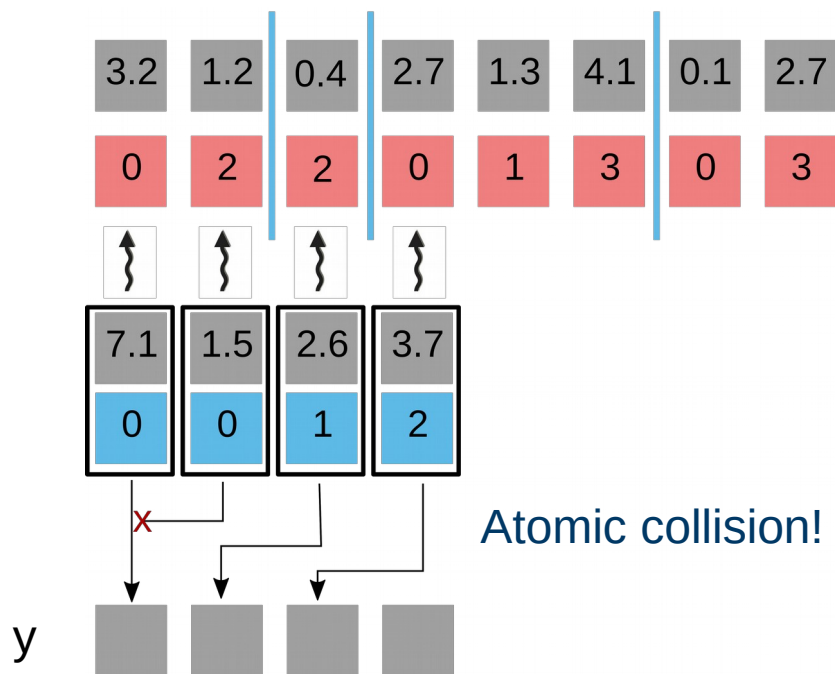


Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Spreading out threads causes strided memory access.

Assign one warp per chunk.

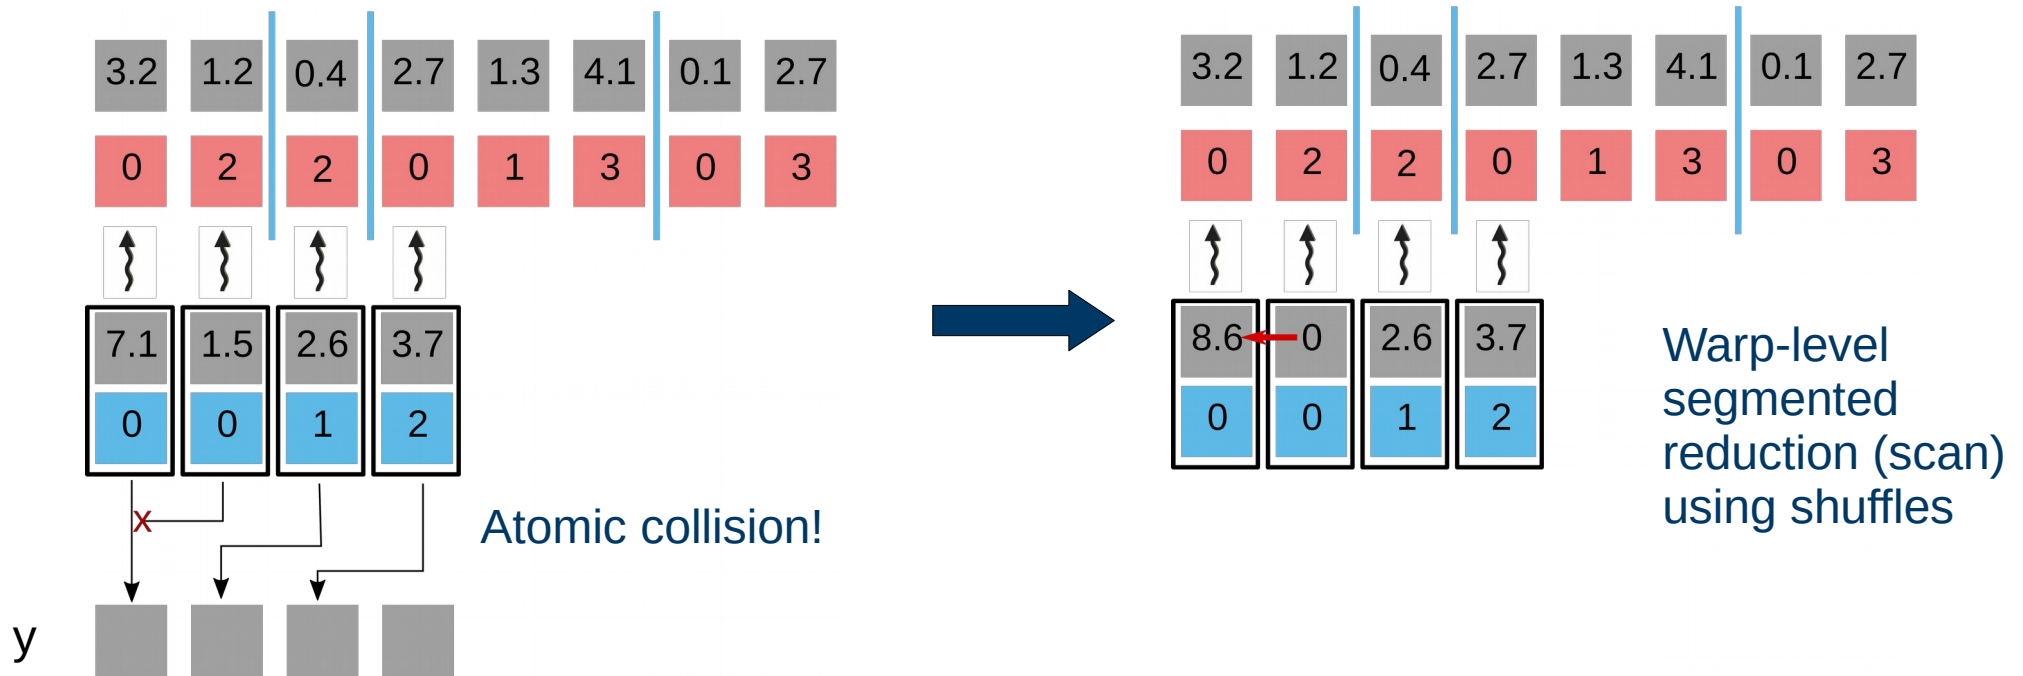


Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Spreading out threads causes strided memory access.

Assign one warp per chunk.

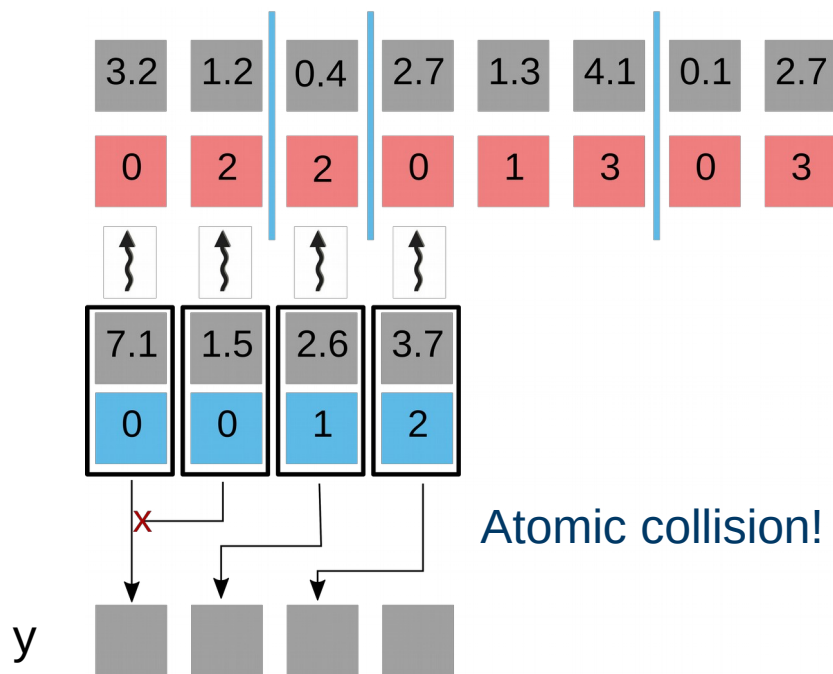


Getting good performance on GPUs

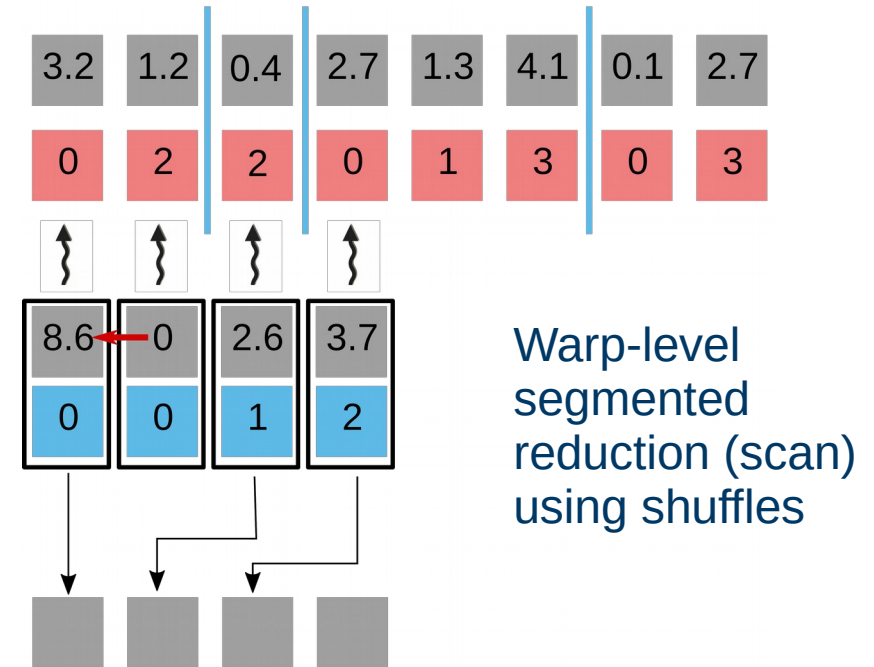
CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Spreading out threads causes strided memory access.

Assign one warp per chunk.



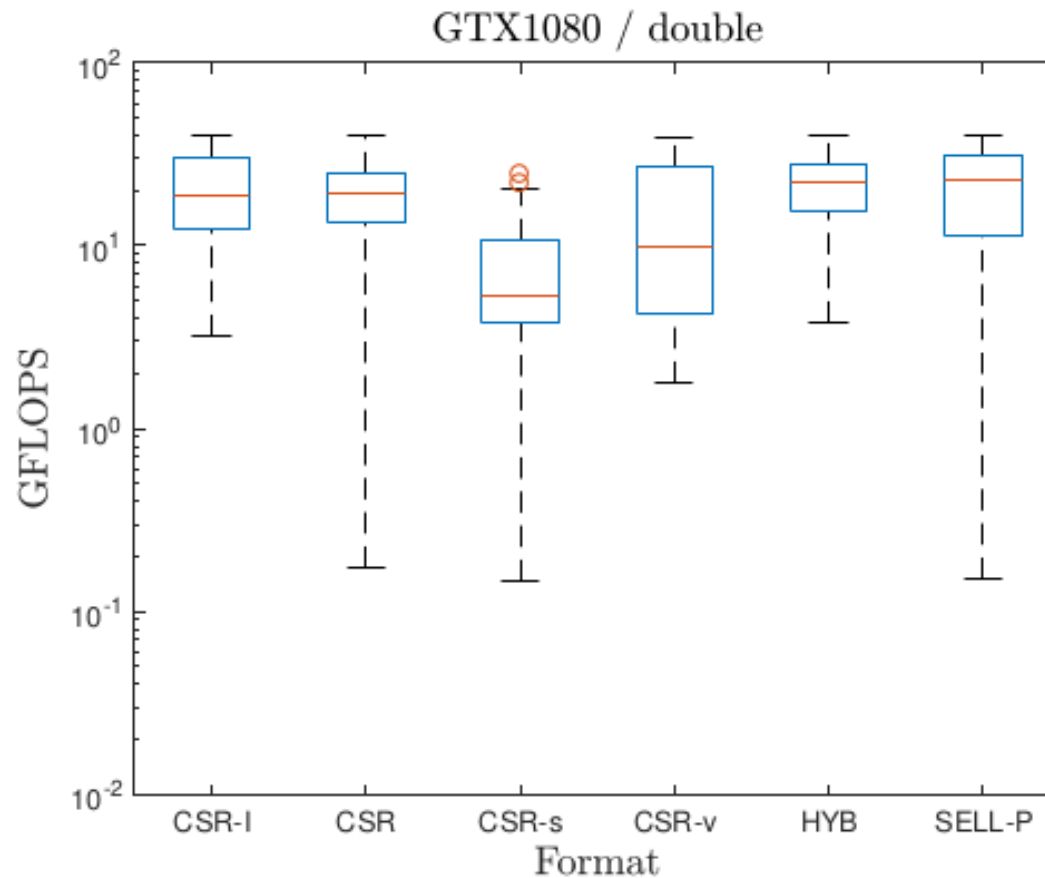
Atomic collision!



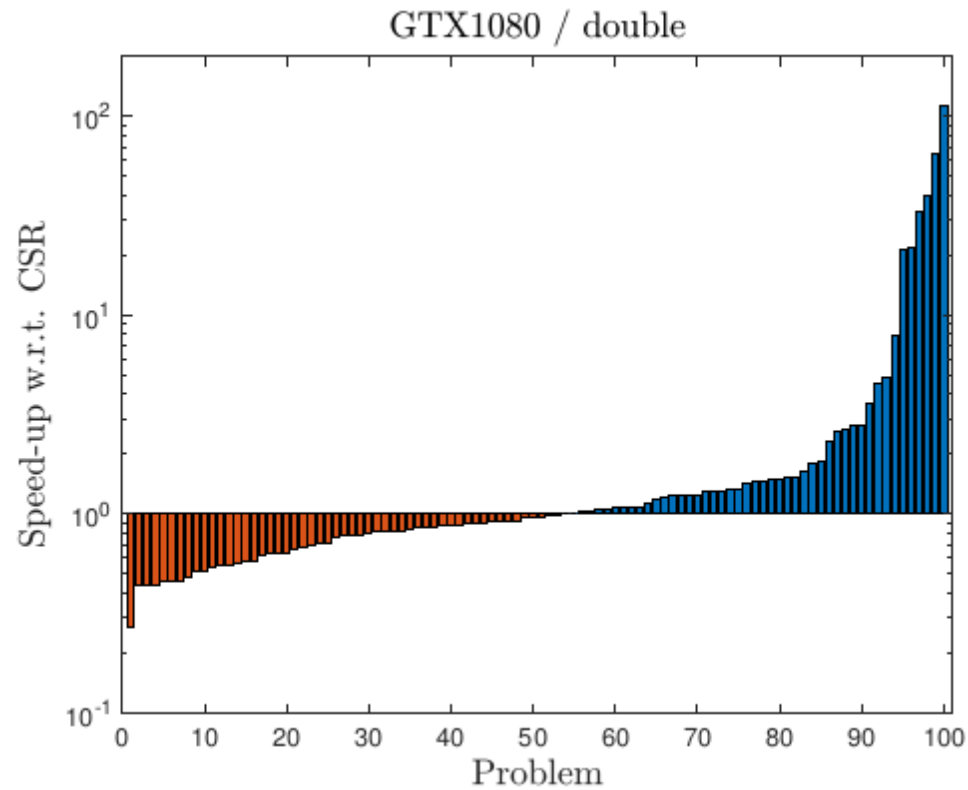
Warp-level segmented reduction (scan) using shuffles

Performance of CSR-I

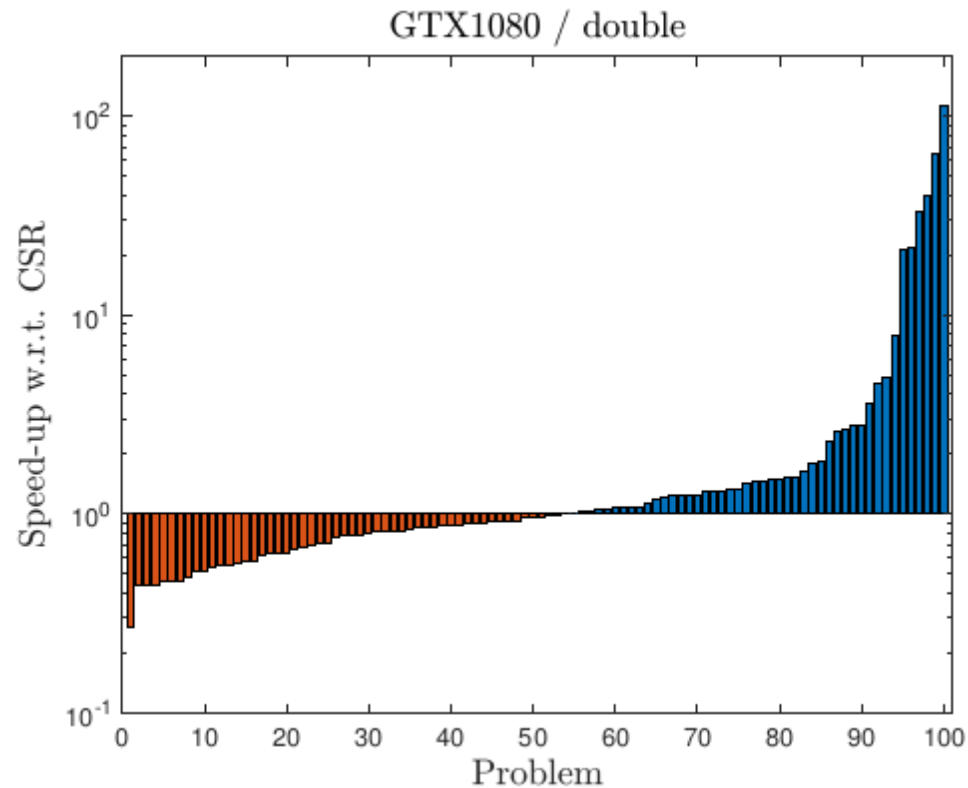
100 matrices from SuiteSparse



Speed-up / slowdown over cuSPARSE CSR



Speed-up / slowdown over cuSPARSE CSR



No format conversion!

- try both, and use the fastest later on!
- sometimes 1 cuSPARSE SpMV = 100 CSR-I SpMV

Choosing the winner a priori

CSR-I designed for irregular patterns

Choosing the winner a priori

CSR-I designed for irregular patterns

How to measure irregularity?

Deviation of row lengths from the mean.

Choosing the winner a priori

CSR-I designed for irregular patterns

How to measure irregularity?

Deviation of row lengths from the mean.

Is “5” regular or irregular?

Depends on the density of the matrix (mean #rows)

Choosing the winner a priori

CSR-I designed for irregular patterns

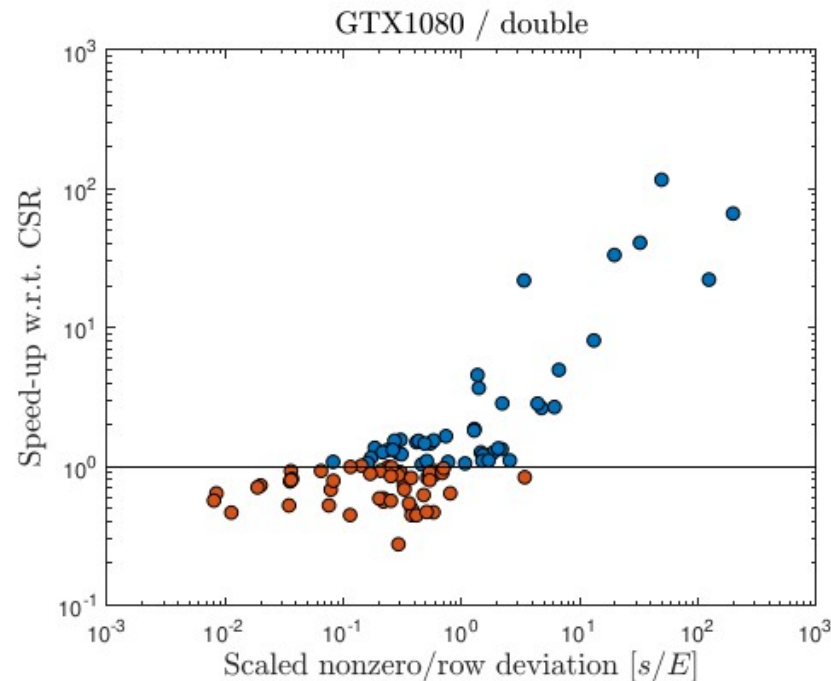
How to measure irregularity?

Deviation of row lengths from the mean.

Is “5” regular or irregular?

Depends on the density of the matrix (mean #rows)

Scatter plot of speedup vs normalized std. dev.



Choosing the winner a priori

CSR-I designed for irregular patterns

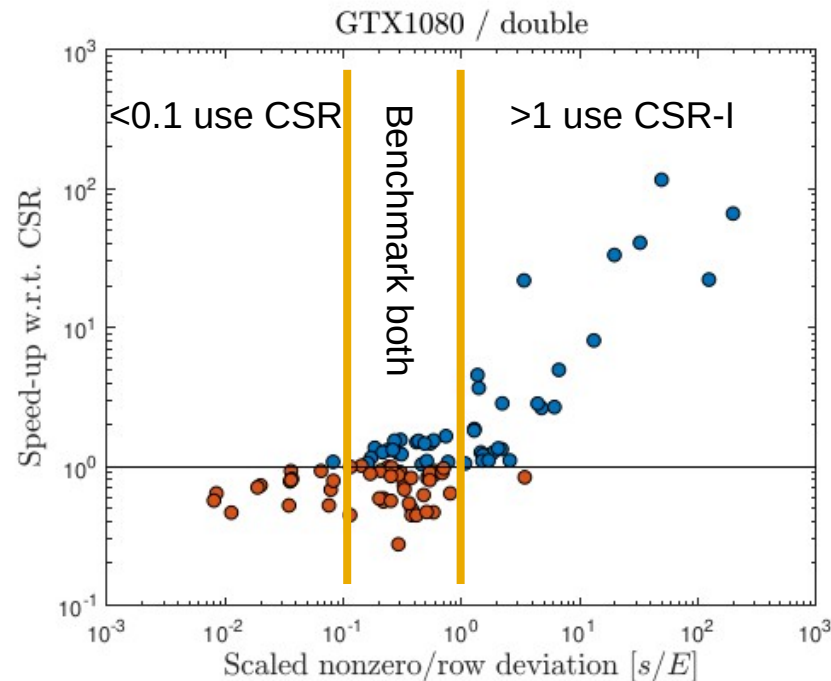
How to measure irregularity?

Deviation of row lengths from the mean.

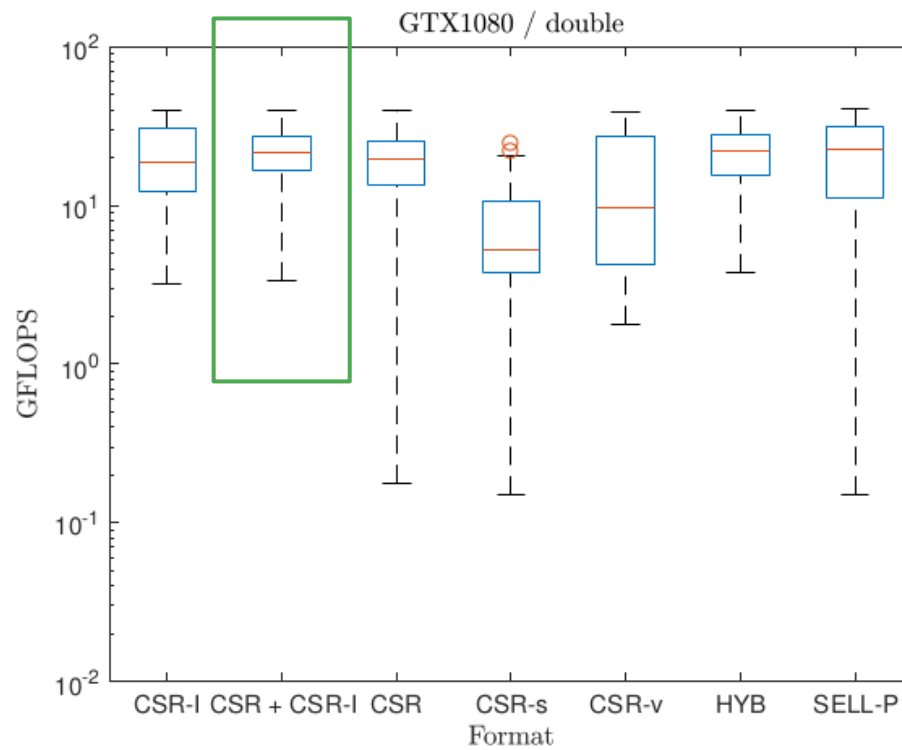
Is “5” regular or irregular?

Depends on the density of the matrix (mean #rows)

Scatter plot of speedup vs normalized std. dev.



Combining both approaches



Choosing the winner a priori

CSR-I designed for irregular patterns

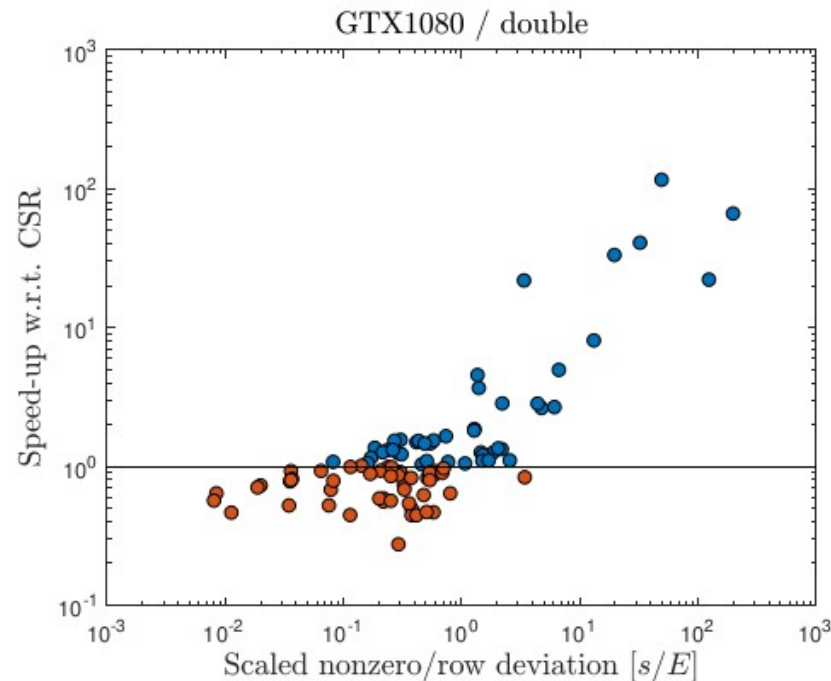
How to measure irregularity?

Deviation of row lengths from the mean.

Is “5” regular or irregular?

Depends on the density of the matrix (mean #rows)

Scatter plot of speedup vs normalized std. dev.



Choosing the winner a priori

CSR-I designed for irregular patterns

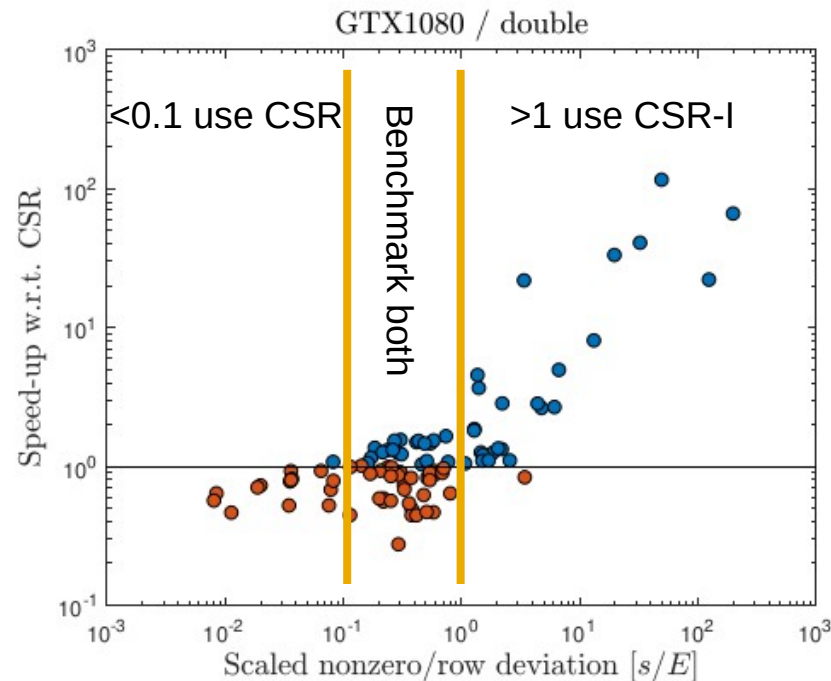
How to measure irregularity?

Deviation of row lengths from the mean.

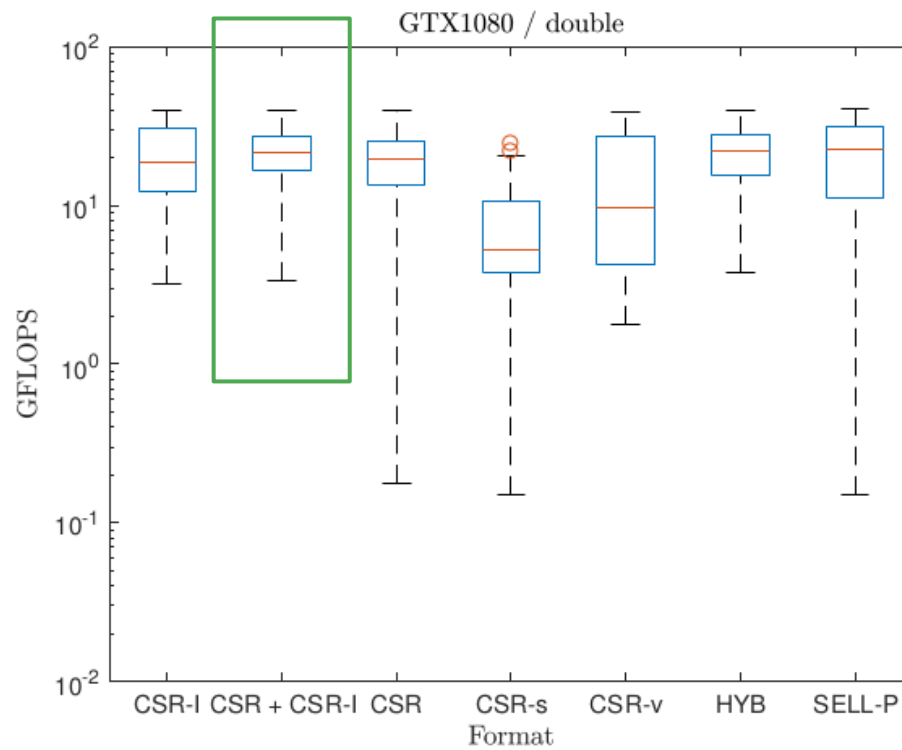
Is “5” regular or irregular?

Depends on the density of the matrix (mean #rows)

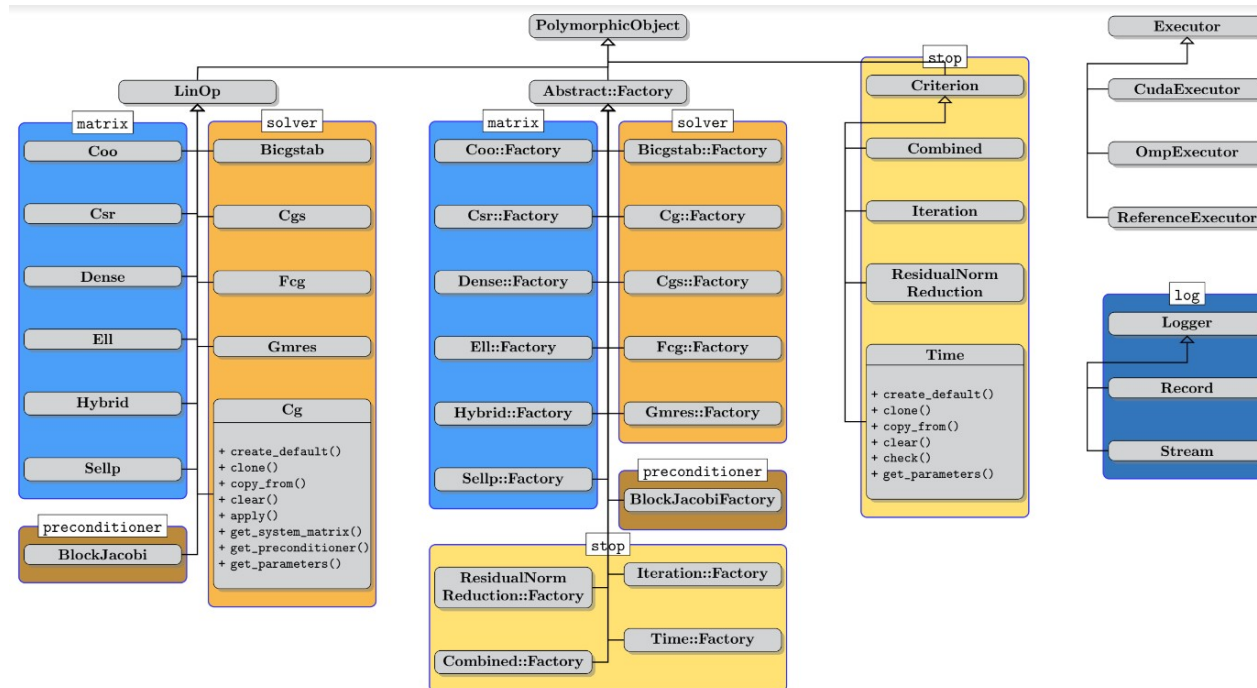
Scatter plot of speedup vs normalized std. dev.



Combining both approaches



Outlook



Choosing the correct combination of

matrix format

solver

preconditioner

... requires expert knowledge or significant trial and error.

Design a tool that does it (semi-)automatically?

Thank you! Questions?

- Krylov-subspace based linear solvers
 - Performance depends on **SpMV** and **Preconditioners**
- **Sparse matrix formats & SpMV**
 - There is no “holy grail” - best format depends on the problem
- **Preconditioners**
 - Accelerate the solution of the solver
 - Only an approximation, can use lower precision storage
- Do not implement from scratch, use a library. For example:



- github.com/ginkgo-project/ginkgo
- Open source license (BSD-3)
- Modern C++
- High performance GPU backend
- reference CPU implementation
- High performance CPU backend on the way

Slides:



github.com/gflegar/talks/tree/master/cmu_pittsburgh_2018_11

