



Optimizing GEMM for manycore architectures

Goran Flegar

SYCL BLAS team

25 August 2017

flegar@uji.es

General matrix-matrix product (GEMM)

$$C = \alpha \operatorname{op}_1(A) \operatorname{op}_2(B) + \beta C$$

- op_i is either identity or transpose
- α and β are scalars
- $\operatorname{op}_1(A)$ is m -by- k , $\operatorname{op}_2(B)$ is k -by- n , C is m -by- n (column major storage)

```
subroutine sgemm ( character      TRANSA,  
                  character      TRANSB,  
                  integer        M,  
                  integer        N,  
                  integer        K,  
                  real            ALPHA,  
                  real, dimension(Lda,*) A,  
                  integer        LDA,  
                  real, dimension(Ldb,*) B,  
                  integer        LDB,  
                  real            BETA,  
                  real, dimension(Ldc,*) C,  
                  integer        LDC  
)
```

General matrix-matrix product (GEMM)

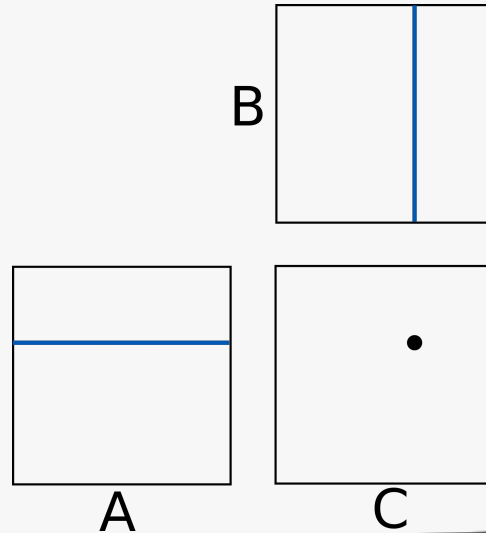
$$C = \alpha \text{op}_1(A) \text{op}_2(B) + \beta C$$

- op_i is either identity or transpose
- α and β are scalars
- $\text{op}_1(A)$ is m -by- k , $\text{op}_2(B)$ is k -by- n , C is m -by- n (column major storage)

```
subroutine sgemm ( character, TRANSA,  
                  character, TRANSB,  
                  integer, M,  
                  integer, N,  
                  integer, K,  
                  real, ALPHA,  
                  real, dimension(Lda,*) A,  
                  integer, LDA,  
                  real, dimension(Ldb,*) B,  
                  integer, LDB,  
                  real, BETA,  
                  real, dimension(Ldc,*) C,  
                  integer, LDC  
)
```

In this talk (for simplicity): $C = AB$

$$c_{ij} = \sum_{l=1}^k a_{il} b_{lj}$$



Naive implementation

$$c_{ij} = \sum_{l=1}^k a_{il} b_{lj}$$

Map one work item to each element of c_{ij} and loop over $a_{i\cdot}$ and $b_{\cdot j}$.

Naive implementation

$$c_{ij} = \sum_{l=1}^k a_{il} b_{lj}$$

Map one work item to each element of c_{ij} and loop over $a_{i\cdot}$ and $b_{\cdot j}$.



AMD R9 Nano

8.19 Tflop/s peak performance
512 GB/s (128 Gfloat/s) bandwidth

SYCL **C** ComputeCpp™



~200 Gflop/s

WHY?

4096-by-4096 matrices

Naive implementation

$$c_{ij} = \sum_{l=1}^k a_{il} b_{lj}$$

Map one work item to each element of c_{ij} and loop over $a_{i\cdot}$ and $b_{\cdot j}$.

Each work item:

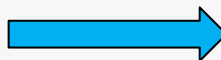
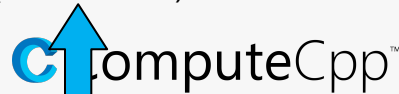
- $2k$ operations
- on $2k$ data elements



AMD R9 Nano

8.19 Tflop/s peak performance

512 GB/s (128 Gfloat/s) bandwidth



~200 Gflop/s



WHY?

Memory bounded kernel!

4096-by-4096 matrices

Naive implementation

$$c_{ij} = \sum_{l=1}^k a_{il} b_{lj}$$

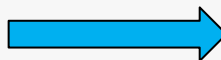
Map one work item to each element of c_{ij} and loop over $a_{i\cdot}$ and $b_{\cdot j}$.



AMD R9 Nano

8.19 Tflop/s peak performance

512 GB/s (128 Gfloat/s) bandwidth



~200 Gflop/s



WHY?

4096-by-4096 matrices

Each work item:

- 2k operations
- on 2k data elements

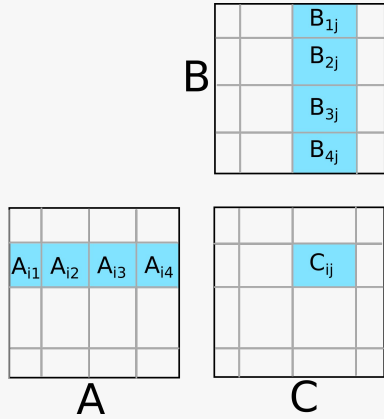
Memory bounded kernel!

Need to reuse data to “escape”
memory bandwidth barrier.

8192 : 128 = 64 : 1

* Need at least 64 operations
for each float fetched!

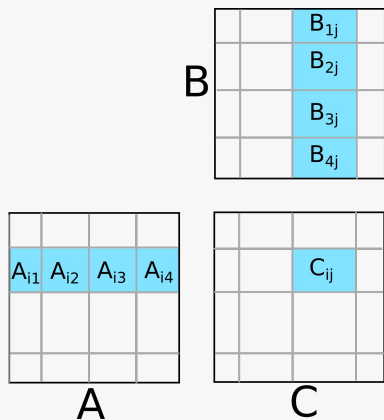
Block matrix multiplication



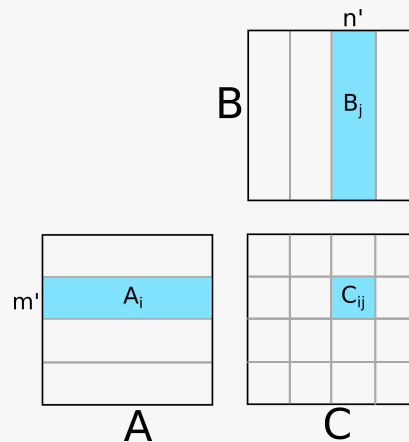
$$C_{ij} = \sum_{l=1}^K A_{il} B_{lj}$$

Block matrix multiplication

Special case: panel multiplication



$$C_{ij} = \sum_{l=1}^K A_{il} B_{lj}$$

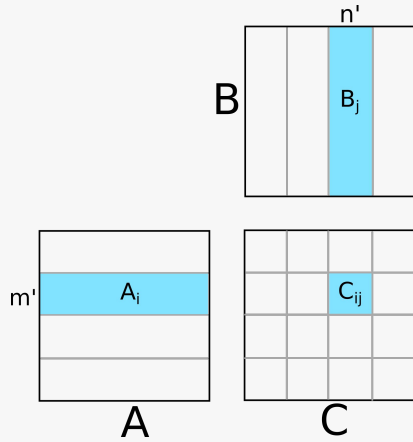


$$C_{ij} = A_i B_j$$

One work item per panel:

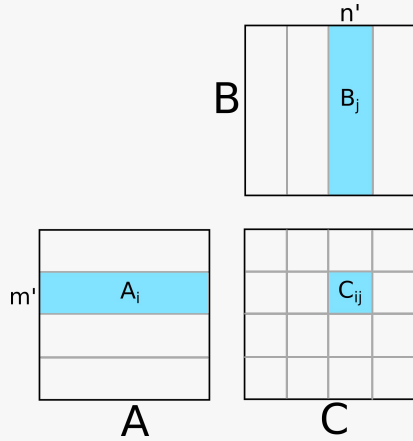
- $2m'n'k$ operations
- on $m'k + kn' + m'n'$ data entities

Maximizing data reuse

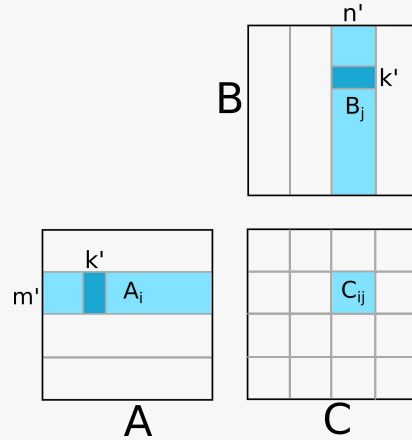


Cannot store the whole panel in caches /
local memory / registers

Maximizing data reuse



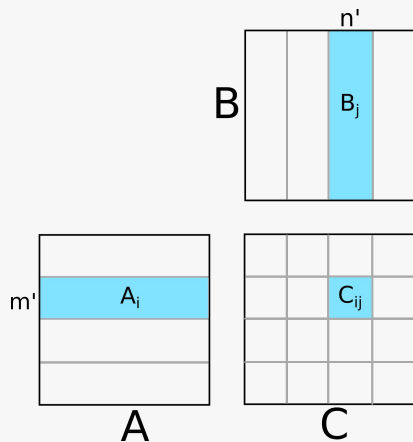
Cannot store the whole panel in caches /
local memory / registers



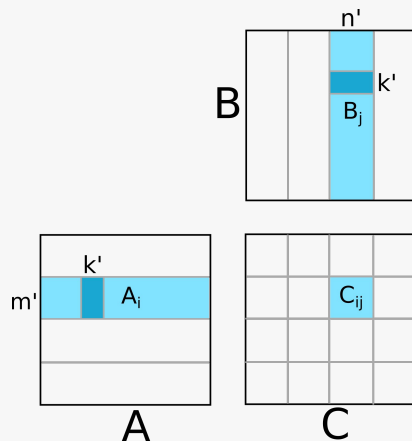
Instead break it into blocks

- Keep C_{ij} in registers
- Load a single *block* of A and B
 - $m'k' + k'n'$ data
- Compute a small gemm with these blocks and add the result to C_{ij}
 - $2m'n'k'$ operations
- Repeat the process for next block

Maximizing data reuse



Cannot store the whole panel in caches / local memory / registers



Instead break it into blocks

- Keep C_{ij} in registers
- Load a single *block* of A and B
 - $m'k' + k'n'$ data
- Compute a small gemm with these blocks and add the result to C_{ij}
 - $2m'n'k'$ operations
- Repeat the process for next block

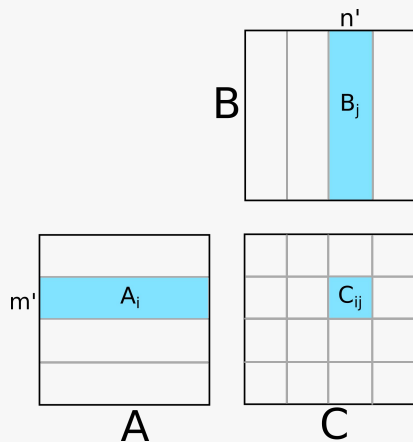
Data reuse:

$$\frac{2m'n'k'}{m'k' + n'k'} = \frac{2m'n'}{m' + n'}$$

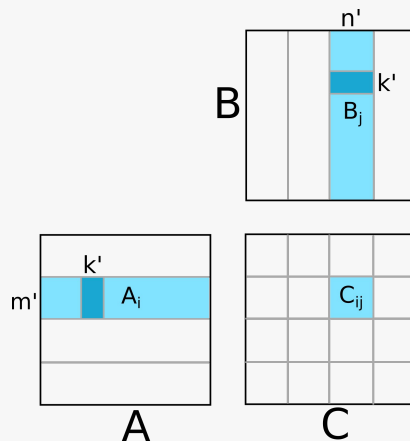
#registers:

$$m'n' + m'k' + k'n'$$

Maximizing data reuse



Cannot store the whole panel in caches / local memory / registers



Instead break it into blocks

Limited amount of registers:

- use k' as small as possible, keeping in mind good memory access
 - (k' = “cache line size”)
- $m' = n'$ is the best choice for constrained number of registers
 - “data reuse” = m'

- Keep C_{ij} in registers
- Load a single *block* of A and B
 - $m'k' + k'n'$ data
- Compute a small gemm with these blocks and add the result to C_{ij}
 - $2m'n'k'$ operations
- Repeat the process for next block

Data reuse:

$$\frac{2m'n'k'}{m'k' + n'k'} = \frac{2m'n'}{m' + n'}$$

#registers:

$$m'n' + m'k' + k'n'$$

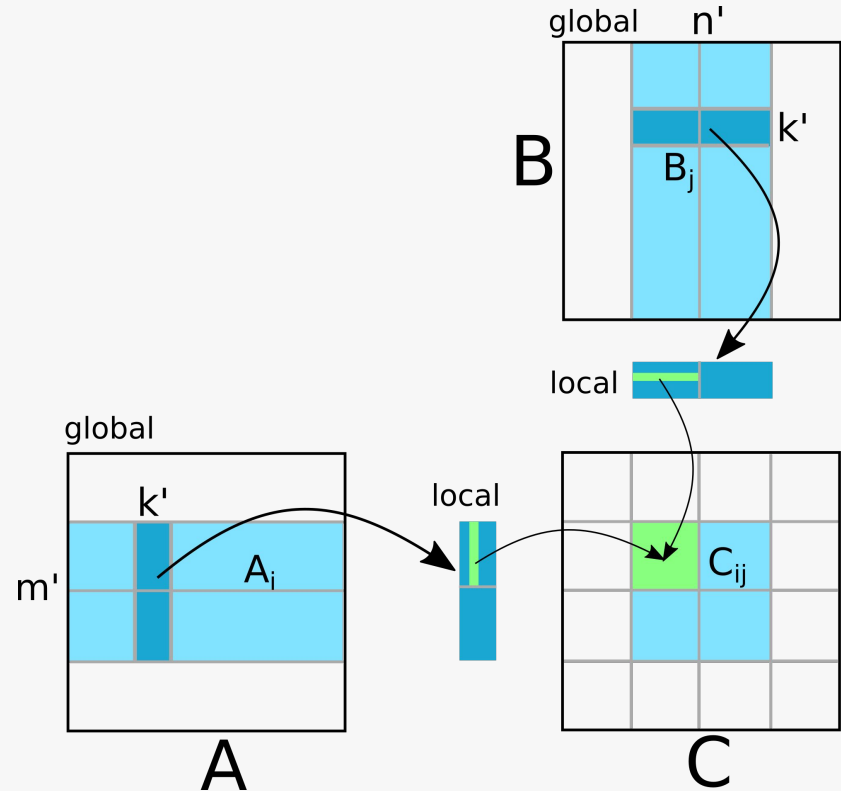
R9 Nano:

- “data reuse” = 8

Collaborate to increase effective data reuse

One work item has only a small amount of available registers.

- Combine the registers of entire workgroup to get more register space.
- Each work item stores only one sub-block of C_{ij} .
- All work items collaborate when reading to local memory.
- Each work item reads from local memory the part it needs.



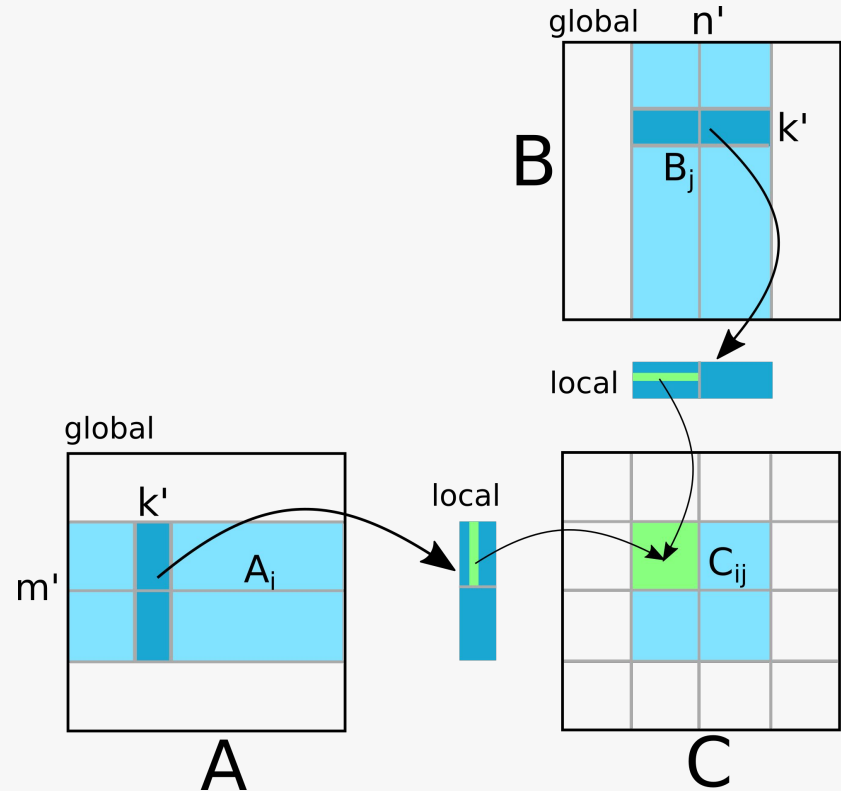
Collaborate to increase effective data reuse

One work item has only a small amount of available registers.

- Combine the registers of entire workgroup to get more register space.
- Each work item stores only one sub-block of C_{ij} .
- All work items collaborate when reading to local memory.
- Each work item reads from local memory the part it needs.

R9 Nano:

- Work group size: 16x16 items
- “local data reuse” = 8
- “global data reuse” = 128



Further optimizations

Memory bandwidth no longer an issue.

Focus on decreasing the volume of “useless” arithmetic instructions.

- Address calculation.
- Bound checking.

Further optimizations

Memory bandwidth no longer an issue.

Focus on decreasing the volume of “useless” arithmetic instructions.

- **Address calculation.**
- Bound checking.

```
template <typename T, typename TernaryOperator>
void matrix_for_each(int m, int n, T *p, int ld, TernaryOperator op) {
    for (int j = 0; j < n; ++j) {
        for (int i = 0; i < m; ++i) {
            op(i, j, p[i + j*ld]);
        }
    }
}
```

Further optimizations


Memory bandwidth no longer an issue.

Focus on decreasing the volume of “useless” arithmetic instructions.

- **Address calculation.**
- Bound checking.

```
template <typename T, typename TernaryOperator>
void matrix_for_each(int m, int n, T *p, int ld, TernaryOperator op) {
    for (int j = 0; j < n; ++j) {
        for (int i = 0; i < m; ++i) {
            op(i, j, p[i + j*ld]);
        }
    }
}
```

Introducing “matrix” abstractions might be tempting, but can have significant overhead.



```
template <typename Matrix, typename TernaryOperator>
void matrix_for_each(Matrix &M, TernaryOperator op) {
    for (int j = 0; j < M.get_num_cols(); ++j) {
        for (int i = 0; i < M.get_num_rows(); ++i) {
            op(i, j, M(i,j));
        }
    }
}
```

Further optimizations

Memory bandwidth no longer an issue.

Focus on decreasing the volume of “useless” arithmetic instructions.

- **Address calculation.**
- Bound checking.

Introducing “matrix” abstractions might be tempting, but can have significant overhead.

```
template <typename T, typename TernaryOperator>
void matrix_for_each(int m, int n, T *p, int ld, TernaryOperator op) {
    for (int j = 0; j < n; ++j) {
        for (int i = 0; i < m; ++i) {
            op(i, j, p[i + j*ld]);
        }
    }
}
```

$3mn$ arithmetic op.

Calculate partial addresses.

```
template <typename Matrix, typename TernaryOperator>
void matrix_for_each(Matrix &M, TernaryOperator op) {
    for (int j = 0; j < M.get_num_cols(); ++j) {
        for (int i = 0; i < M.get_num_rows(); ++i) {
            op(i, j, M(i,j));
        }
    }
}
```

```
template <typename T, typename TernaryOperator>
void matrix_for_each(int m, int n, T *p, int ld, TernaryOperator op) {
    for (int j = 0; j < n; ++j) {
        for (int i = 0; i < m; ++i) {
            op(i, j, p[i]);
        }
        p += ld;
    }
}
```

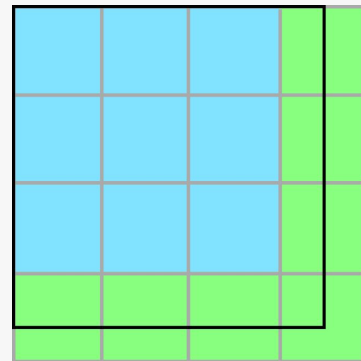
$(m+1)n$ arithmetic op.

Further optimizations

Memory bandwidth no longer an issue.

Focus on decreasing the volume of “useless” arithmetic instructions.

- Address calculation.
- **Bound checking.**
 - Skip bound checking in **internal tiles**.
 - Bound check in **external tiles**.



C

Naive implementation

$$c_{ij} = \sum_{l=1}^k a_{il} b_{lj}$$

Map one work item to each element of c_{ij} and loop over $a_{i\cdot}$ and $b_{\cdot j}$.

Map one work group per block of C + further optimizations (16-by-16 work group, with 8-by-8 sub-block per work item)

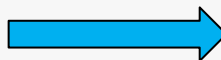


AMD R9 Nano

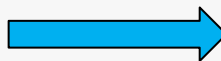
8 Tflop/s peak performance

500 GB/s (125 Gfloat/s) bandwidth

SYCL **C** ComputeCpp™



~ 200 Gflop/s



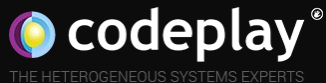
~ 4 Tflop/s

4096-by-4096 matrices

What next?

- Vectorization: possible performance improvement with vectorized access (vload / vstore).
- Different matrix sizes: If C is small, the number of matrix blocks might be too small to utilize the GPU.
 - Use smaller blocks? Less data reuse!
 - Use multiple work groups per block? Race conditions! (need atomic operations)
- Implement other BLAS 3 routines (optimization ideas should be similar)

We're
Hiring!
codeplay.com/careers/



Thank you! Questions?



[@codeplaysoft](https://twitter.com/codeplaysoft)



info@codeplay.com



codeplay.com