

Balanced CSR Sparse Matrix-Vector Product on Graphics Processors

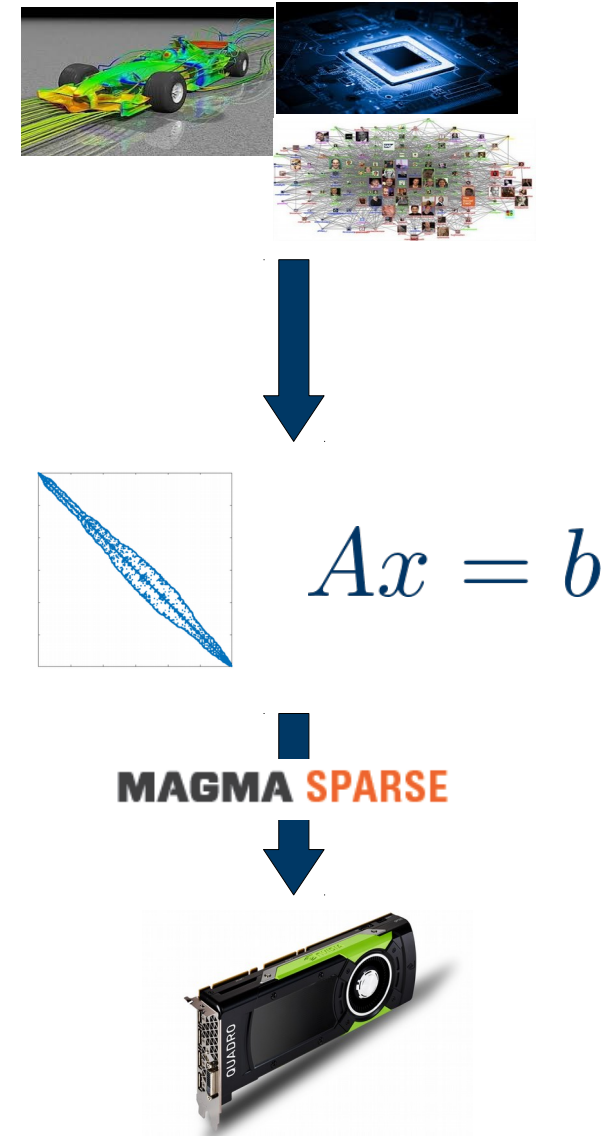
Goran Flegar, Enrique S. Quintana-Ortí



Scan me
for slides!

MAGMA-sparse software library

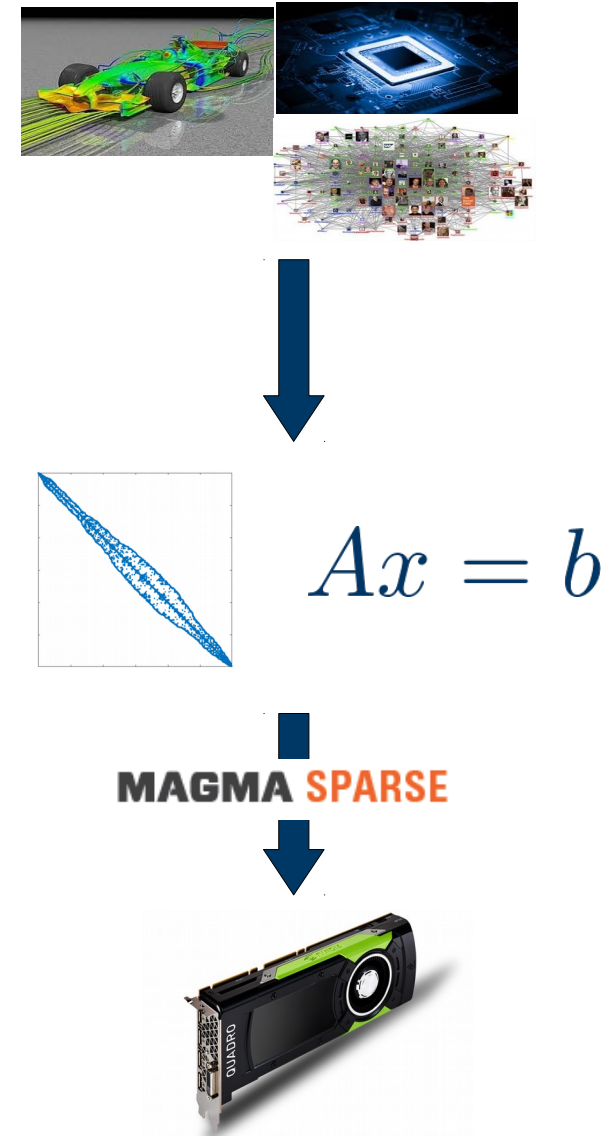
- GPU-accelerated sparse linear algebra library
 - Focus: linear systems



Joint effort: Innovative Computing Lab at University of Tennessee, Knoxville; Karlsruhe Institute of Technology; Universidad Jaume I

MAGMA-sparse software library

- GPU-accelerated sparse linear algebra library
 - Focus: linear systems
 - Iterative, Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations

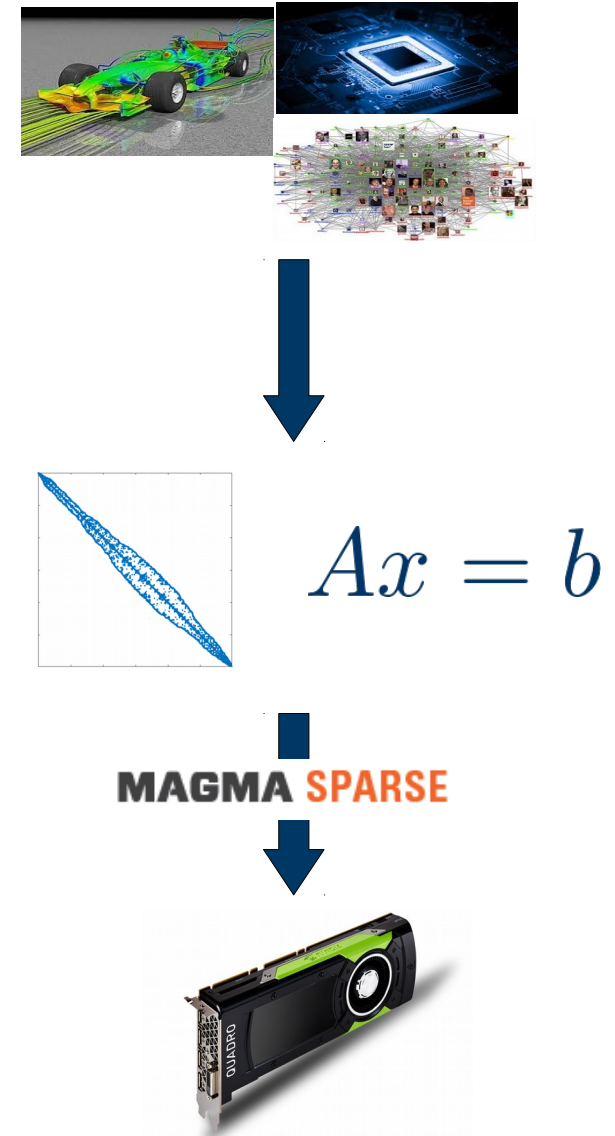


Joint effort: Innovative Computing Lab at University of Tennessee, Knoxville; Karlsruhe Institute of Technology; University Jaume I



MAGMA-sparse software library

- GPU-accelerated sparse linear algebra library
 - Focus: linear systems
 - Iterative, Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver

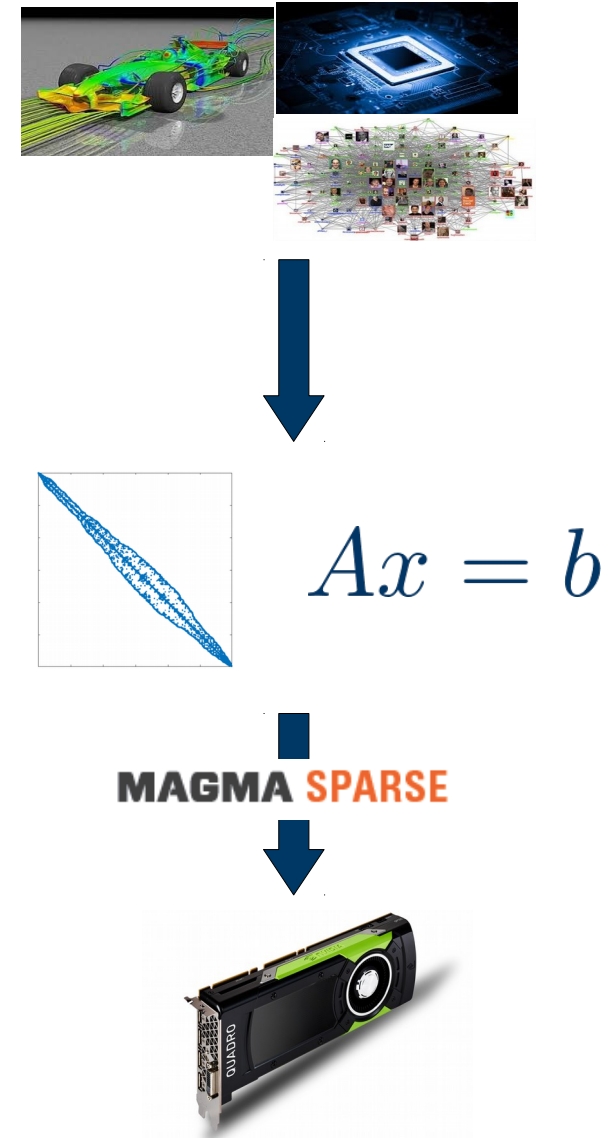


Joint effort: Innovative Computing Lab at University of Tennessee, Knoxville; Karlsruhe Institute of Technology; University Jaume I



MAGMA-sparse software library

- GPU-accelerated sparse linear algebra library
 - Focus: linear systems
 - Iterative, Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver
 - Preconditioners
 - reduce the number of iterations

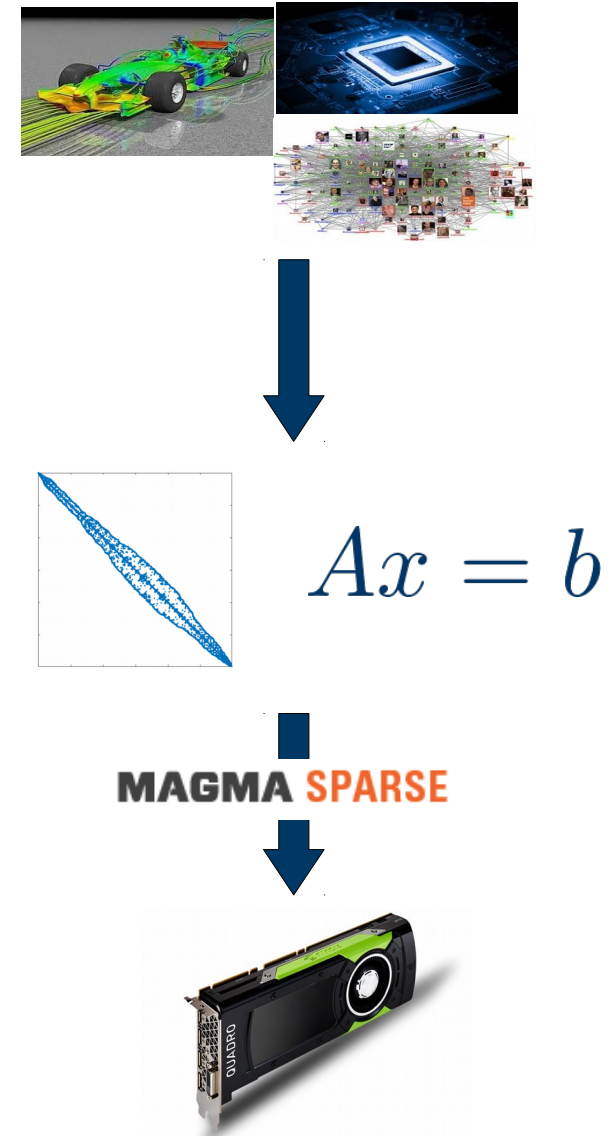


Joint effort: Innovative Computing Lab at University of Tennessee, Knoxville; Karlsruhe Institute of Technology; University Jaume I



MAGMA-sparse software library

- GPU-accelerated sparse linear algebra library
 - Focus: linear systems
 - Iterative, Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver
 - Preconditioners
 - reduce the number of iterations

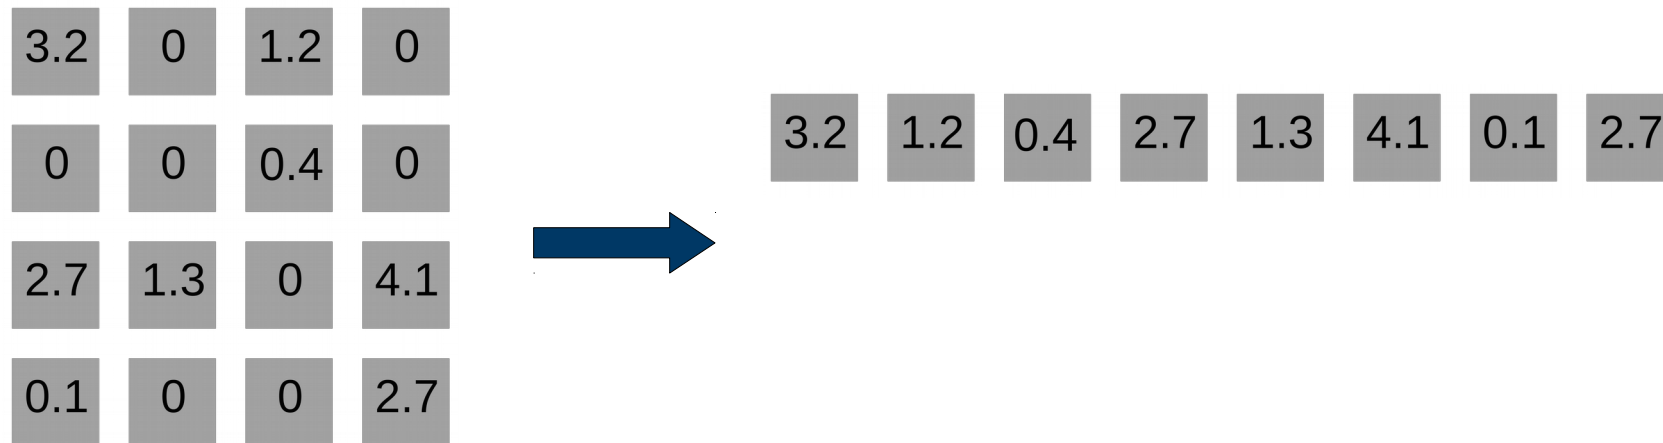


Joint effort: Innovative Computing Lab at University of Tennessee, Knoxville; Karlsruhe Institute of Technology; University Jaume I

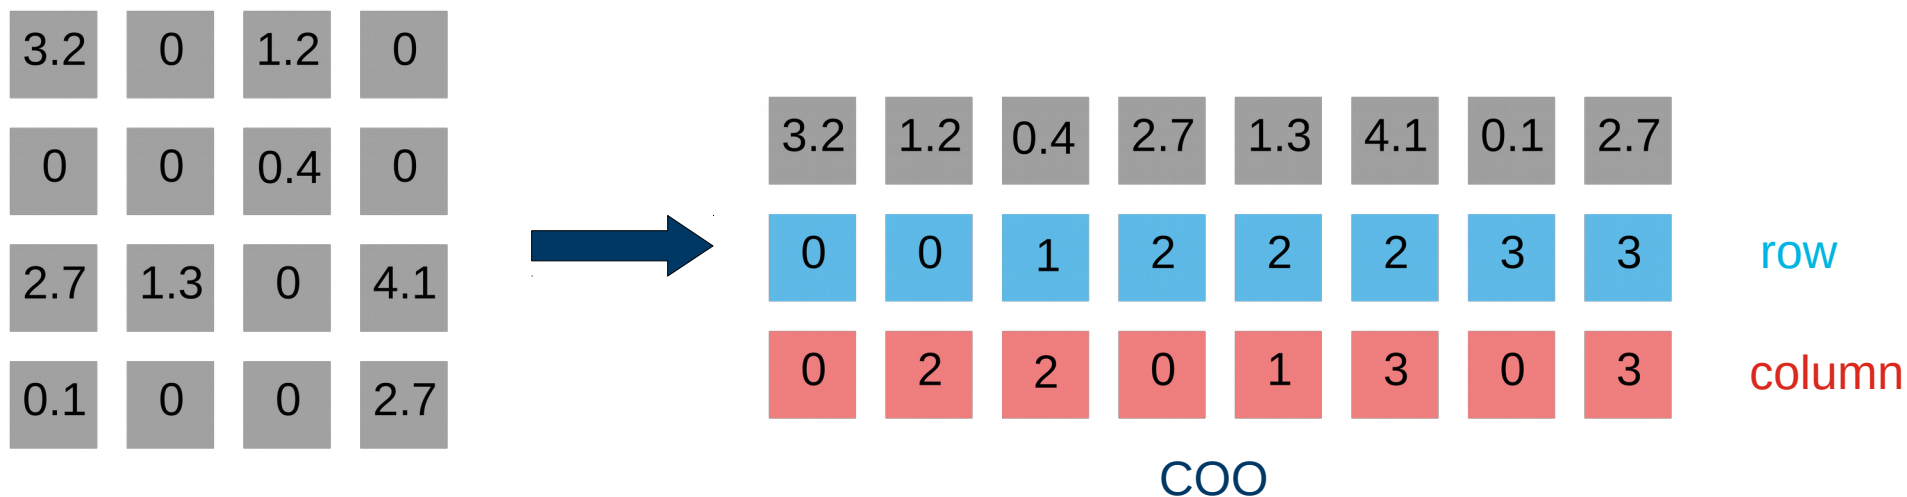
Sparse matrix formats

3.2	0	1.2	0
0	0	0.4	0
2.7	1.3	0	4.1
0.1	0	0	2.7

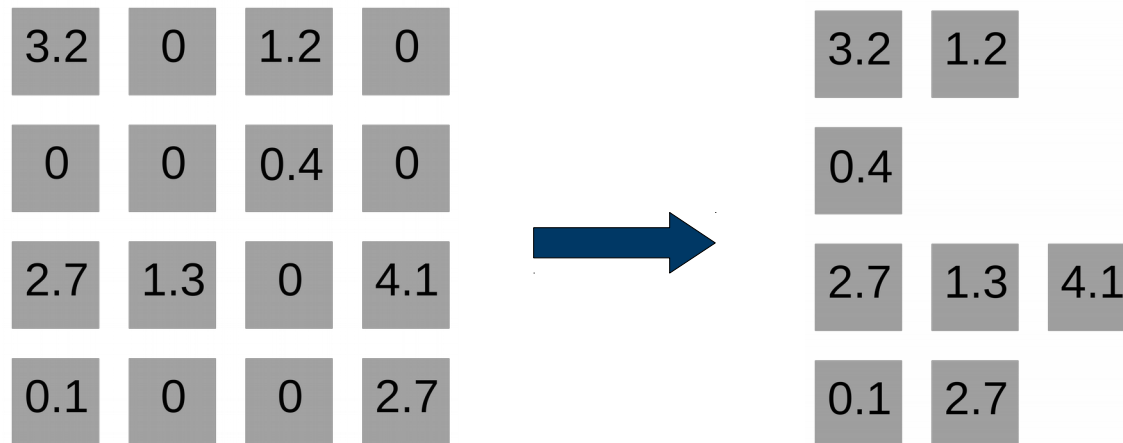
Sparse matrix formats



Sparse matrix formats



Sparse matrix formats




Sparse matrix formats

3.2	0	1.2	0
0	0	0.4	0
2.7	1.3	0	4.1
0.1	0	0	2.7



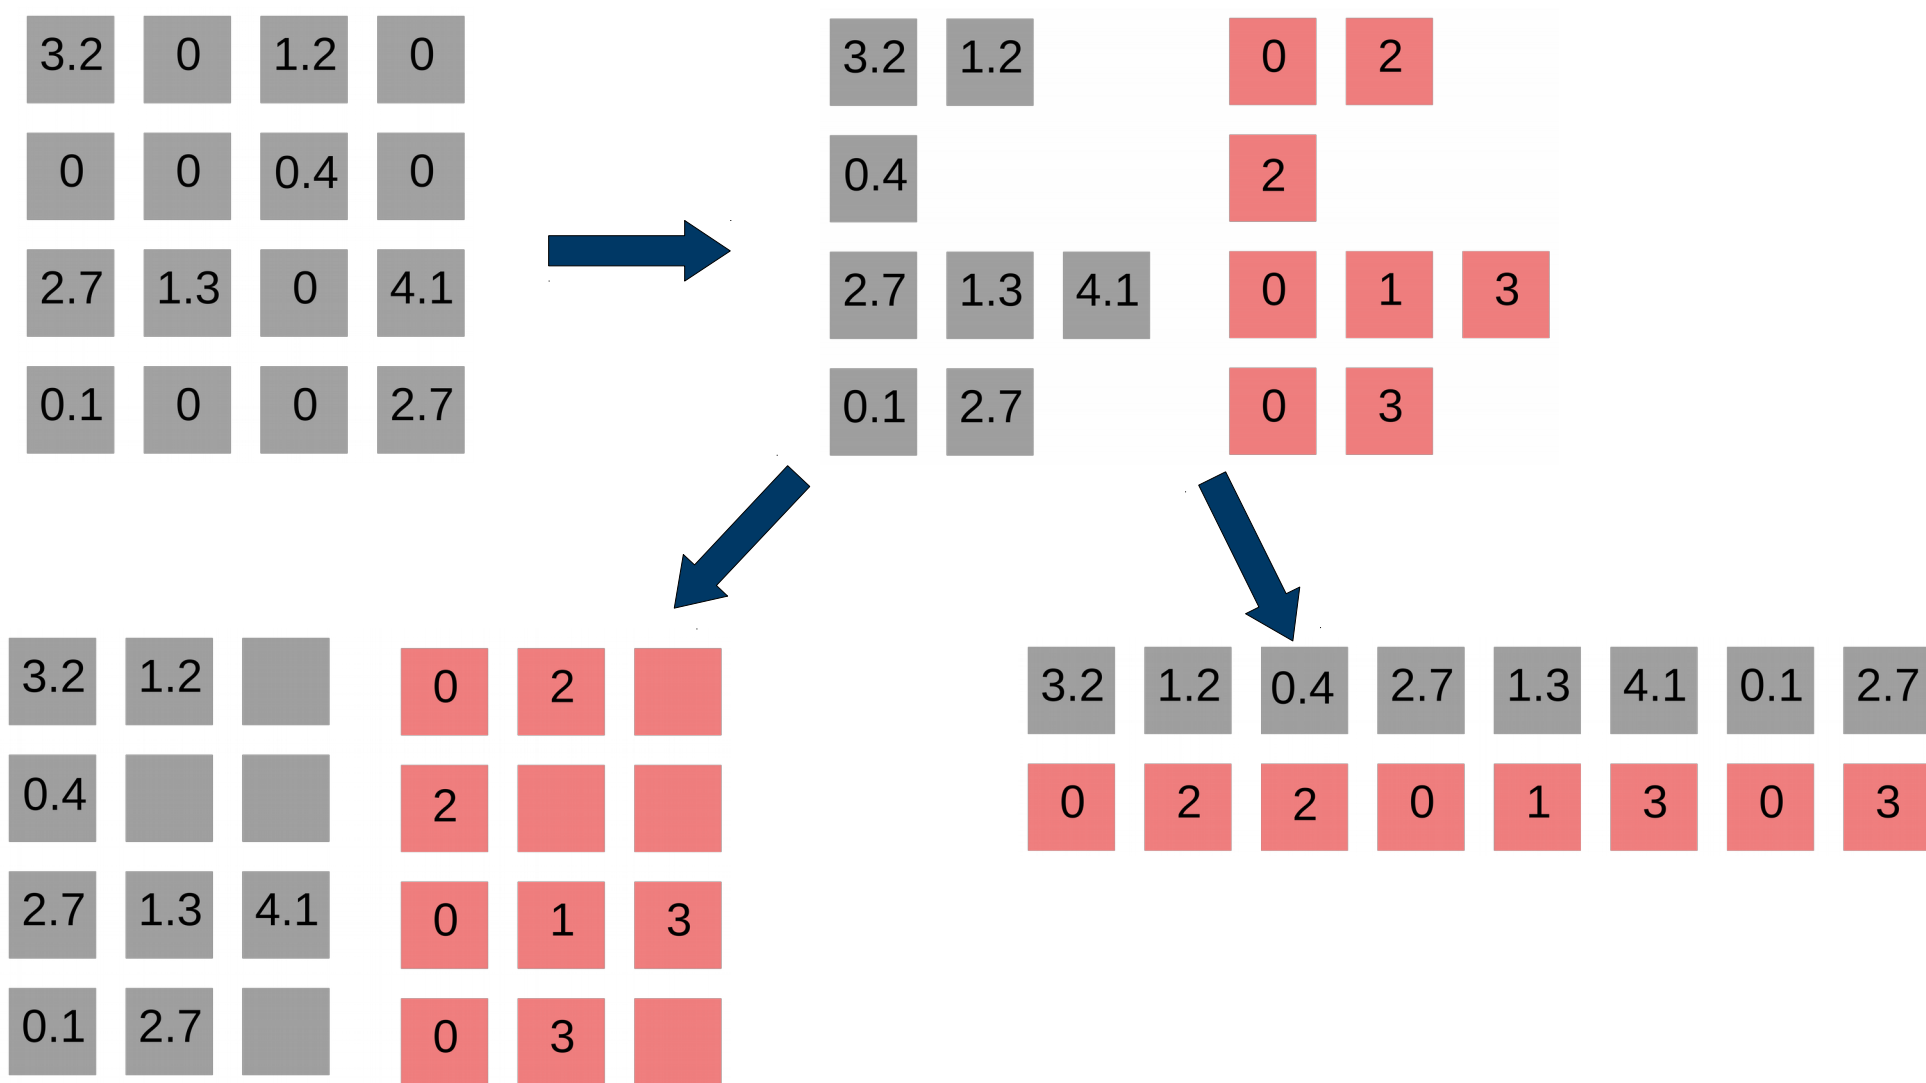
3.2	1.2			0	2	
0.4				2		
2.7	1.3	4.1		0	1	3
0.1	2.7			0	3	

Sparse matrix formats




ELL

Sparse matrix formats




ELL


Sparse matrix formats




Sparse matrix formats



Sparse matrix formats



... leads to CSC



CSR

“Standard” approach

CSR SpMV

3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7	Values (val)
0	2	2	0	1	3	0	3	Column indexes (colidx)
0	2	3	6					Row pointers (rowptr)

CSR SpMV

3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7	Values (val)
0	2	2	0	1	3	0	3	Column indexes (colidx)
0	2	3	6					Row pointers (rowptr)

$$y := Ax$$

CSR SpMV

3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7	Values (val)
0	2	2	0	1	3	0	3	Column indexes (colidx)
0	2	3	6					Row pointers (rowptr)

$$y := Ax$$

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

CSR SpMV

3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7
-----	-----	-----	-----	-----	-----	-----	-----

Values (val)

0	2	2	0	1	3	0	3
---	---	---	---	---	---	---	---

Column indexes (colidx)

0	2	3	6
---	---	---	---

Row pointers (rowptr)

$$y := Ax$$

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2   for (int i = 0; i < m; ++i) {  
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)  
4       y[i] += val[j] * x [ colidx[j] ];  
5   }  
6 }
```

Bell & Garland '08

- parallelize outer loop

~ cuSPARSE SpMV

CSR SpMV

3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7
-----	-----	-----	-----	-----	-----	-----	-----

Values (val)

0	2	2	0	1	3	0	3
---	---	---	---	---	---	---	---

Column indexes (colidx)

0	2	3	6
---	---	---	---

Row pointers (rowptr)

$$y := Ax$$

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2   for (int i = 0; i < m; ++i) {  
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)  
4       y[i] += val[j] * x [ colidx[j] ];  
5   }  
6 }
```

Bell & Garland '08

- parallelize outer loop

~ cuSPARSE SpMV

Load imbalance!
Non-coalescence!

CSR SpMV

3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7
-----	-----	-----	-----	-----	-----	-----	-----

Values (val)

0	2	2	0	1	3	0	3
---	---	---	---	---	---	---	---

Column indexes (colidx)

0	2	3	6
---	---	---	---

Row pointers (rowptr)

$$y := Ax$$

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2   for (int i = 0; i < m; ++i) {  
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)  
4       y[i] += val[j] * x [ colidx[j] ];  
5   }  
6 }
```

Bell & Garland '08

- parallelize outer loop

~ cuSPARSE SpMV

Load imbalance!
Non-coalescence!

Specialized formats

- HYB – ELL + COO
- SELL-P – good memory access, parallelizes well
- ... a few new ones every year

CSR SpMV

3.2	1.2	0.4	2.7	1.3	4.1	0.1	2.7	Values (val)
0	2	2	0	1	3	0	3	Column indexes (colidx)
0	2	3	6					Row pointers (rowptr)

$$y := Ax$$

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2   for (int i = 0; i < m; ++i) {  
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)  
4       y[i] += val[j] * x [ colidx[j] ];  
5   }  
6 }
```

Bell & Garland '08

- parallelize outer loop

~ cuSPARSE SpMV

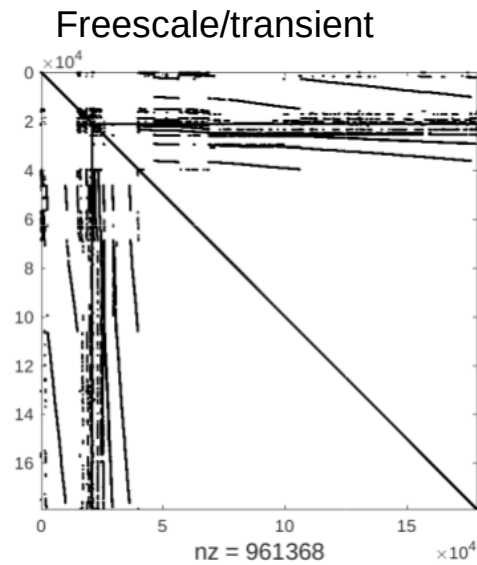
Load imbalance!
Non-coalescence!

Specialized formats

- HYB – ELL + COO
- SELL-P – good memory access, parallelizes well
- ... a few new ones every year

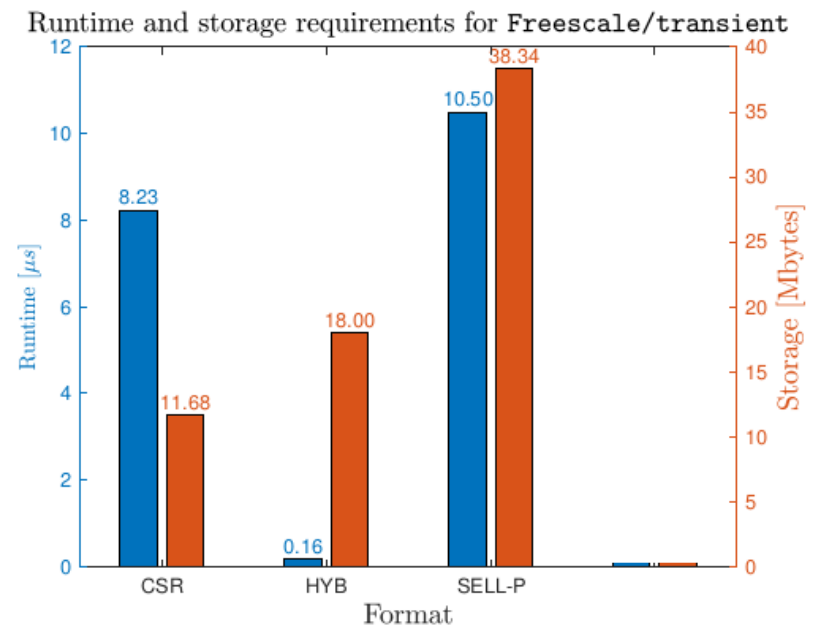
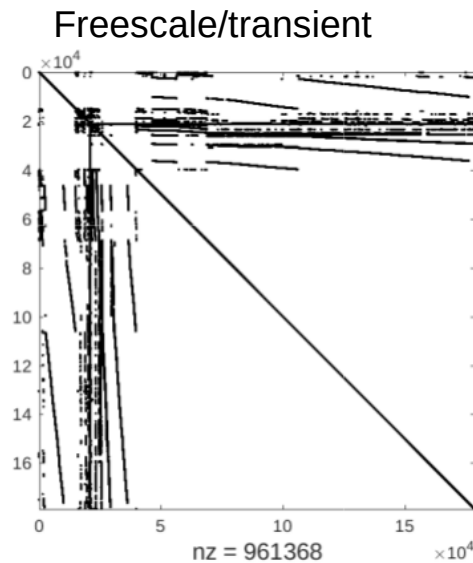
Library code – SpMV is not the only kernel!

Motivating example



Motivating example

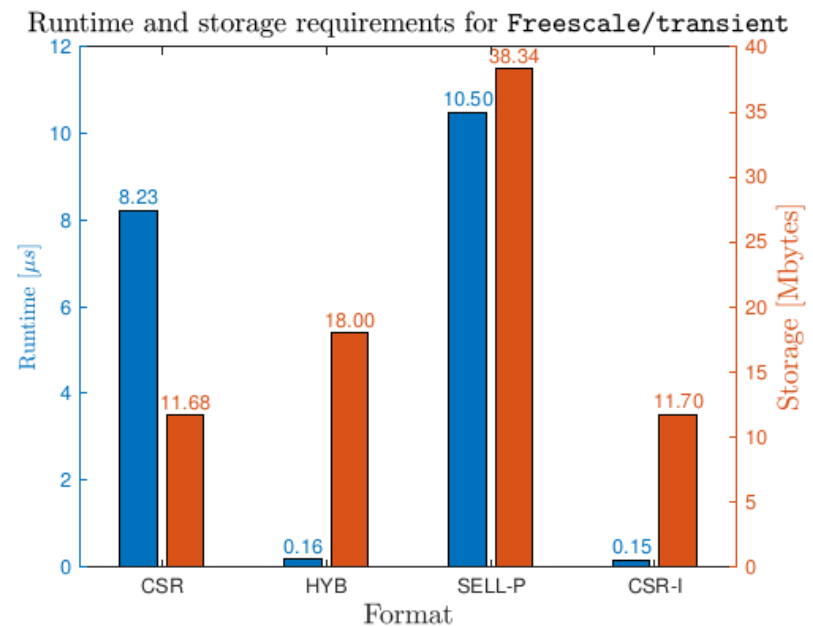
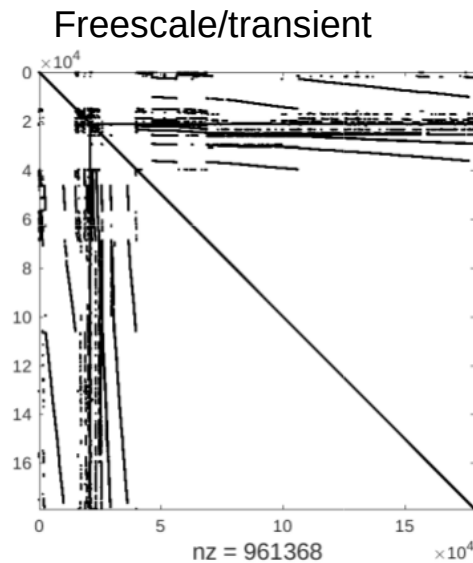
* GTX 1080



Can we do better than HYB using CSR?

Motivating example

* GTX 1080





Can we do better than HYB using CSR?

55x speedup

YES!

Motivating example

* GTX 1080



Can we do better than HYB using CSR?

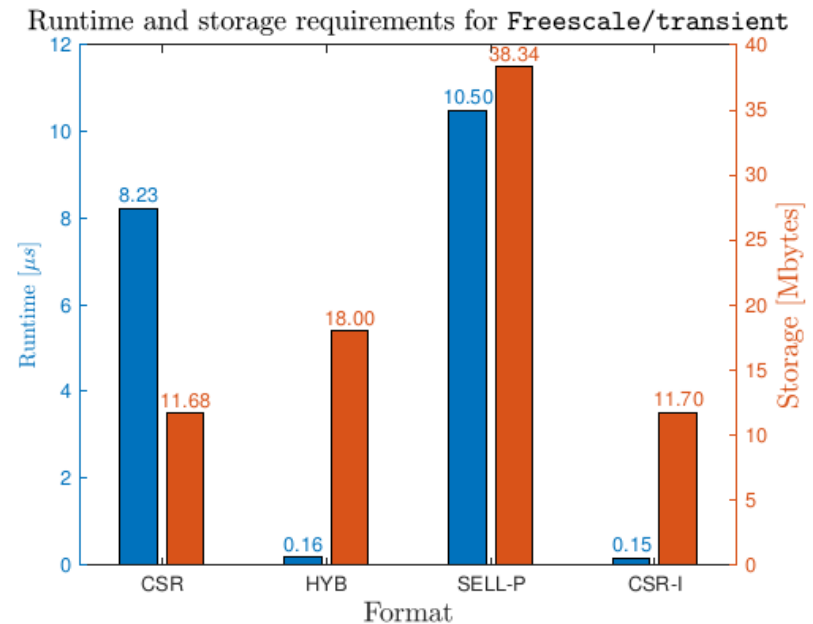
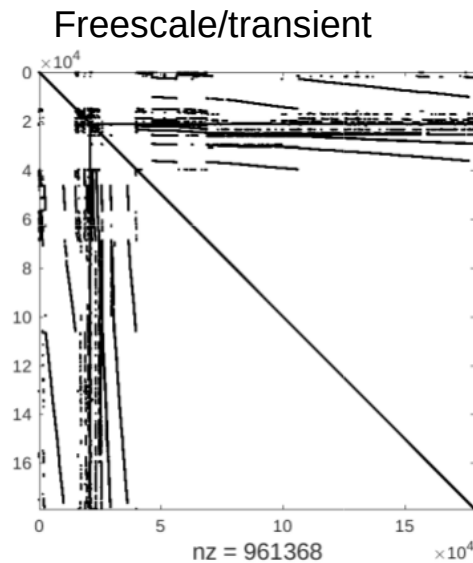
55x speedup

YES!

New **CSR-I** algorithm (and format).

Motivating example

* GTX 1080



Can we do better than HYB using CSR?

55x speedup

YES!

New **CSR-I** algorithm (and format).

Only need to add a small amount of data to vanilla CSR. Add - but not modify the existing format!

How to do it? By force!

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2     for (int i = 0; i < m; ++i) {  
3         for (int j = rowptr[i]; j < rowptr[i+1]; ++j)  
4             y[i] += val[j] * x [ colidx[j] ];  
5     }  
6 }
```

How to do it? By force!

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2   for (int i = 0; i < m; ++i) {  
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)  
4       y[i] += val[j] * x[ colidx[j] ];  
5   }  
6 }
```



Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {  
2   int row = -1, next_row = 0, nnz = rowptr[m];  
3   for (int i = 0; i < nnz; ++i) {  
4     while (i >= next_row) next_row = rowptr[++row+1];  
5     y[row] += val[i] * x[ colidx[i] ];  
6   }}
```

How to do it? By force!

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   int row = -1, next_row = 0, nnz = rowptr[m];
3   for (int i = 0; i < nnz; ++i) {
4     while (i >= next_row) next_row = rowptr[++row+1];
5     y[row] += val[i] * x[colidx[i]];
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
3   int row = -1, next_row = 0, nnz = rowptr[m];
4   for (int k = 0; k < T; ++k) {
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {
6       while (i >= next_row) next_row = rowptr[++row+1];
7       y[row] += val[i] * x[colidx[i]];
8     }}
```

How to do it? By force!

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   int row = -1, next_row = 0, nnz = rowptr[m];
3   for (int i = 0; i < nnz; ++i) {
4     while (i >= next_row) next_row = rowptr[++row+1];
5     y[row] += val[i] * x[colidx[i]];
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
3   int row = -1, next_row = 0, nnz = rowptr[m];
4   for (int k = 0; k < T; ++k) {
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {
6       while (i >= next_row) next_row = rowptr[++row+1];
7       y[row] += val[i] * x[colidx[i]];
8     }}
```


How to do it? By force!

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   int row = -1, next_row = 0, nnz = rowptr[m];
3   for (int i = 0; i < nnz; ++i) {
4     while (i >= next_row) next_row = rowptr[++row+1];
5     y[row] += val[i] * x[colidx[i]];
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
3   int row = -1, next_row = 0, nnz = rowptr[m];
4   for (int k = 0; k < T; ++k) {
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {
6       while (i >= next_row) next_row = rowptr[++row+1];
7       y[row] += val[i] * x[colidx[i]];
8     }}
```

Race conditions!

How to do it? By force!

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   int row = -1, next_row = 0, nnz = rowptr[m];
3   for (int i = 0; i < nnz; ++i) {
4     while (i >= next_row) next_row = rowptr[++row+1];
5     y[row] += val[i] * x[colidx[i]];
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
3   int row = -1, next_row = 0, nnz = rowptr[m];
4   for (int k = 0; k < T; ++k) {
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {
6       while (i >= next_row) next_row = rowptr[++row+1];
7       y[row] += val[i] * x[colidx[i]];
8     }}
```

Race conditions!

- Use atomics
- Accumulate partial result into registers

How to do it? By force!

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   int row = -1, next_row = 0, nnz = rowptr[m];
3   for (int i = 0; i < nnz; ++i) {
4     while (i >= next_row) next_row = rowptr[++row+1];
5     y[row] += val[i] * x[colidx[i]];
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
3   int row = -1, next_row = 0, nnz = rowptr[m];
4   for (int k = 0; k < T; ++k) {
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {
6       while (i >= next_row) next_row = rowptr[++row+1];
7       y[row] += val[i] * x[colidx[i]];
8     }}
```

Parallelize this!

State between outer loop iterations!

Race conditions!

- Use atomics
- Accumulate partial result into registers

How to do it? By force!

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   for (int i = 0; i < m; ++i) {
3     for (int j = rowptr[i]; j < rowptr[i+1]; ++j)
4       y[i] += val[j] * x[colidx[j]];
5   }
6 }
```

Merge the two loops into one.

```
1 void SpMV_CSR(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
2   int row = -1, next_row = 0, nnz = rowptr[m];
3   for (int i = 0; i < nnz; ++i) {
4     while (i >= next_row) next_row = rowptr[++row+1];
5     y[row] += val[i] * x[colidx[i]];
6   }}
```

Split the loop into equal chunks.

```
1 const int T = thread_count;
2 void SpMV_CSRI(int m, int *rowptr, int *colidx, float *val, float *x, float *y) {
3   int row = -1, next_row = 0, nnz = rowptr[m];
4   for (int k = 0; k < T; ++k) {
5     for (int i = k*nnz / T; i < (k+1)*nnz / T; ++i) {
6       while (i >= next_row) next_row = rowptr[++row+1];
7       y[row] += val[i] * x[colidx[i]];
8     }}
```

Parallelize this!

State between outer loop iterations!

Race conditions!

- Use atomics
- Accumulate partial result into registers

Precompute starting value of "row" for each thread.

Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Spreading out threads causes strided memory access.


Assign one warp per chunk.

Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Spreading out threads causes strided memory access.

Assign one warp per chunk.




Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Spreading out threads causes strided memory access.

Assign one warp per chunk.




Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Spreading out threads causes strided memory access.

Assign one warp per chunk.

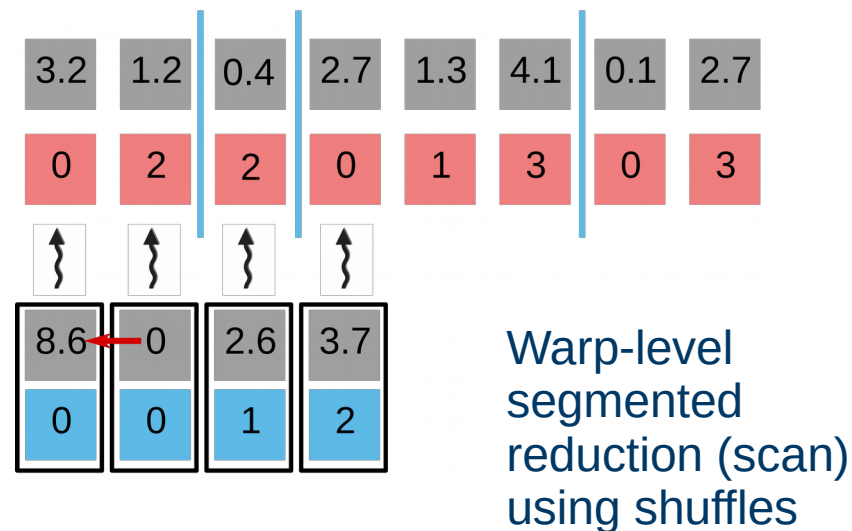



Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)

Spreading out threads causes strided memory access.

Assign one warp per chunk.

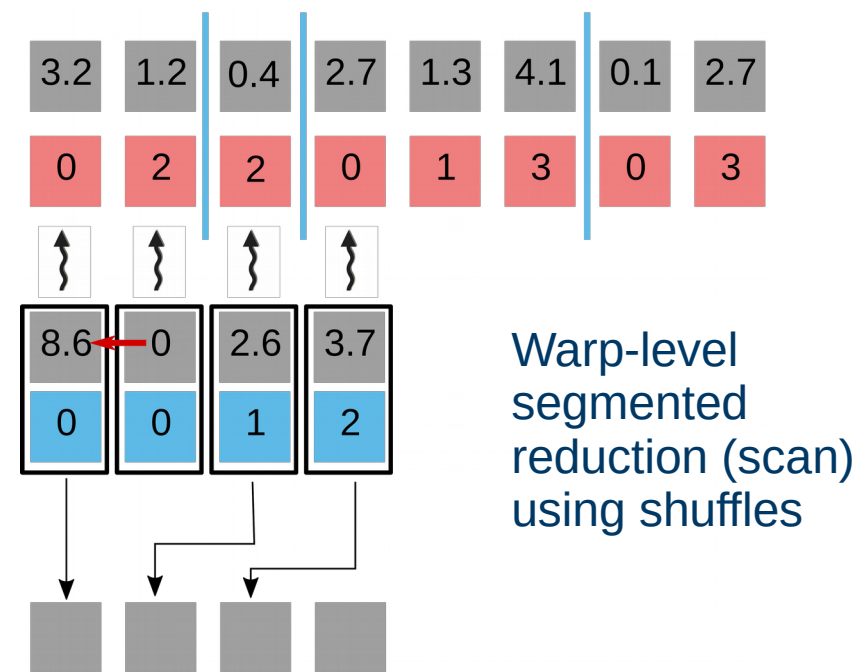



Getting good performance on GPUs

CUDA thread = 1 lane of a 32-wide SIMD unit (warp)


Spreading out threads causes strided memory access.

Assign one warp per chunk.

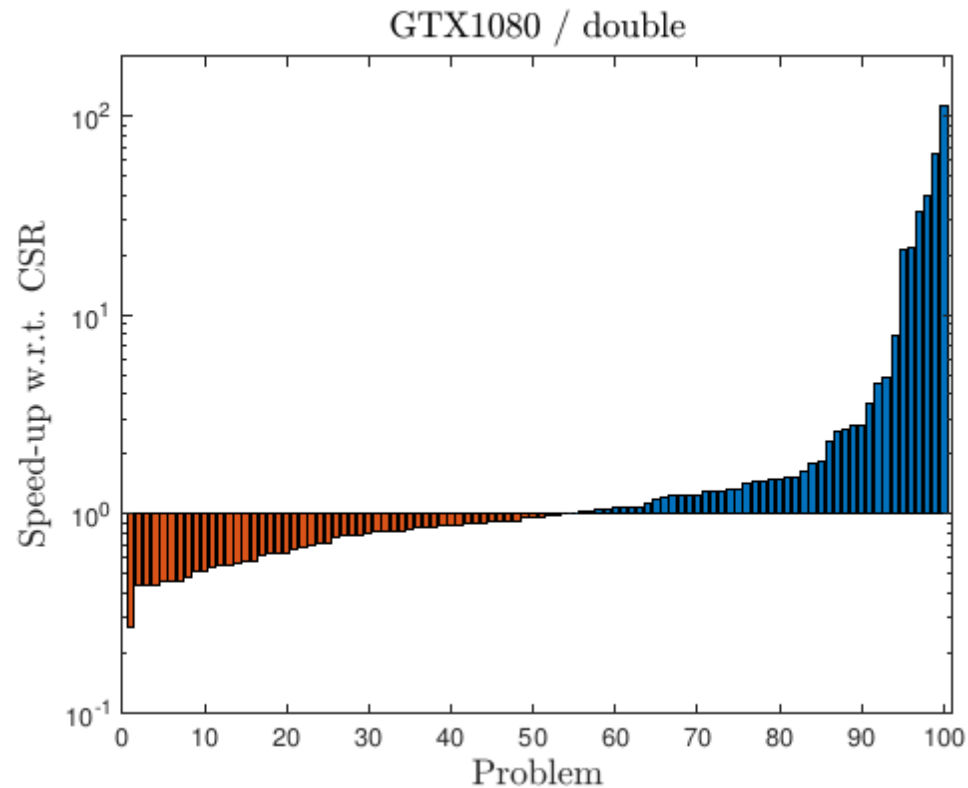


Performance of CSR-I


100 matrices from SuiteSparse



Speed-up / slowdown over cuSPARSE CSR



Speed-up / slowdown over cuSPARSE CSR



No format conversion!

- try both, and use the fastest later on!
- sometimes 1 cuSPARSE SpMV = 100 CSR-I SpMVs

Choosing the winner a priori

CSR-I designed for irregular patterns

Choosing the winner a priori

CSR-I designed for irregular patterns

How to measure irregularity?

Deviation of row lengths from the mean.

Choosing the winner a priori

CSR-I designed for irregular patterns

How to measure irregularity?

Deviation of row lengths from the mean.

Is “5” regular or irregular?

Depends on the density of the matrix (mean #rows)

Choosing the winner a priori

CSR-I designed for irregular patterns


How to measure irregularity?

Deviation of row lengths from the mean.

Is “5” regular or irregular?

Depends on the density of the matrix (mean #rows)

Scatter plot of speedup vs normalized std. dev.



Choosing the winner a priori

CSR-I designed for irregular patterns


How to measure irregularity?

Deviation of row lengths from the mean.


Is “5” regular or irregular?

Depends on the density of the matrix (mean #rows)

Scatter plot of speedup vs normalized std. dev.



Combining both approaches



Conclusion

Use **atomics and warp shuffles** to tackle **irregular sparsity patterns**.

Determine a priori when CSR-I is faster than standard algorithm.

- keep cuSPARSE performance for regular patterns
- CSR-I for irregular ones

Thank you! Questions?

All functionalities are part of the MAGMA-sparse project.

MAGMA SPARSE

ROUTINES BiCG, BiCGSTAB, Block-Asynchronous Jacobi, CG, CGS, GMRES, IDR, Iterative refinement, LOBPCG, LSQR, QMR, TFQMR

PRECONDITIONERS ILU / IC, Jacobi, ParILU, ParILUT, Block Jacobi, ISAI

KERNELS SpMV, SpMM

DATA FORMATS CSR, ELL, SELL-P, CSR5, HYB

<http://icl.cs.utk.edu/magma/>

Scan me
for slides!



github.com/gflegar/talks/europar_2017

This research is based on a cooperation between Hartwig Anzt, Jack Dongarra (University of Tennessee), Goran Flegar and Enrique S. Quintana-Ortí (Universitat Jaume I).



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE



UNIVERSITAT
JAUME I