

Batched Generation of Block-Jacobi Preconditioners for Iterative Sparse Linear System Solvers on GPUs

Hartwig Anzt, Jack Dongarra, Goran Flegar, Enrique S. Quintana-Ortí,
Andrés E. Tomás



Problem setting

- Solve sparse linear system using an iterative Krylov method

$$Ax = b, \quad A \in \mathbb{R}^{n \times n}$$

Problem setting

- Solve sparse linear system using an iterative Krylov method
- Convergence typically benefits from using a preconditioner

$$Ax = b, \quad A \in \mathbb{R}^{n \times n}$$

$$M^{-1}Ax = M^{-1}b$$

Problem setting

- Solve sparse linear system using an iterative Krylov method
- Convergence typically benefits from using a preconditioner
- **Need high degree of parallelism** to use a GPU effectively
 - 56 SMs x 64 cores = 3584 cores!
 - Oversubscribe to hide memory latency
- Use a preconditioner with high *parallelization potential*

$$Ax = b, \quad A \in \mathbb{R}^{n \times n}$$

$$M^{-1}Ax = M^{-1}b$$

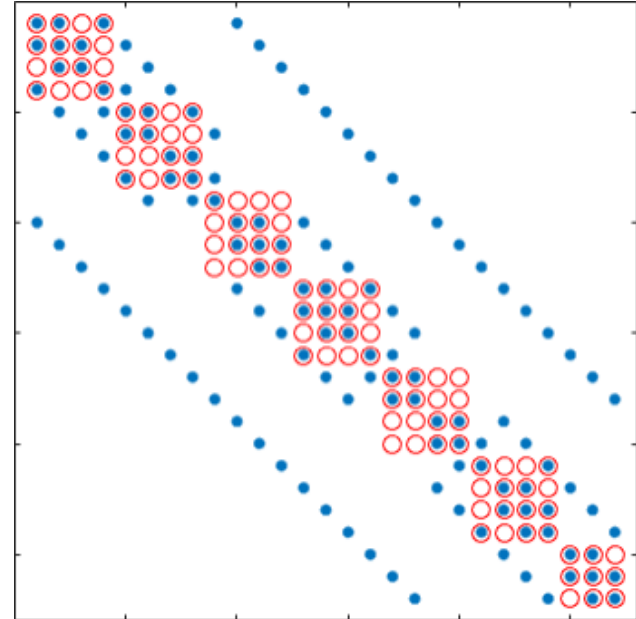
NVIDIA GP100



source: devblogs.nvidia.com/parallelforall/

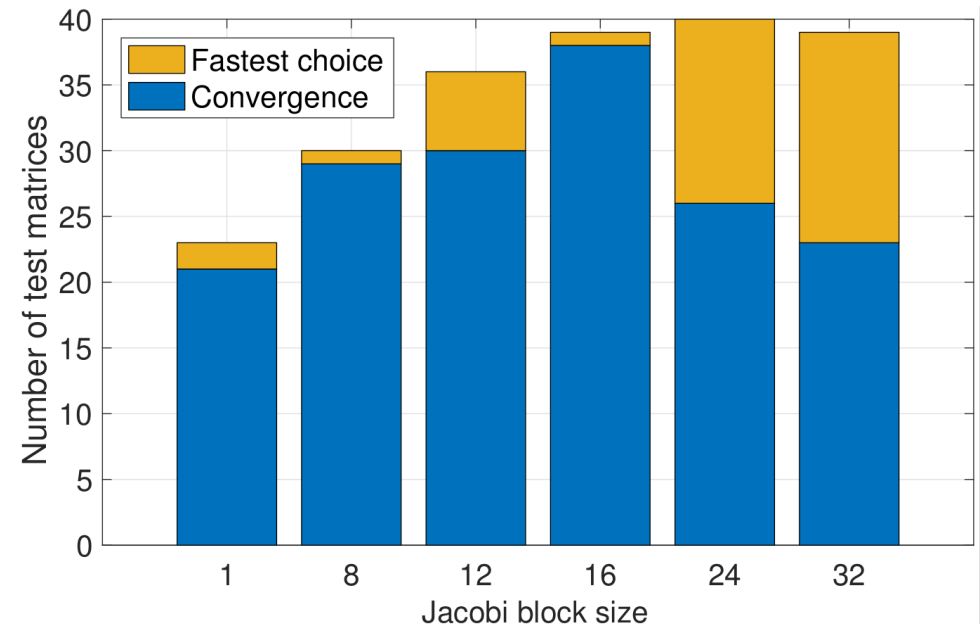
Block-Jacobi preconditioning

- Scalar Jacobi
 - Scale with inverse of main diagonal
- Block-Jacobi
 - Scale with inverses of diagonal blocks (possibly of different sizes!)
 - Can reflect the block structure of the problem
 - Often superior to scalar Jacobi
- Can process each block independently!



Benefits of block-Jacobi

- 40 matrices from SuiteSparse
- MAGMA-sparse open source library
 - IDR solver
 - Jacobi preconditioner
 - Supervariable blocking
- Block-Jacobi improves the robustness of the solver
 - More problems converge
- Decreases time-to-solution



- Restrict block size to 32×32
 - Large block sizes require more memory to store the preconditioner matrix

General Ideas

- Restrict block size to 32x32
 - Large block sizes require more memory to store the preconditioner matrix
- Use a single warp to process the whole block (one thread per column)
 - No need for explicit synchronization
- Use the large register file to store the entire block
 - Read/write from mem. once
 - Comm. via warp shuffles
 - Avoids load/store instructions



General Ideas

- Restrict block size to 32x32
 - Large block sizes require more memory to store the preconditioner matrix
- Use a single warp to process the whole block (one thread per column)
 - No need for explicit synchronization




General Ideas

- Restrict block size to 32x32
 - Large block sizes require more memory to store the preconditioner matrix
- Use a single warp to process the whole block (one thread per column)
 - No need for explicit synchronization
- Use the large register file to store the entire block
 - Read/write from mem. once
 - Comm. via warp shuffles
 - Avoids load/store instructions




Implementation options

- Inversion in preconditioner setup + matrix-vector product in application
 - (FLOPS: $2n^3$ setup, $2n^2$ app.)
 - Batched Gauss-Jordan elimination (BGJE)
 - Each step consists of column scaling and a rank-1 update of the whole matrix
 - Easily achievable load balancing
 - H. Anzt et al., “Batched Gauss-Jordan Elimination for Block-Jacobi Preconditioner Generation on GPUs”, PMAM’17

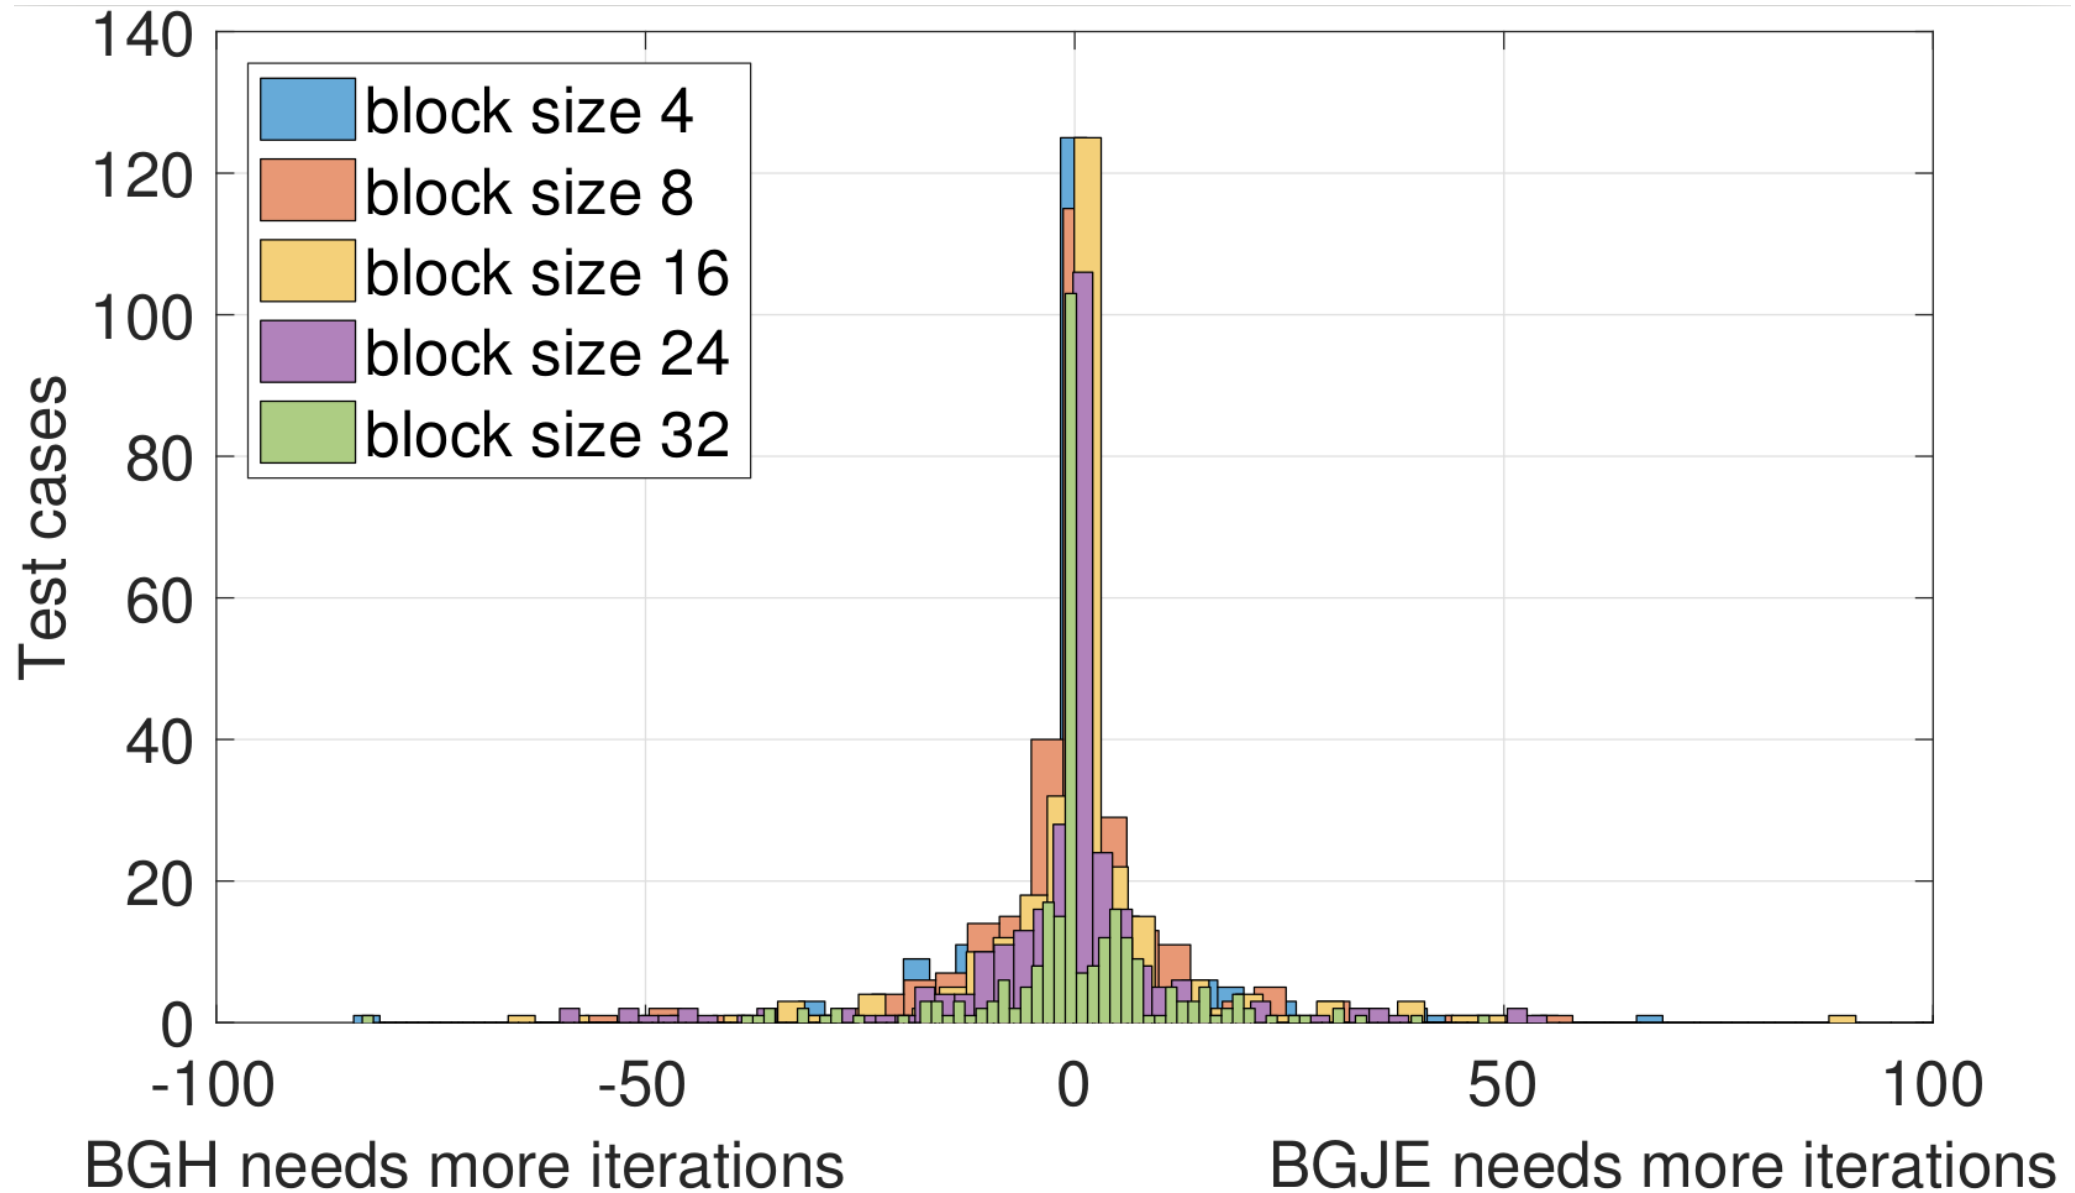


Implementation options

- Matrix decomposition in setup + solve in application
 - (FLOPS: $\frac{2}{3}n^3$ setup, $2n^2$ app.)
 - Gauss-Huard decomposition




Inversion?!





Gauss-Huard decomposition

- Decomposition
 - GEMV ($G = G - RR$)
 - SCAL ($O = O / B$)
 - GER ($L = L - BO$)
 - Column pivoting
 - Do not swap the columns, just remember which thread holds which column of the result



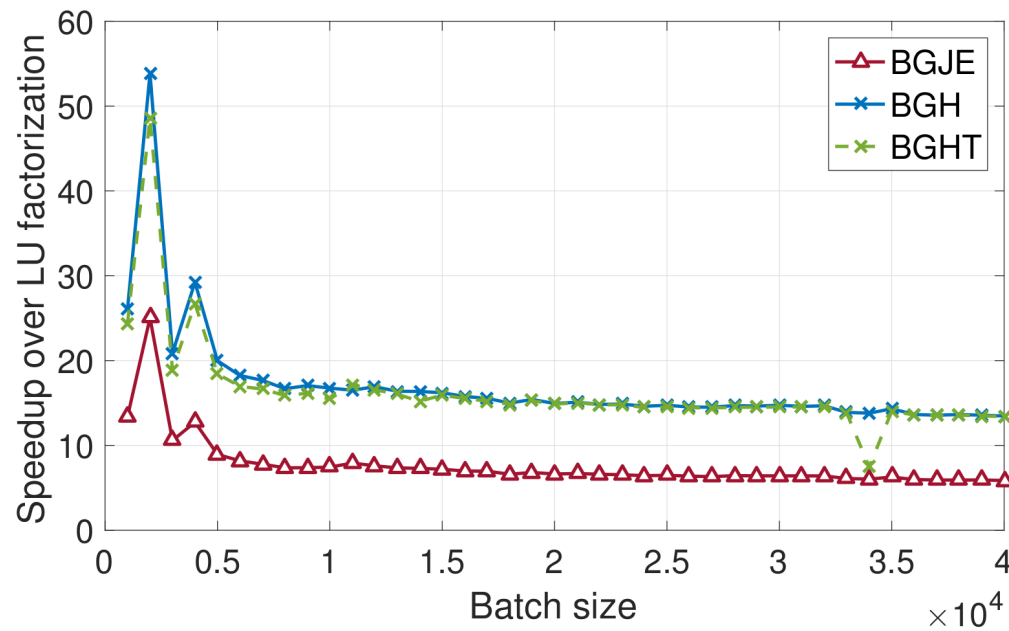
Gauss-Huard decomposition

- Decomposition
 - GEMV ($G = G - RR$)
 - SCAL ($O = O / B$)
 - GER ($L = L - BO$)
 - Column pivoting
 - Do not swap the columns, just remember which thread holds which column of the result
- Solve
 - Load only the solution vector into registers
 - DOT ($G = G - RR$)
 - SCAL ($O = O / B$)
 - AXPY ($L = L - BO$)
 - Write lower part transposed wrp. to anti-diagonal for coalesced mem. access (GHT)

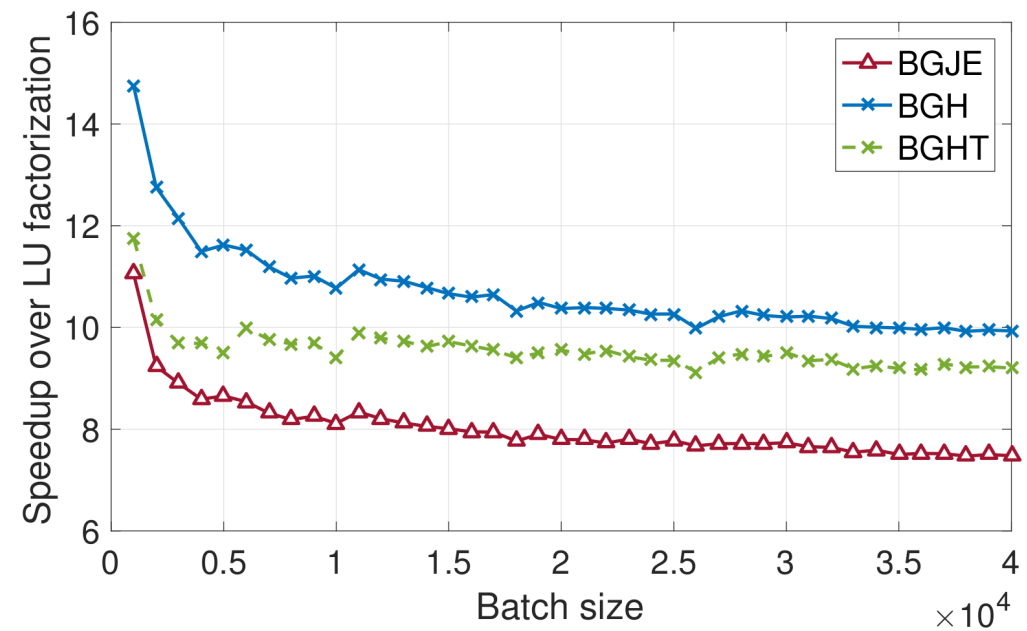


Decomposition comparison to batched LU (MAGMA)


Block size 16



Block size 32





Application time





Total runtime of block-Jacobi preconditioned BiCGSTAB

Block size 16



Block size 32



Thank you! Questions?

All functionalities are part of the MAGMA-sparse project.

MAGMA SPARSE

ROUTINES BiCG, BiCGSTAB, Block-Asynchronous Jacobi, CG, CGS, GMRES, IDR, Iterative refinement, LOBPCG, LSQR, QMR, TFQMR

PRECONDITIONERS ILU / IC, Jacobi, ParILU, ParILUT, Block Jacobi, ISAI

KERNELS SpMV, SpMM

DATA FORMATS CSR, ELL, SELL-P, CSR5, HYB

<http://icl.cs.utk.edu/magma/>



This research is based on a cooperation between Hartwig Anzt, Jack Dongarra (University of Tennessee), Goran Flegar, Enrique S. Quintana-Ortí and Adrés E. Tomás (Universidad Jaume I).

