

# Batched Generation of Block-Jacobi Preconditioners for Iterative Sparse Linear System Solvers on GPUs

---

Hartwig Anzt, Jack Dongarra, Goran Flegar, Enrique S. Quintana-Ortí,  
Andrés E. Tomás



# Problem setting

---

- Solve sparse linear system using an iterative Krylov method

$$Ax = b, \quad A \in \mathbb{R}^{n \times n}$$

# Problem setting

---

- Solve sparse linear system using an iterative Krylov method
- Convergence typically benefits from using a preconditioner

$$Ax = b, \quad A \in \mathbb{R}^{n \times n}$$

$$M^{-1}Ax = M^{-1}b$$

# Problem setting

- Solve sparse linear system using an iterative Krylov method
- Convergence typically benefits from using a preconditioner
- **Need high degree of parallelism** to use a GPU effectively
  - 56 SMs x 64 cores = 3584 cores!
  - Oversubscribe to hide memory latency
- Use a preconditioner with high *parallelization potential*

$$Ax = b, \quad A \in \mathbb{R}^{n \times n}$$

$$M^{-1}Ax = M^{-1}b$$

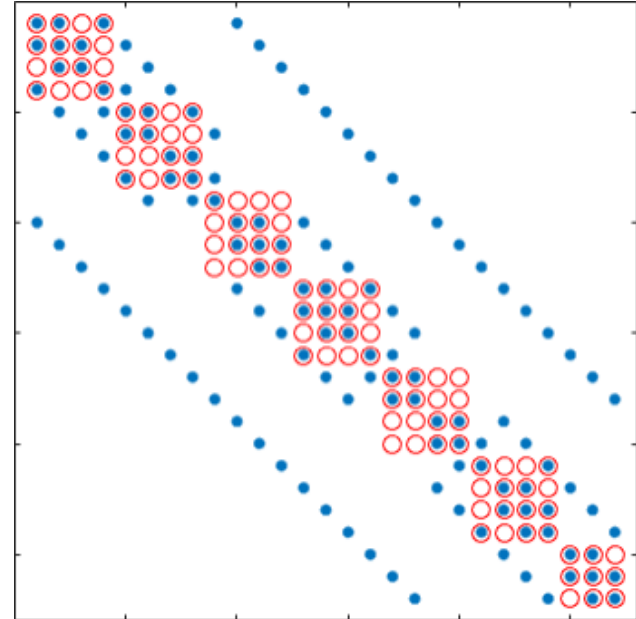
NVIDIA GP100



source: [devblogs.nvidia.com/parallelforall/](https://devblogs.nvidia.com/parallelforall/)

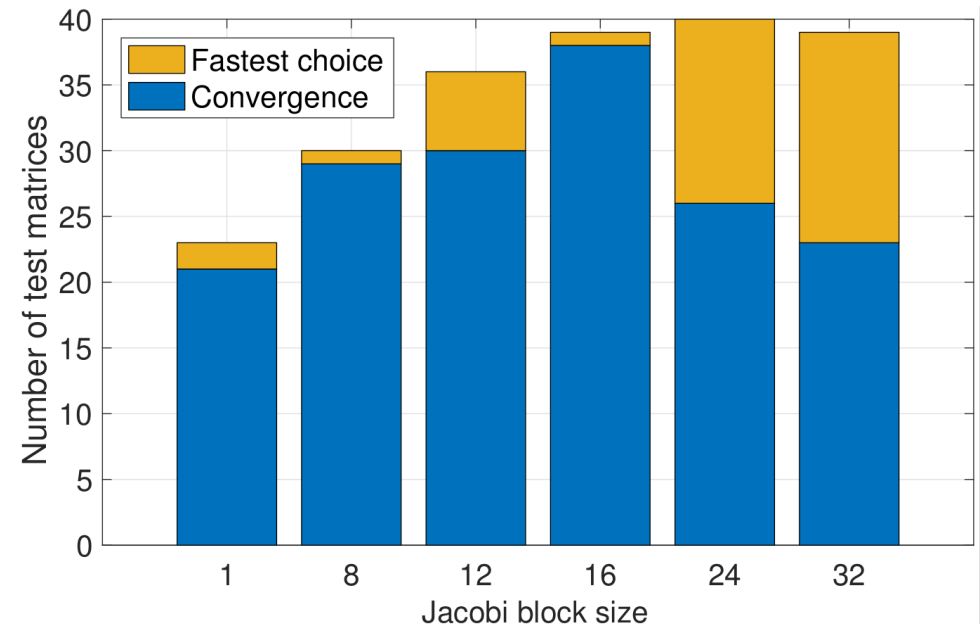
# Block-Jacobi preconditioning

- Scalar Jacobi
  - Scale with inverse of main diagonal
- Block-Jacobi
  - Scale with inverses of diagonal blocks (possibly of different sizes!)
  - Can reflect the block structure of the problem
  - Often superior to scalar Jacobi
- Can process each block independently!



# Benefits of block-Jacobi

- 40 matrices from SuiteSparse
- MAGMA-sparse open source library
  - IDR solver
  - Jacobi preconditioner
  - Supervariable blocking
- Block-Jacobi improves the robustness of the solver
  - More problems converge
- Decreases time-to-solution



- Restrict block size to  $32 \times 32$ 
  - Large block sizes require more memory to store the preconditioner matrix

# General Ideas

- Restrict block size to 32x32
  - Large block sizes require more memory to store the preconditioner matrix
- Use a single warp to process the whole block (one thread per column)
  - No need for explicit synchronization





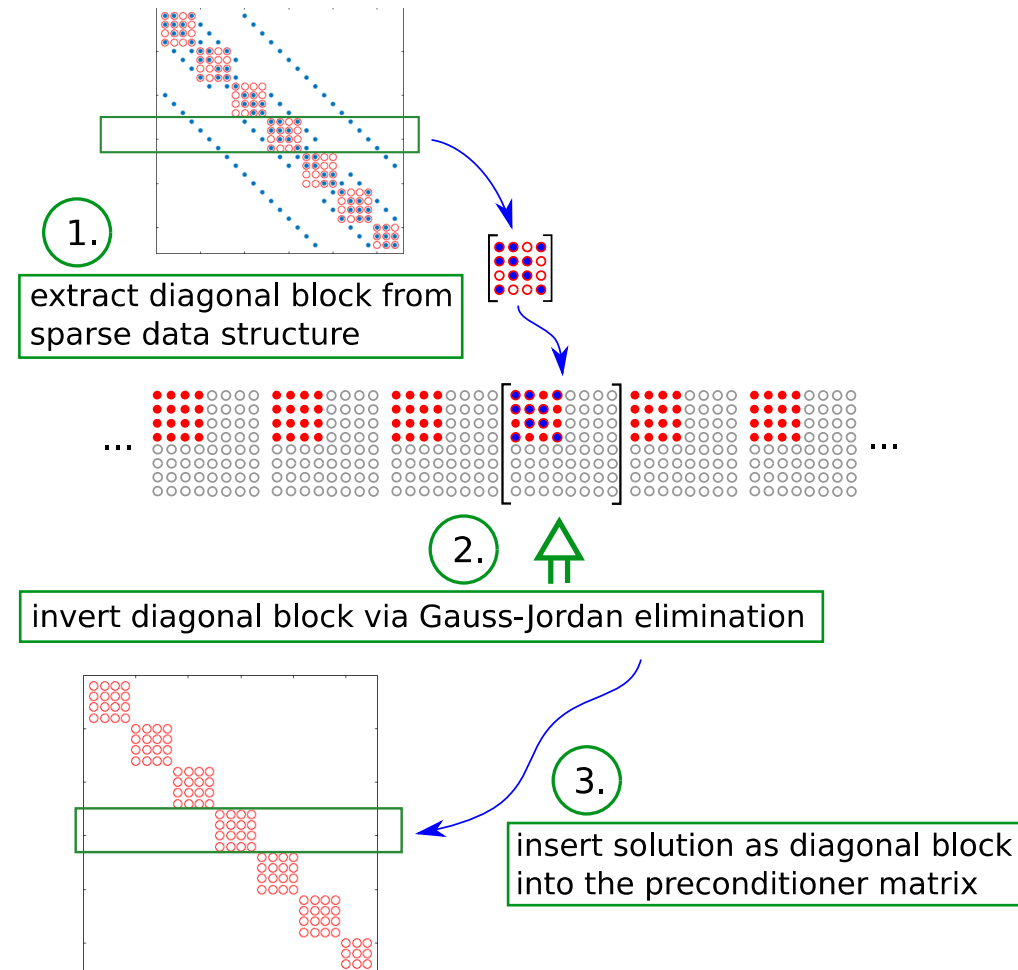
# General Ideas

- Restrict block size to 32x32
  - Large block sizes require more memory to store the preconditioner matrix
- Use a single warp to process the whole block (one thread per column)
  - No need for explicit synchronization
- Use the large register file to store the entire block
  - Read/write from mem. once
  - Comm. via warp shuffles
  - Avoids load/store instructions



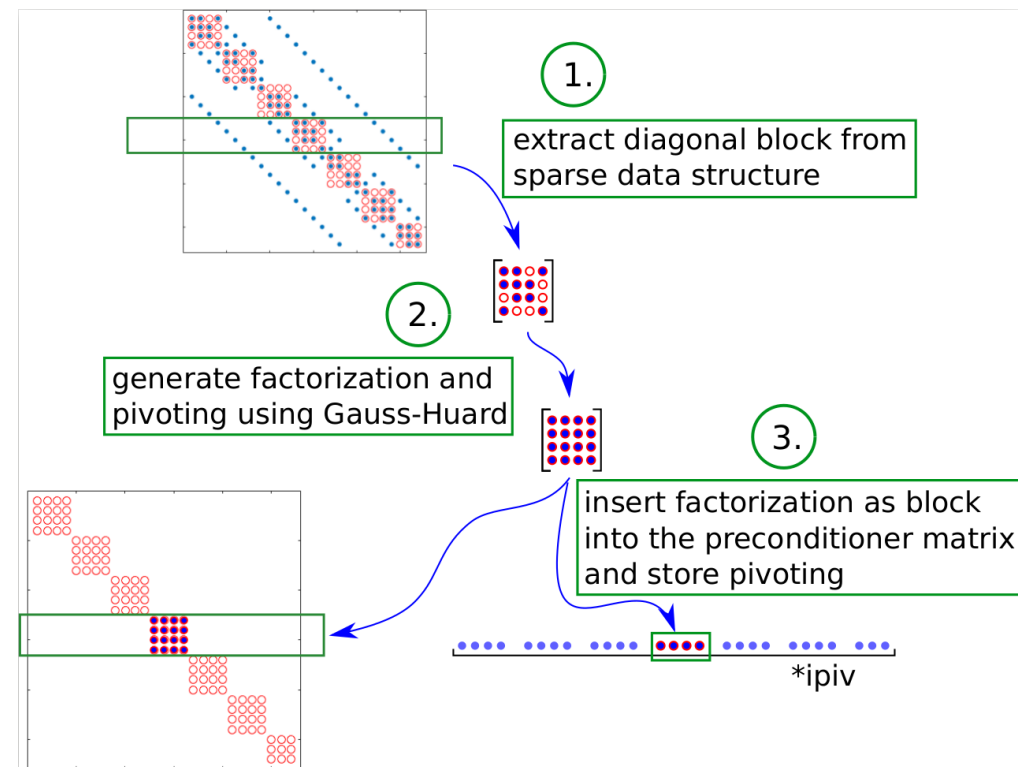
# Implementation options

- Inversion in preconditioner setup + matrix-vector product in application
  - (FLOPS:  $2n^3$  setup,  $2n^2$  app.)
  - Batched Gauss-Jordan elimination (BGJE)
    - Each step consists of column scaling and a rank-1 update of the whole matrix
    - Easily achievable load balancing
  - H. Anzt et al., “Batched Gauss-Jordan Elimination for Block-Jacobi Preconditioner Generation on GPUs”, PMAM’17

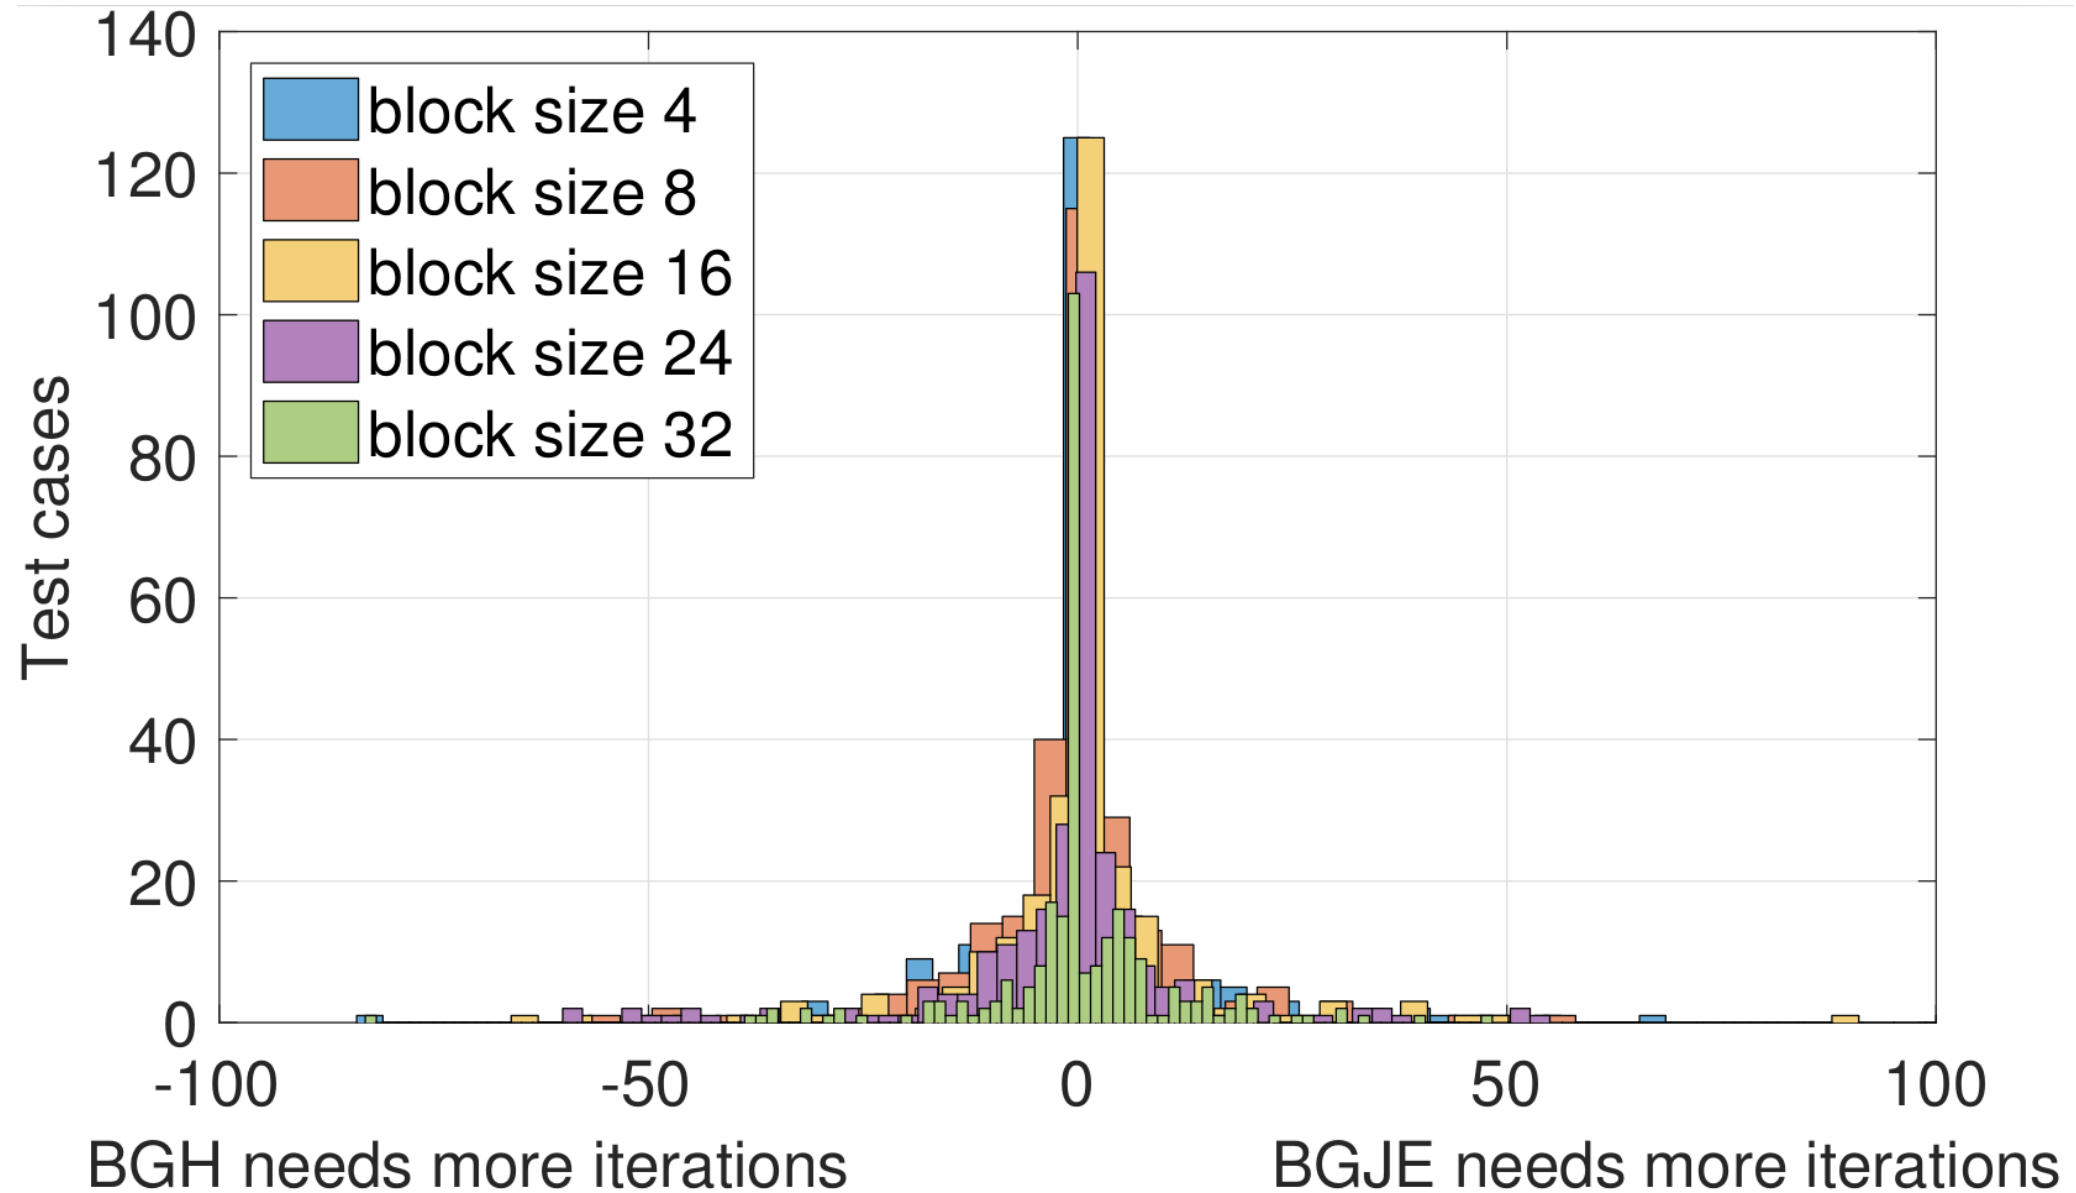


# Implementation options

- Matrix decomposition in setup + solve in application
  - (FLOPS:  $\frac{2}{3}n^3$  setup,  $2n^2$  app.)
  - Gauss-Huard decomposition

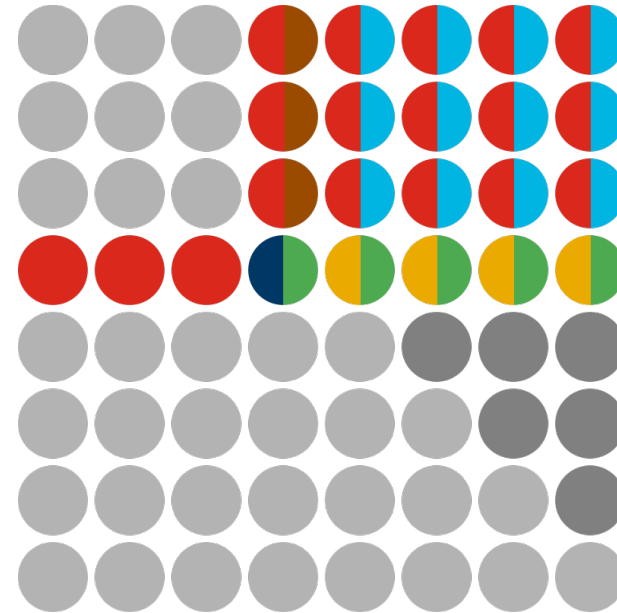


# Inversion?!



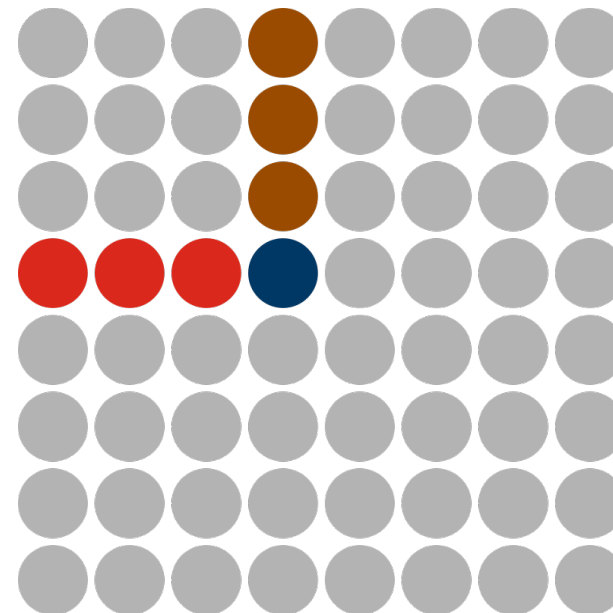
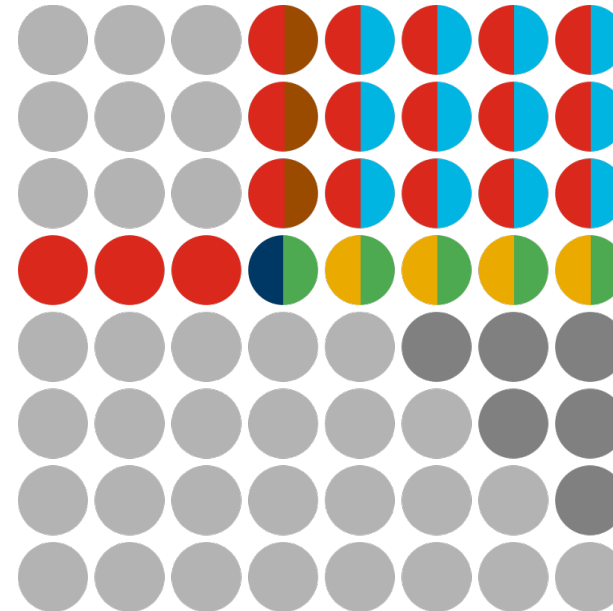
# Gauss-Huard decomposition

- Decomposition
  - GEMV ( $G = G - RR$ )
  - SCAL ( $O = O / B$ )
  - GER ( $L = L - BO$ )
  - Column pivoting
    - Do not swap the columns, just remember which thread holds which column of the result



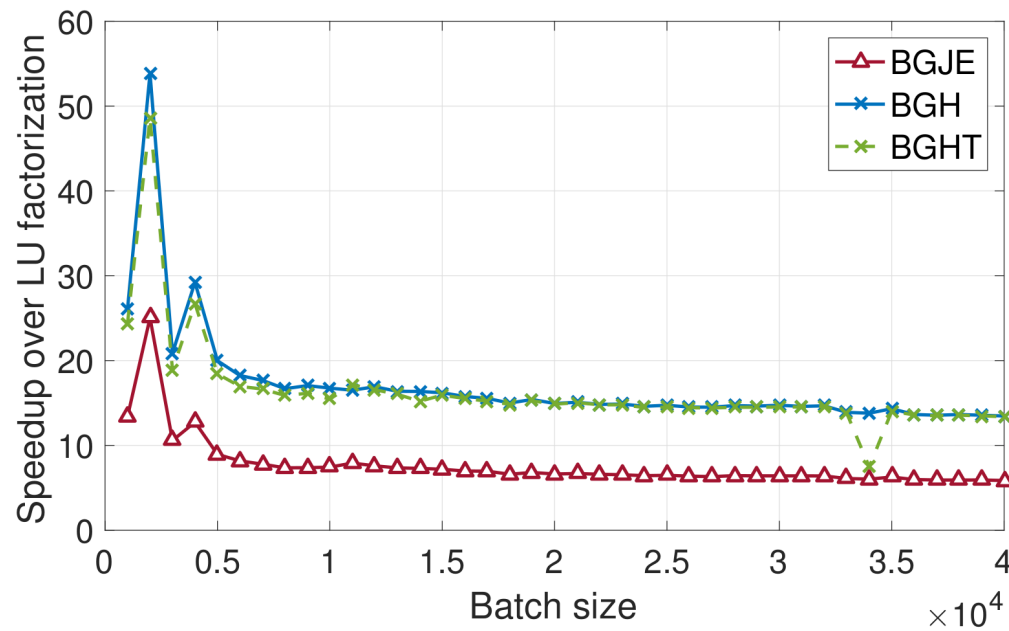
# Gauss-Huard decomposition

- Decomposition
  - GEMV ( $G = G - RR$ )
  - SCAL ( $O = O / B$ )
  - GER ( $L = L - BO$ )
  - Column pivoting
    - Do not swap the columns, just remember which thread holds which column of the result
- Solve
  - Load only the solution vector into registers
  - DOT ( $G = G - RR$ )
  - SCAL ( $O = O / B$ )
  - AXPY ( $L = L - BO$ )
  - Write lower part transposed wrp. to anti-diagonal for coalesced mem. access ( $GHT$ )

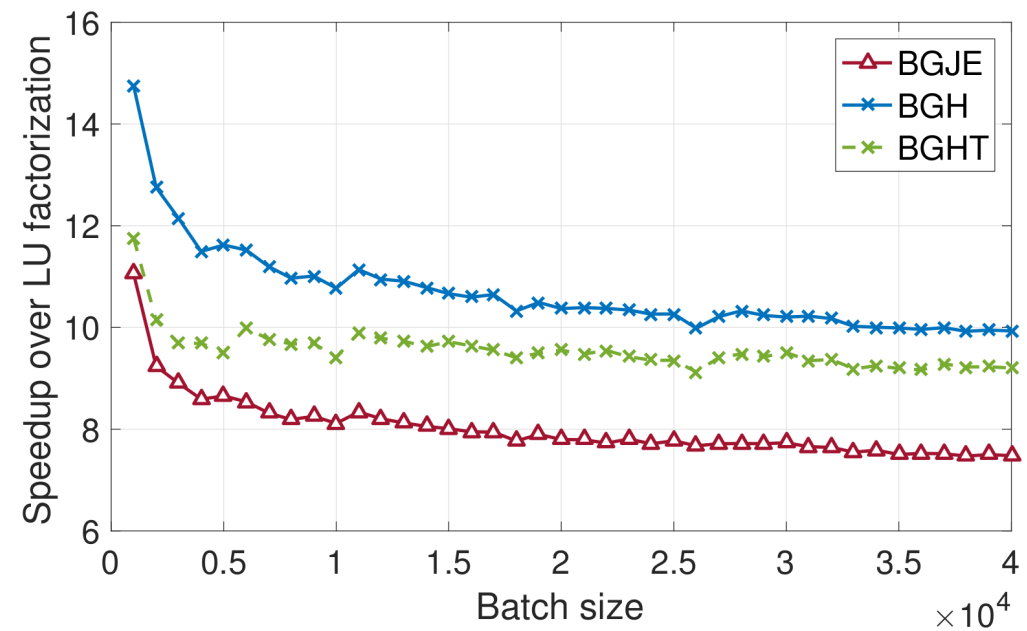


# Decomposition comparison to batched LU (MAGMA)

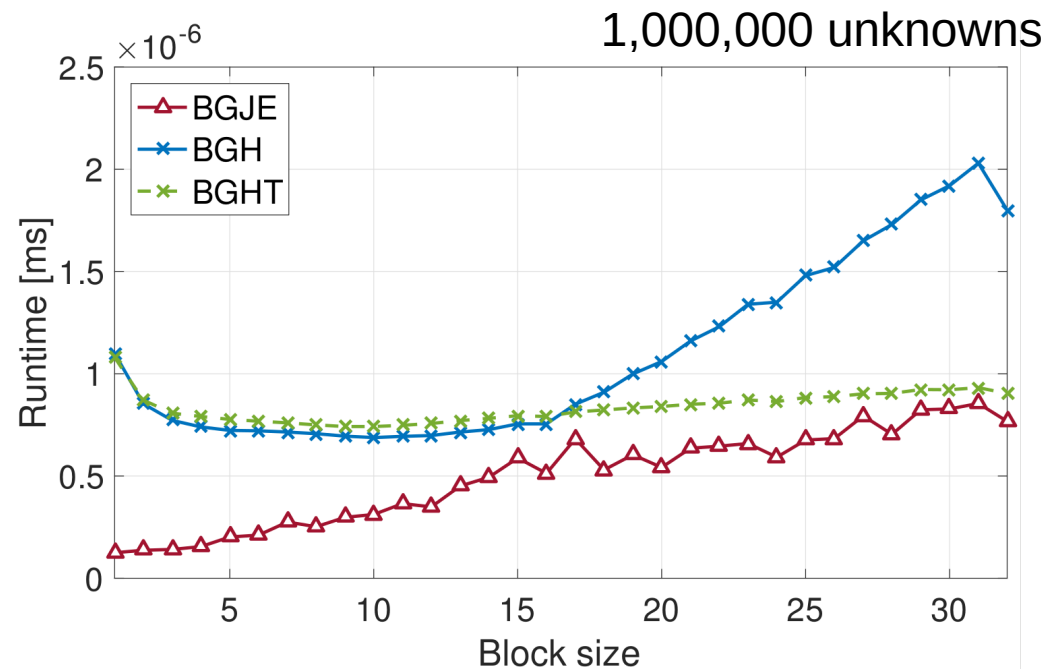
Block size 16



Block size 32



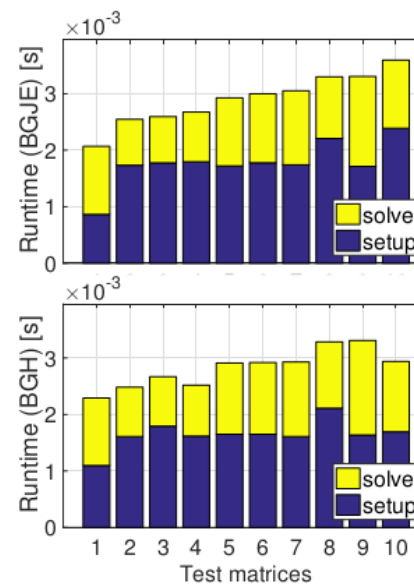
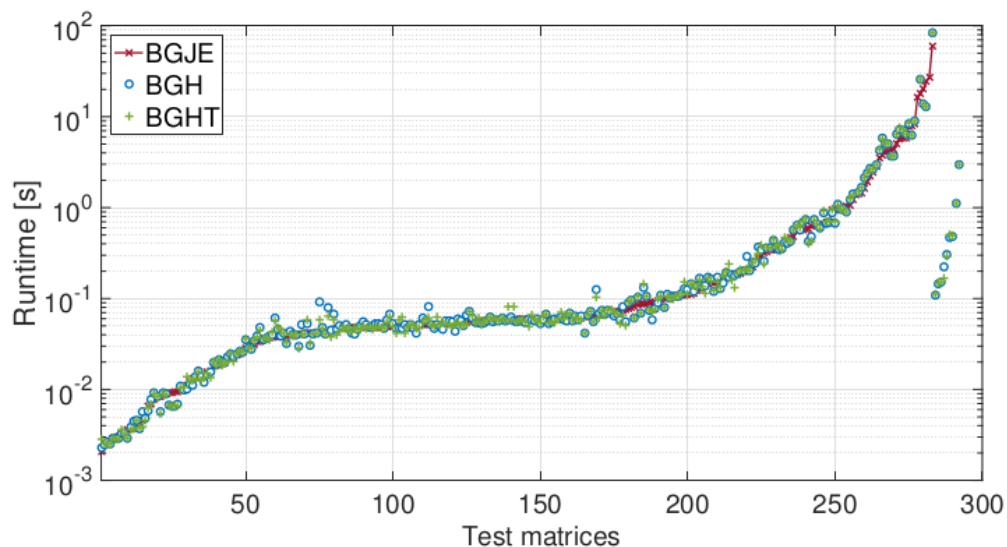
# Application time



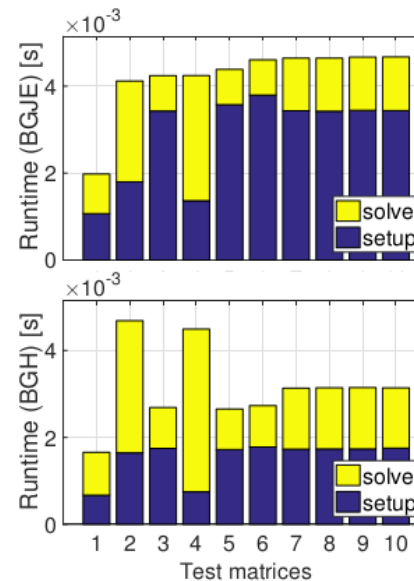
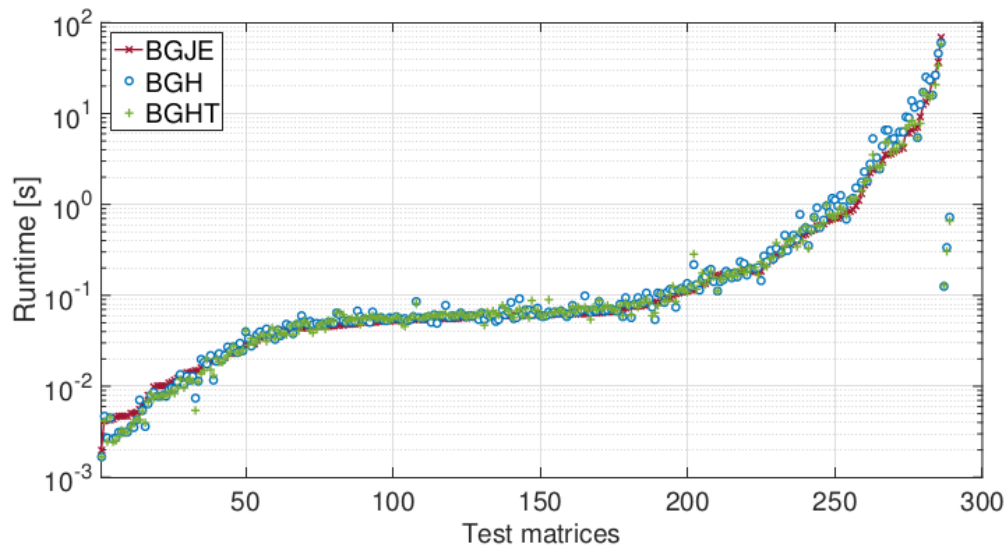


# Total runtime of block-Jacobi preconditioned BiCGSTAB

Block size 16



Block size 32



# Thank you! Questions?

All functionalities are part of the MAGMA-sparse project.

## MAGMA SPARSE

**ROUTINES** BiCG, BiCGSTAB, Block-Asynchronous Jacobi, CG, CGS, GMRES, IDR, Iterative refinement, LOBPCG, LSQR, QMR, TFQMR

**PRECONDITIONERS** ILU / IC, Jacobi, ParILU, ParILUT, Block Jacobi, ISAI

**KERNELS** SpMV, SpMM

**DATA FORMATS** CSR, ELL, SELL-P, CSR5, HYB

<http://icl.cs.utk.edu/magma/>



*This research is based on a cooperation between Hartwig Anzt, Jack Dongarra (University of Tennessee), Goran Flegar, Enrique S. Quintana-Ortí and Adrés E. Tomás (Universidad Jaume I).*

