

Block-Jacobi preconditioning in Ginkgo

Goran Flegar

Introduction

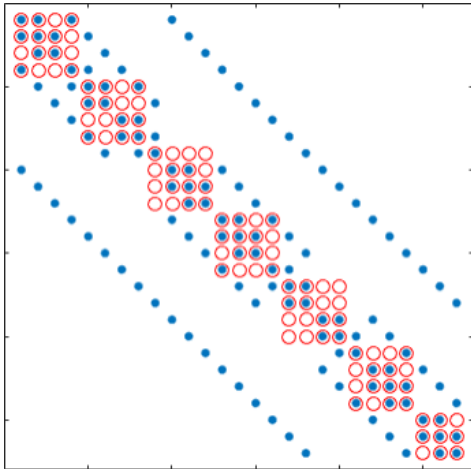
- I'm assuming we are all familiar with:
 - Iterative solution of linear systems
 - Why it's useful
 - Operations within Krylov solvers (BLAS 1, SpMV, preconditioning)
 - Preconditioning
 - Replacing the original system with the preconditioned system
 - Changes needed within the Krylov solver to incorporate preconditioning

Block-Jacobi preconditioning

- Improve performance for problems with inherent block structure
 - Usually blocks are smaller than 30×30 (blocks can be of different sizes!)

Block-Jacobi preconditioning

- Improve performance for problems with inherent block structure
 - Usually blocks are smaller than 30x30 (**blocks can be of different sizes!**)



- Block-Jacobi preconditioning
 - Use only diagonal blocks for approximation

$$\text{diag}(A) = [D_1, \dots, D_k]$$

$$M := \text{diag}(D_1, \dots, D_k)$$

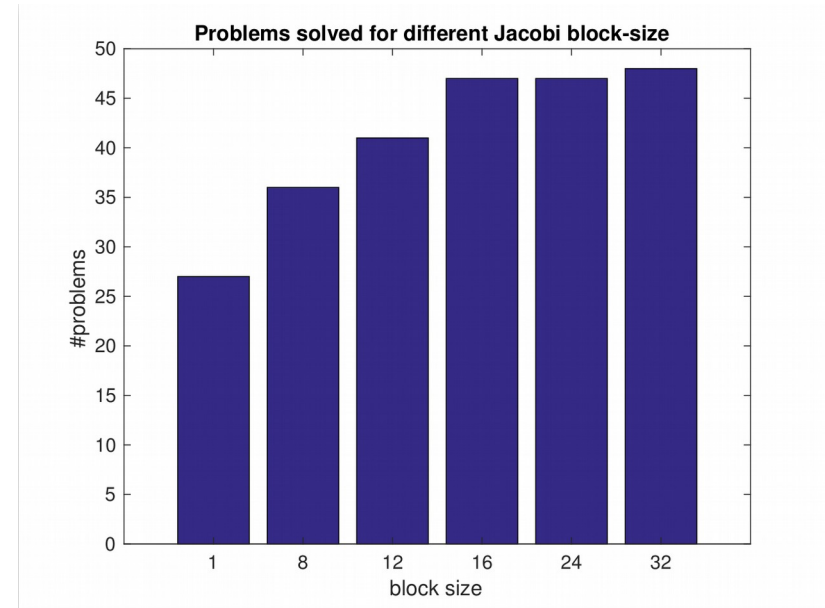
Benefits of block-Jacobi

- 56 matrices from SuiteSparse
- MAGMA-sparse open source library
 - IDR solver
 - Scalar Jacobi preconditioner
 - Supervariable agglomeration
 - Detects block structure of the matrix

M. Goetz and H. Anzt, "Machine Learning-Aided Numerical Linear Algebra: Convolutional Neural Networks for the Efficient Preconditioner Generation," 2018 IEEE/ACM 9th Workshop on Latest Advances in Scalable Algorithms for Large-Scale Systems (scalA), Dallas, TX, USA, 2018, pp. 49-56.

Benefits of block-Jacobi

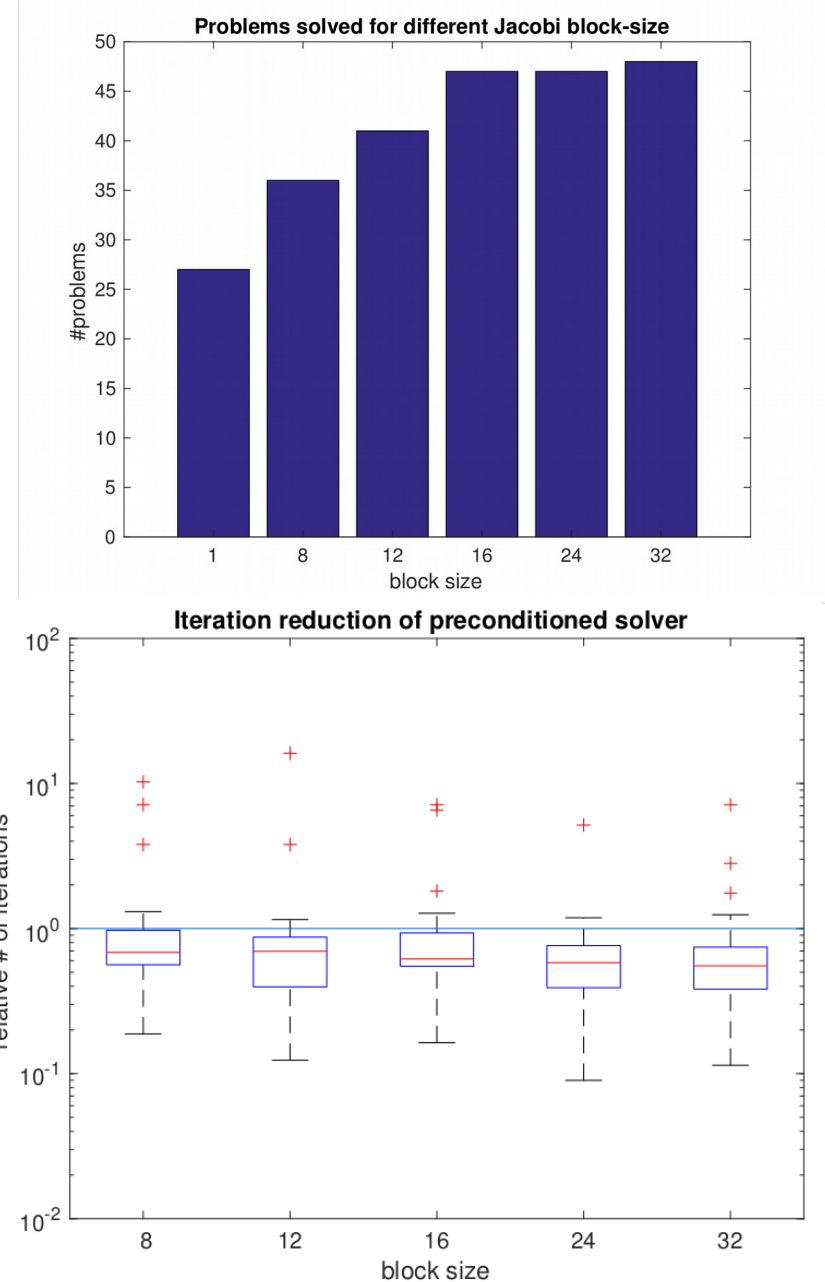
- 56 matrices from SuiteSparse
- MAGMA-sparse open source library
 - IDR solver
 - Scalar Jacobi preconditioner
 - Supervariable agglomeration
 - Detects block structure of the matrix
- Improves the robustness of the solver



M. Goetz and H. Anzt, "Machine Learning-Aided Numerical Linear Algebra: Convolutional Neural Networks for the Efficient Preconditioner Generation," 2018 IEEE/ACM 9th Workshop on Latest Advances in Scalable Algorithms for Large-Scale Systems (scalA), Dallas, TX, USA, 2018, pp. 49-56.

Benefits of block-Jacobi

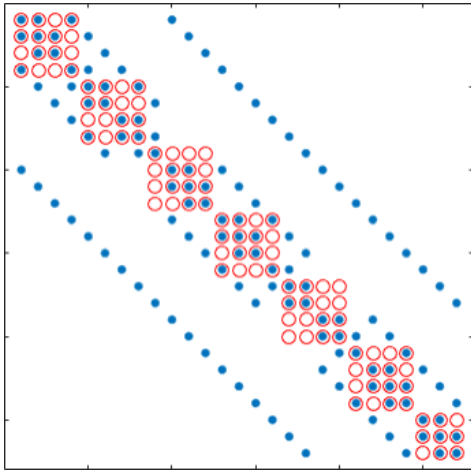
- 56 matrices from SuiteSparse
- MAGMA-sparse open source library
 - IDR solver
 - Scalar Jacobi preconditioner
 - Supervariable agglomeration
 - Detects block structure of the matrix
- Improves the robustness of the solver
- Improves convergence of the solver



M. Goetz and H. Anzt, "Machine Learning-Aided Numerical Linear Algebra: Convolutional Neural Networks for the Efficient Preconditioner Generation," 2018 IEEE/ACM 9th Workshop on Latest Advances in Scalable Algorithms for Large-Scale Systems (scalA), Dalla, TX, USA, 2018, pp. 49-56.

Block-Jacobi preconditioning

- Improve performance for problems with inherent block structure
 - Usually blocks are smaller than 30x30 (**blocks can be of different sizes!**)



- Block-Jacobi preconditioning
 - Use only diagonal blocks for approximation

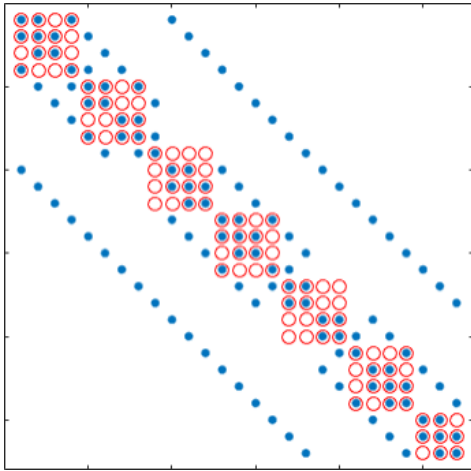
$$\text{diag}(A) = [D_1, \dots, D_k]$$

$$M := \text{diag}(D_1, \dots, D_k)$$

$$y := M^{-1}z \quad \longrightarrow \quad y_i := D_i^{-1}z_i, \quad \forall i$$

Block-Jacobi preconditioning

- Improve performance for problems with inherent block structure
 - Usually blocks are smaller than 30x30 (**blocks can be of different sizes!**)



- Block-Jacobi preconditioning
 - Use only diagonal blocks for approximation

$$\text{diag}(A) = [D_1, \dots, D_k]$$

$$M := \text{diag}(D_1, \dots, D_k)$$

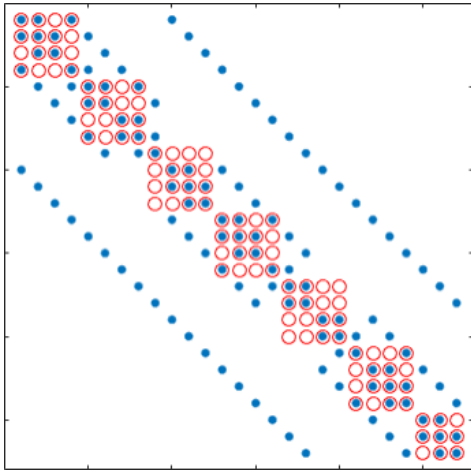
$$y := M^{-1}z \quad \longrightarrow \quad y_i := D_i^{-1}z_i, \quad \forall i$$

$$\longrightarrow \quad \tilde{D}_i := \text{inv}(D_i)$$
$$y_i := \tilde{D}_i z_i$$

inv = Gauss-Jordan elimination

Block-Jacobi preconditioning

- Improve performance for problems with inherent block structure
 - Usually blocks are smaller than 30x30 (**blocks can be of different sizes!**)



- Block-Jacobi preconditioning
 - Use only diagonal blocks for approximation

$$\text{diag}(A) = [D_1, \dots, D_k]$$

$$M := \text{diag}(D_1, \dots, D_k)$$

$$y := M^{-1}z \quad \longrightarrow \quad y_i := D_i^{-1}z_i, \quad \forall i$$

$$\tilde{D}_i := \text{inv}(D_i)$$

Setup

$$y_i := \tilde{D}_i z_i$$

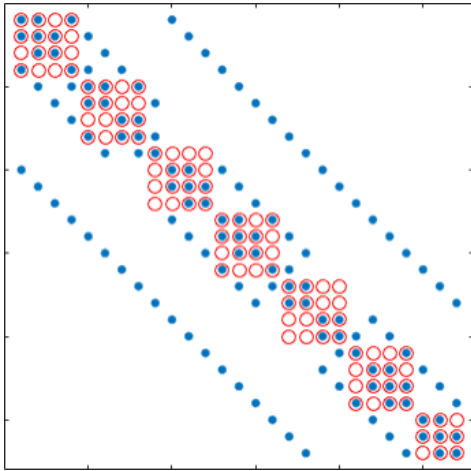
Application

Inversion-based block-Jacobi

inv = Gauss-Jordan elimination

Block-Jacobi preconditioning

- Current focus: improve performance for problems with inherent block structure
 - Usually blocks are smaller than 30x30 (**blocks can be of different sizes!**)



- Block-Jacobi preconditioning
 - Use only diagonal blocks for approximation

$$\text{diag}(A) = [D_1, \dots, D_k]$$

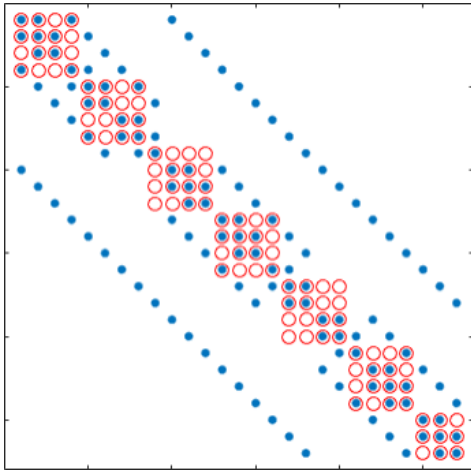
$$M := \text{diag}(D_1, \dots, D_k)$$

$$y := M^{-1}z \quad \longrightarrow \quad y_i := D_i^{-1}z_i, \quad \forall i$$

$$\longrightarrow \quad D_i y_i = z_i$$

Block-Jacobi preconditioning

- Current focus: improve performance for problems with inherent block structure
 - Usually blocks are smaller than 30x30 (**blocks can be of different sizes!**)



- Block-Jacobi preconditioning
 - Use only diagonal blocks for approximation

$$\text{diag}(A) = [D_1, \dots, D_k]$$

$$M := \text{diag}(D_1, \dots, D_k)$$

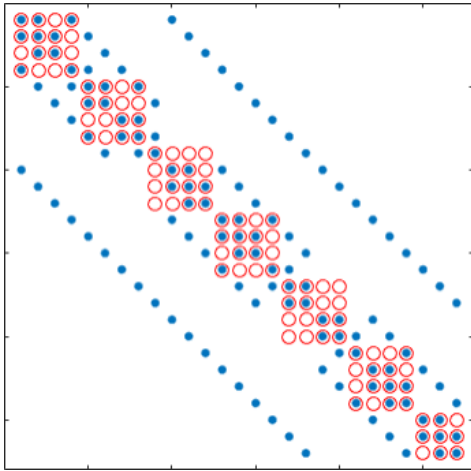
$$y := M^{-1}z \quad \longrightarrow \quad y_i := D_i^{-1}z_i, \quad \forall i$$

$$D_i = L_i U_i$$

$$\longrightarrow D_i y_i = z_i \quad \longrightarrow \quad \begin{aligned} U_i y_i &= w_i \\ L_i w_i &= z_i \end{aligned}$$

Block-Jacobi preconditioning

- Current focus: improve performance for problems with inherent block structure
 - Usually blocks are smaller than 30x30 (**blocks can be of different sizes!**)



- Block-Jacobi preconditioning
 - Use only diagonal blocks for approximation

$$\text{diag}(A) = [D_1, \dots, D_k]$$

$$M := \text{diag}(D_1, \dots, D_k)$$

$$y := M^{-1}z \longrightarrow y_i := D_i^{-1}z_i, \quad \forall i$$

$$D_i = L_i U_i$$

Setup

$$\longrightarrow D_i y_i = z_i \longrightarrow$$

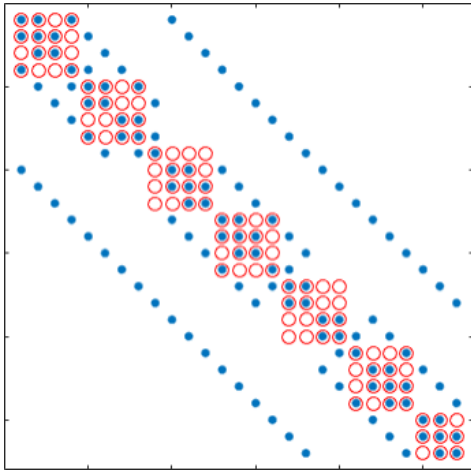
$$\begin{aligned} U_i y_i &= w_i \\ L_i w_i &= z_i \end{aligned}$$

Application

Decomposition-based block-Jacobi

Block-Jacobi preconditioning

- Current focus: improve performance for problems with inherent block structure
 - Usually blocks are smaller than 30x30 (**blocks can be of different sizes!**)



- Block-Jacobi preconditioning
 - Use only diagonal blocks for approximation

$$\text{diag}(A) = [D_1, \dots, D_k]$$

$$M := \text{diag}(D_1, \dots, D_k)$$

$$y := M^{-1}z \longrightarrow y_i := D_i^{-1}z_i, \quad \forall i$$

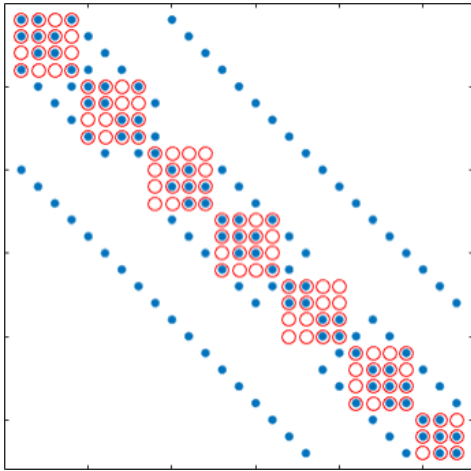
$$\tilde{D}_i := \text{gh}_d(D_i)$$

$$\longrightarrow D_i y_i = z_i \longrightarrow z_i := \text{gh}_a(\tilde{D}_i, y_i)$$

gh = Gauss-Huard decomposition / application

Block-Jacobi preconditioning

- Current focus: improve performance for problems with inherent block structure
 - Usually blocks are smaller than 30x30 (**blocks can be of different sizes!**)



- Block-Jacobi preconditioning
 - Use only diagonal blocks for approximation

$$\text{diag}(A) = [D_1, \dots, D_k]$$

$$M := \text{diag}(D_1, \dots, D_k)$$

$$y := M^{-1}z \longrightarrow y_i := D_i^{-1}z_i, \quad \forall i$$

$$\longrightarrow D_i y_i = z_i \longrightarrow$$

$$\tilde{D}_i := \text{gh}_d(D_i)$$

Setup

$$z_i := \text{gh}_a(\tilde{D}_i, y_i)$$

Application

Decomposition-based block-Jacobi

Block-Jacobi setup & application ecosystem


Setup

Gauss-Jordan inversion



Inversion-based
($2n^3 + 2n^2$ FLOPS)


matrix-vector multiply



H. Anzt et al. "Variable-size batched Gauss-Jordan elimination for block-Jacobi preconditioning on graphics processors", ParCo'19


Application

Gauss-Huard decomposition



Decomposition-based
($2/3n^3 + 2n^2$ FLOPS)

Gauss-Huard solve




H. Anzt et al. "Variable-Size Batched Gauss-Huard for Block-Jacobi Preconditioning", ICCS'17

LU factorization



2x triangular solve



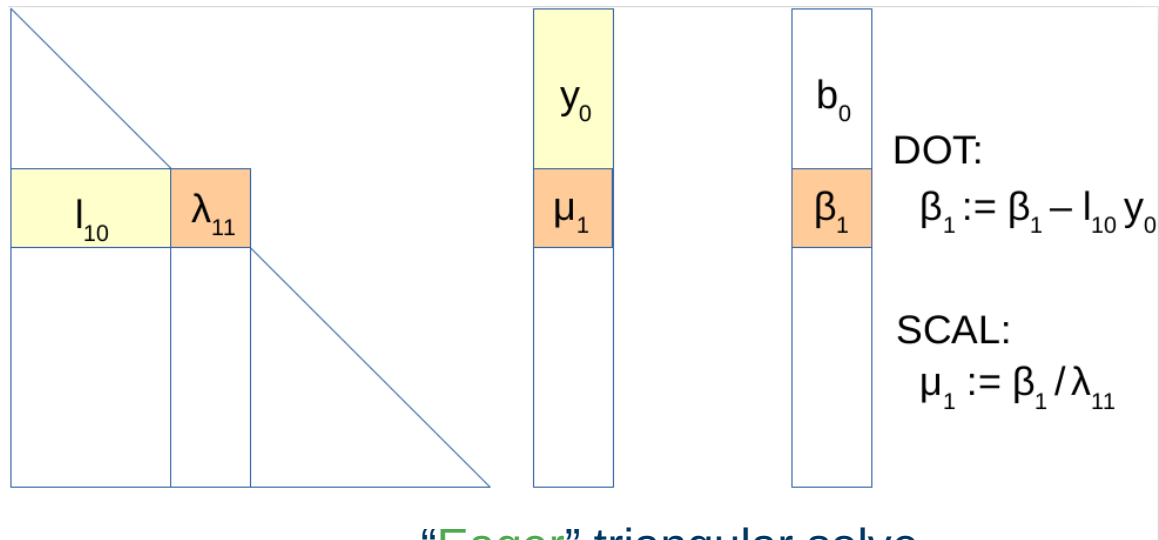
H. Anzt et al. "Flexible-Size Batched LU for Small Matrices and its Integration into Block-Jacobi Preconditioning", ICPP'17

● - read ● - write

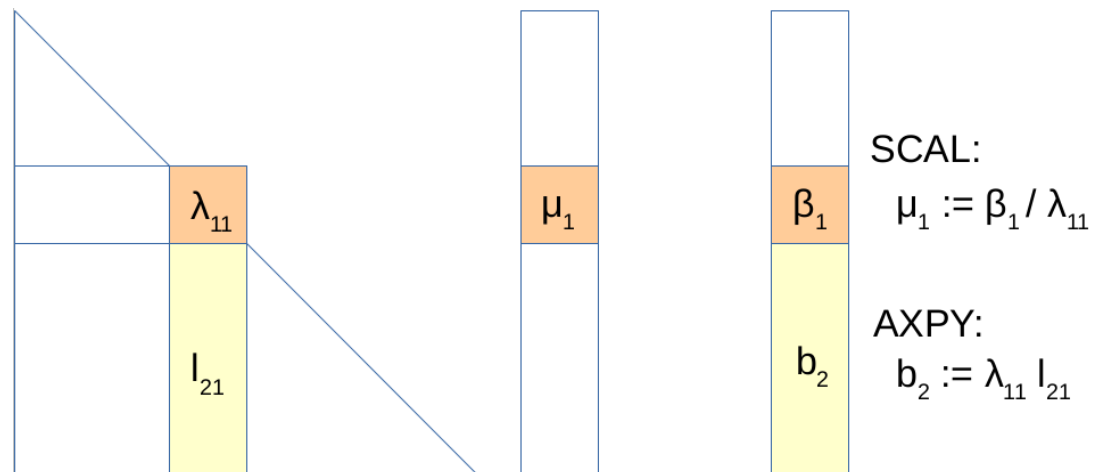
Triangular solves

- Use “eager” triangular solves
 - Cast solution vector updates in terms of axpy, not in terms of dot product!

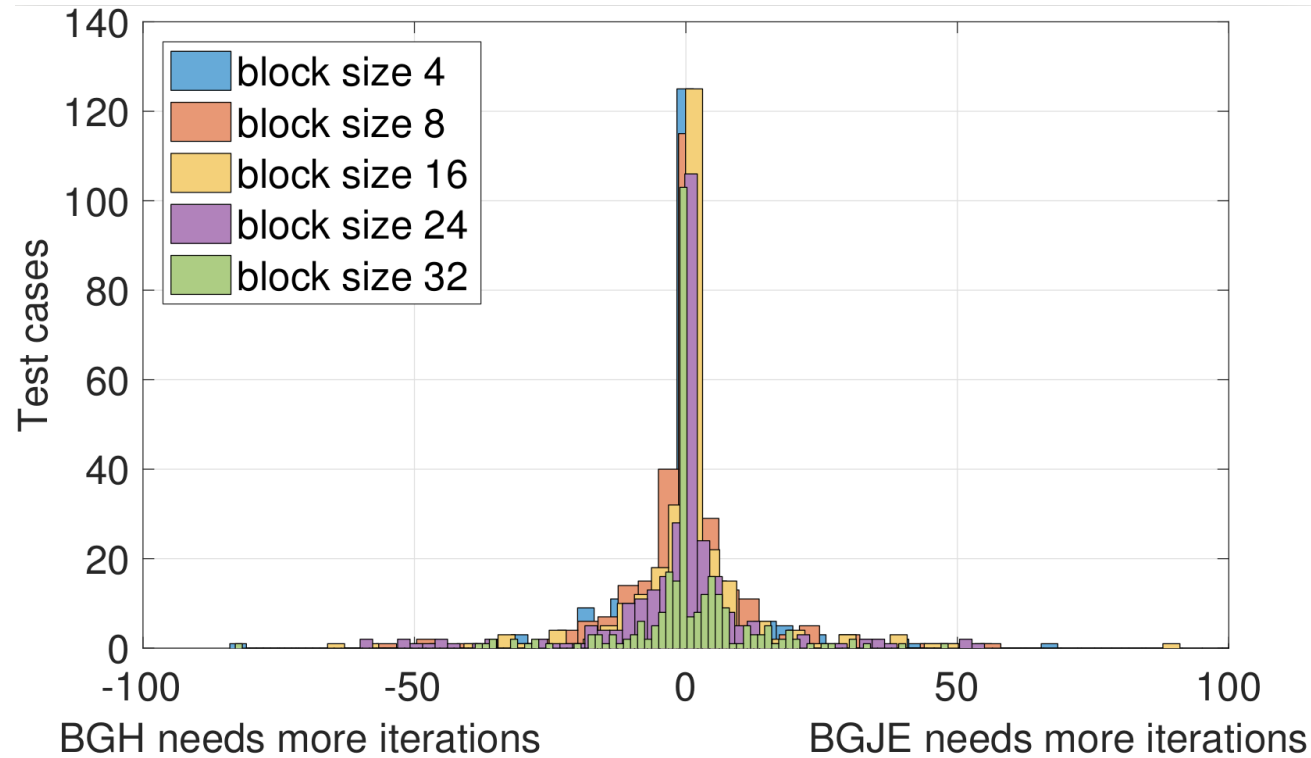
“Lazy” triangular solve



“Eager” triangular solve



No signs of instability in practice



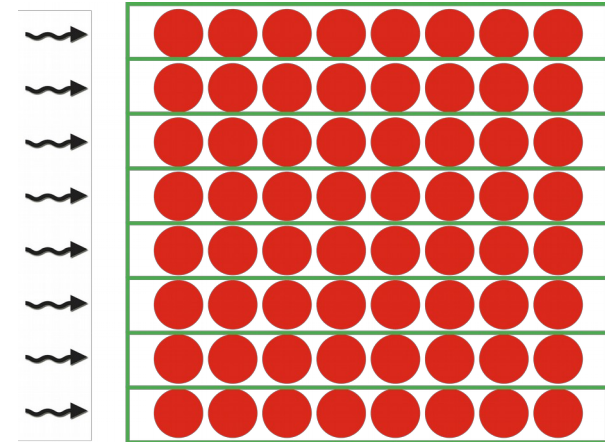
Batched inversion and decomposition

- Assign one warp to each problem
 - hardware SIMD unit, represented as a group of 32 threads in CUDA



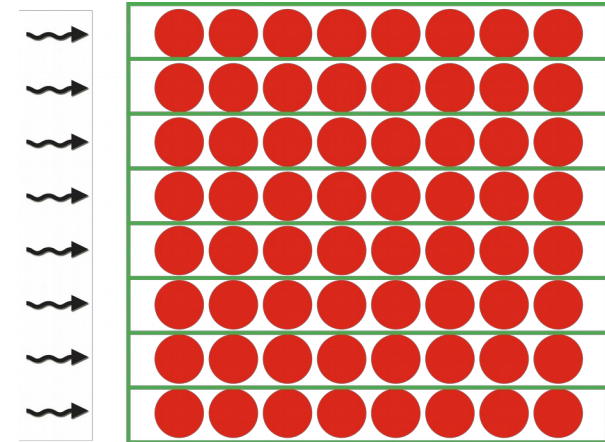
Batched inversion and decomposition

- Assign **one warp to each problem**
 - hardware SIMD unit, represented as a group of 32 threads in CUDA
- Process each row (column) by a single thread
 - Able to support problems of size up to 32-by-32
 - Keep the entire row (column) in thread's **registers**
 - Communicate data between rows via **warp-shuffles**
 - Current implementation: **use padding for problems of smaller sizes**
 - Future work: **multiple smaller problems per warp**

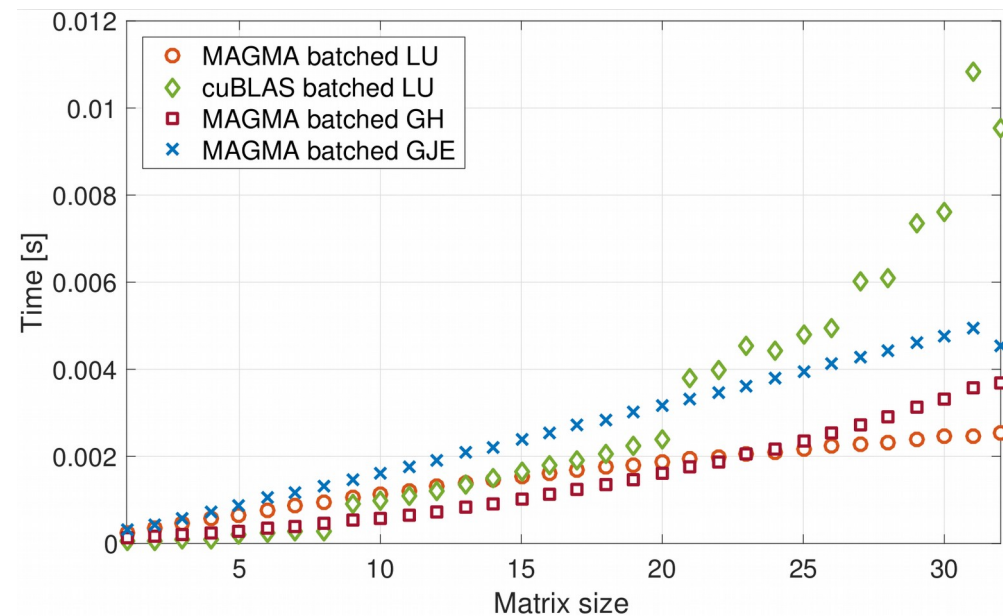
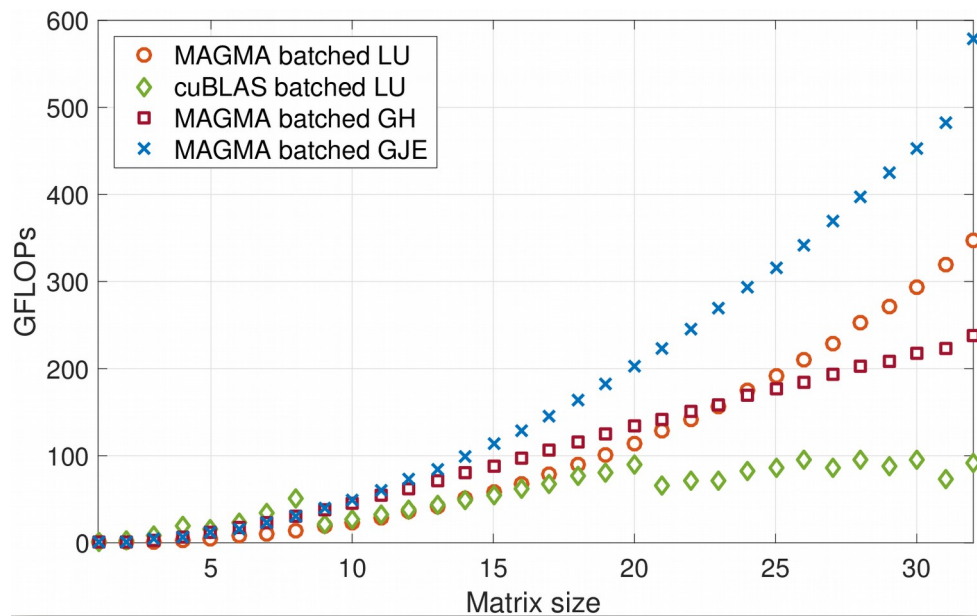


Batched inversion and decomposition


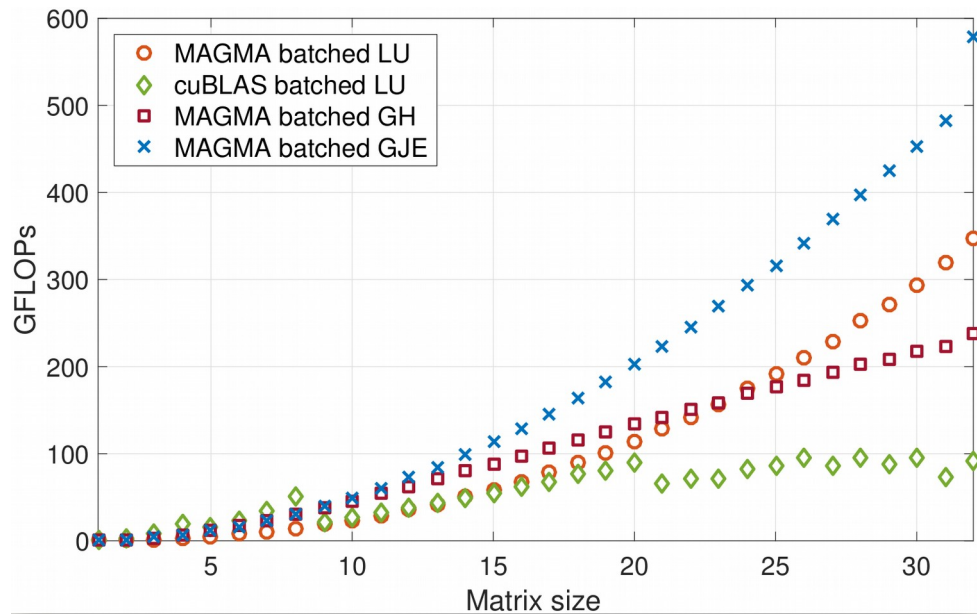
- Assign **one warp to each problem**
 - hardware SIMD unit, represented as a group of 32 threads in CUDA
- Process each row (column) by a single thread
 - Able to support problems of size up to 32-by-32
 - Keep the entire row (column) in thread's **registers**
 - Communicate data between rows via **warp-shuffles**
 - Current implementation: **use padding for problems of smaller sizes**
 - Future work: **multiple smaller problems per warp**
- Use implicit pivoting
 - Do not explicitly swap rows (column), “re-assign” the threads instead



Batched routines performance




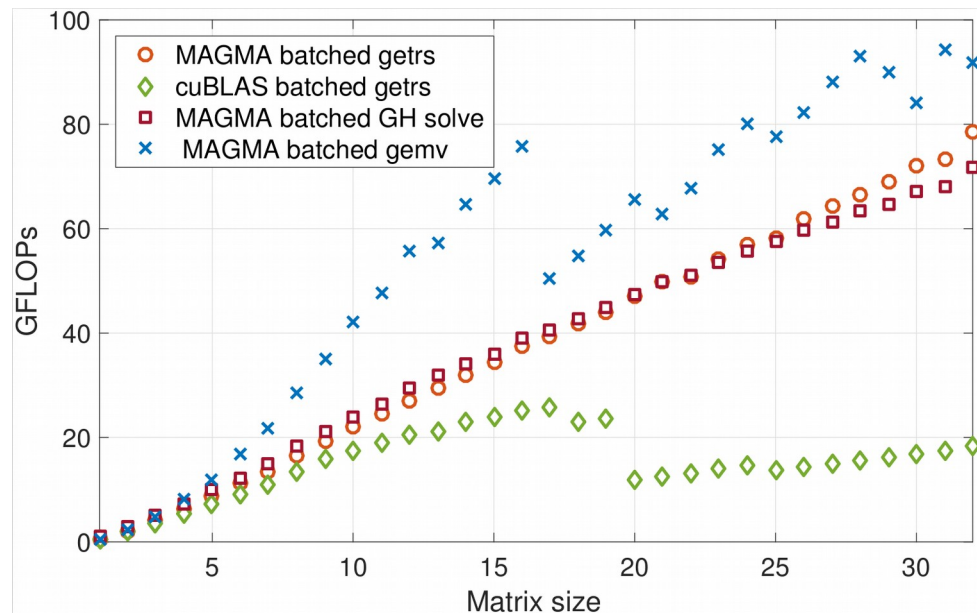

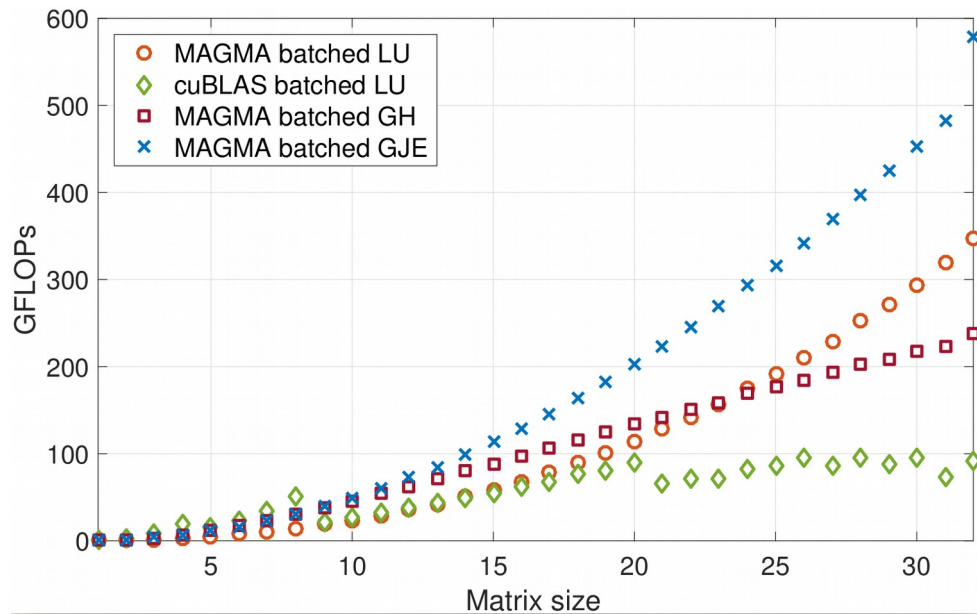
Batched routines performance



MAGMA-sparse routines can also:

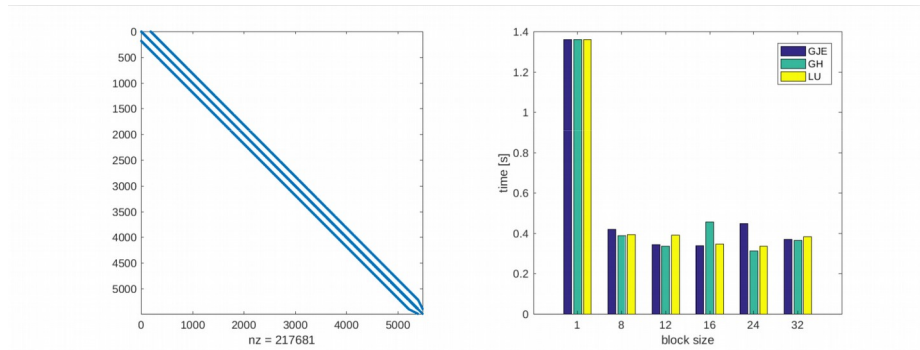
- Handle problems of different sizes
- Integrate diagonal block extraction and diagonal block decomposition / inversion into a single kernel

Batched routines performance

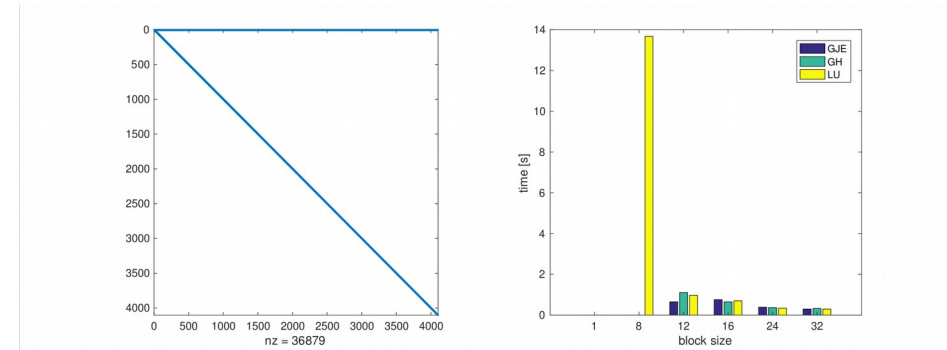


Complete solver runtime

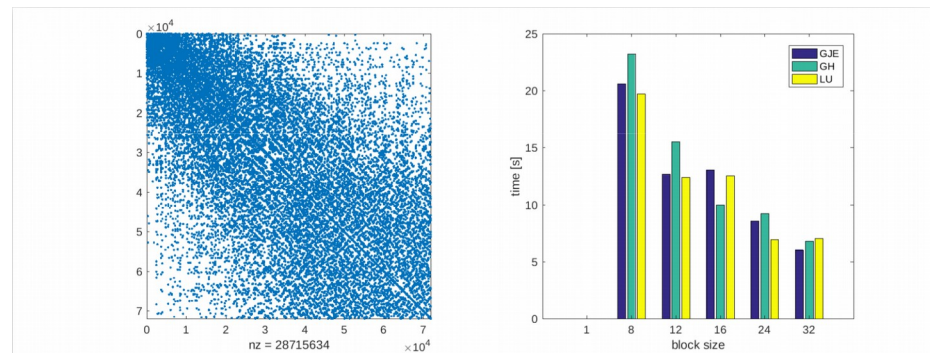
s2rmt3m1



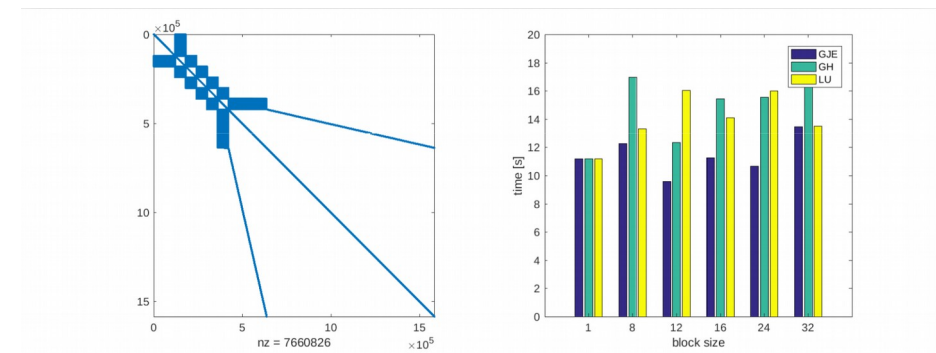
Chebyshev3



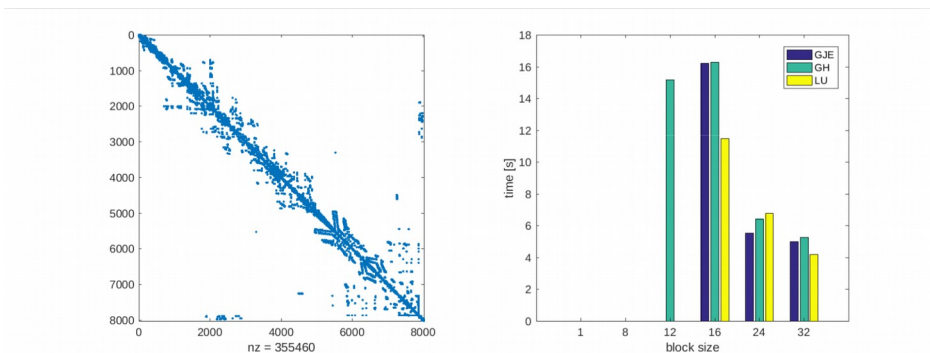
nd24k



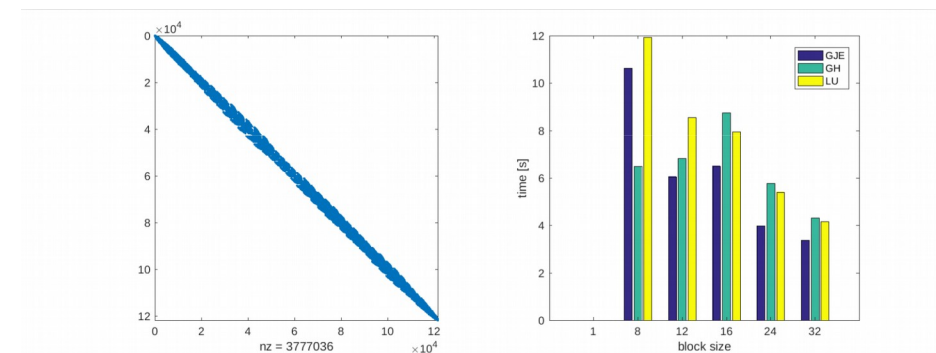
G3_circuit



bcsstk38

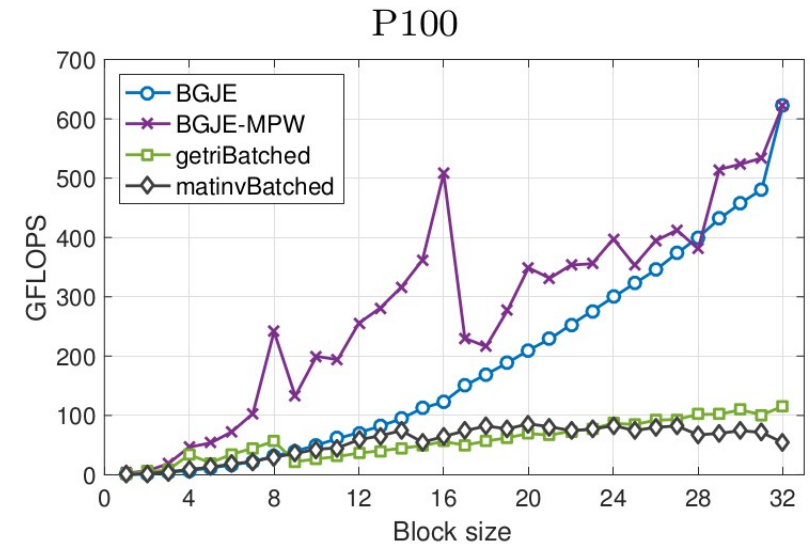


ship_003




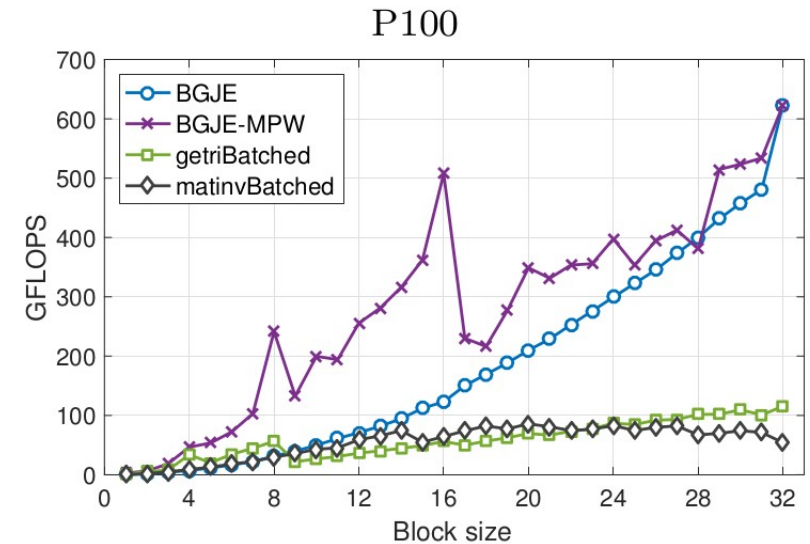
Flexible-size batched routines

- Problems can be too small to effectively use one warp
 - Solution: assign multiple problems per warp



Flexible-size batched routines

- Problems can be too small to effectively use one warp
 - Solution: assign multiple problems per warp
- Allow batches where problems are of different sizes (flexible-size)
 - How to combine this with multiple problems per warp?
 - Remember: entire warp executes the same instruction!
 - Current solution: padding



Adaptive precision block-Jacobi

Preconditioner is an **approximation** of the system matrix

- Applying a preconditioner inherently carries an **error**
- For block-Jacobi the relative error of z is usually around 0.01 – 0.1

$$z := M^{-1}y \approx A^{-1}y$$

Anzt, et al. "Adaptive precision in block-Jacobi preconditioning for iterative sparse linear system solvers." Concurrency and Computation: Practice and Experience 31.6 (2019): e4460.

Adaptive precision block-Jacobi

Preconditioner is an **approximation** of the system matrix

- Applying a preconditioner inherently carries an **error**
- For block-Jacobi the relative error of z is usually around 0.01 – 0.1

$$z := M^{-1}y \approx A^{-1}y$$

Preconditioner application is **memory bounded**

- Most of the cost comes from reading the matrix from memory
- Idea: use **lower precision** to **store** the matrix

Anzt, et al. "Adaptive precision in block-Jacobi preconditioning for iterative sparse linear system solvers." Concurrency and Computation: Practice and Experience 31.6 (2019): e4460.

Adaptive precision block-Jacobi

Preconditioner is an **approximation** of the system matrix

- Applying a preconditioner inherently carries an **error**
- For block-Jacobi the relative error of z is usually around 0.01 – 0.1

$$z := M^{-1}y \approx A^{-1}y$$

Preconditioner application is **memory bounded**

- Most of the cost comes from reading the matrix from memory
- Idea: use **lower precision** to **store** the matrix

Adaptive precision in inversion-based block-Jacobi:

- All **computation** is done in **double precision**
- Preconditioner matrix is **stored** in **lower precision**, with roundoff error “ u ”
- Error bound:

$$\frac{\|\delta z_i\|}{\|z_i\|} \lesssim (c_i \kappa(D_i) u_d + u) \kappa(D_i)$$

Anzt, et al. "Adaptive precision in block-Jacobi preconditioning for iterative sparse linear system solvers." Concurrency and Computation: Practice and Experience 31.6 (2019): e4460.

Adaptive precision block-Jacobi

Preconditioner is an **approximation** of the system matrix

- Applying a preconditioner inherently carries an **error**
- For block-Jacobi the relative error of z is usually around 0.01 – 0.1

$$z := M^{-1}y \approx A^{-1}y$$

Preconditioner application is **memory bounded**

- Most of the cost comes from reading the matrix from memory
- Idea: use **lower precision** to **store** the matrix

Adaptive precision in inversion-based block-Jacobi:

- All **computation** is done in **double precision**
- Preconditioner matrix is **stored** in **lower precision**, with roundoff error “ u ”
- Error bound:

$$\frac{\|\delta z_i\|}{\|z_i\|} \lesssim (c_i \kappa(D_i) u_d + u) \kappa(D_i)$$

Assuming the preconditioner block is relatively well conditioned

- The error is determined by the product of u , and the condition number
- **Choose the precision for each block independently, such that at least 1 digit of the result is correct**

Experimental results

Determining the precision:

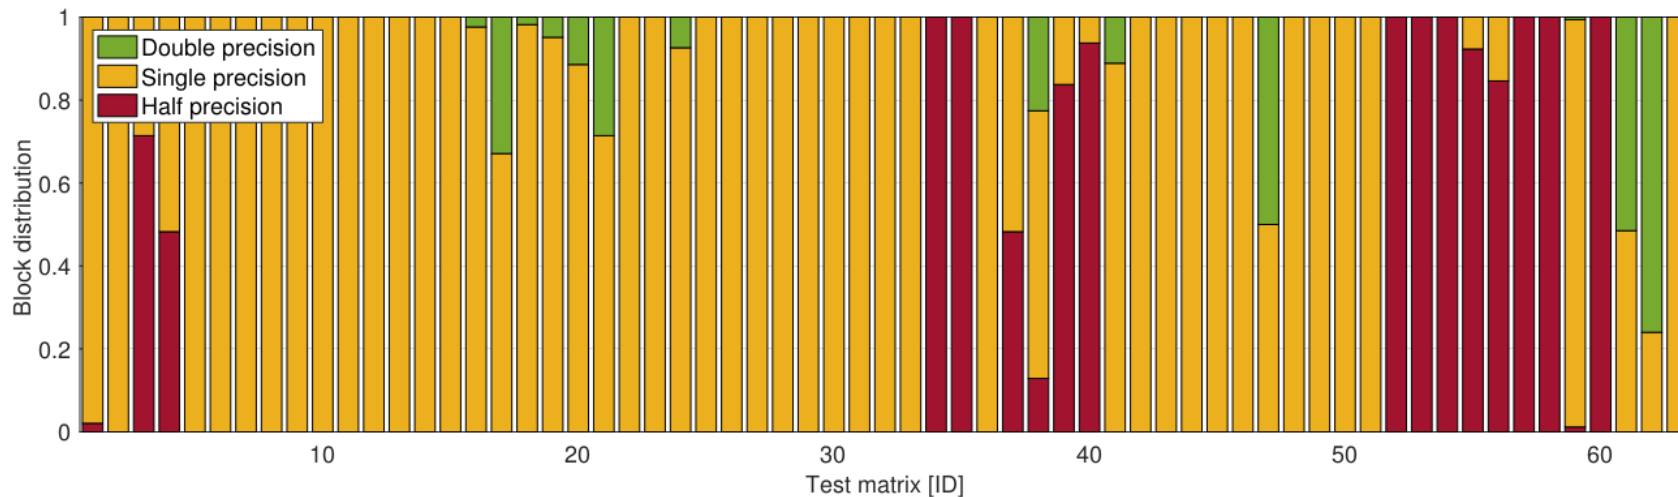
$$\begin{cases} \text{fp16} & \text{if } 0 < \kappa_1(D_i) \leq 10^2, \\ \text{fp32} & \text{if } 10^2 < \kappa_1(D_i) \leq 10^6, \text{ and} \\ \text{fp64} & \text{otherwise,} \end{cases}$$

Experimental results

Determining the precision:

$$\begin{cases} \text{fp16} & \text{if } 0 < \kappa_1(D_i) \leq 10^2, \\ \text{fp32} & \text{if } 10^2 < \kappa_1(D_i) \leq 10^6, \text{ and} \\ \text{fp64} & \text{otherwise,} \end{cases}$$

% of diagonal blocks stored in each precision *:

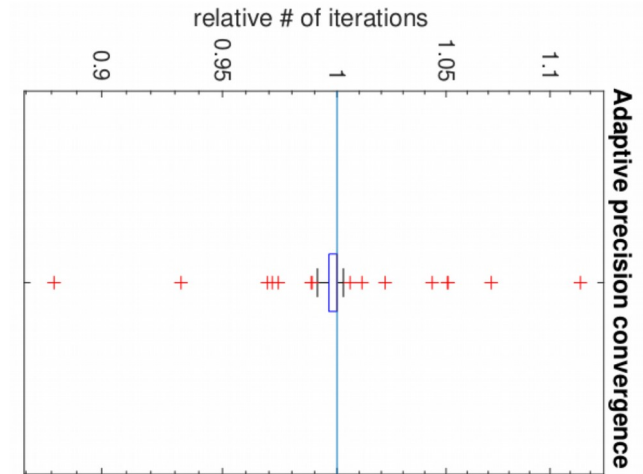


* Prototype implementation in MATLAB,
Results on 63 matrices from SuiteSparse


Experimental results

Determining the precision:

$$\begin{cases} \text{fp16} & \text{if } 0 < \kappa_1(D_i) \leq 10^2, \\ \text{fp32} & \text{if } 10^2 < \kappa_1(D_i) \leq 10^6, \text{ and} \\ \text{fp64} & \text{otherwise,} \end{cases}$$



% of diagonal blocks stored in each precision *:



* Prototype implementation in MATLAB,
Results on 63 matrices from SuiteSparse

Overflow and underflow


- Storing the block in lower precision can cause underflow or overflow in some values.
- This is not accounted for by the numerical analysis shown in the previous slides

Overflow and underflow


- Storing the block in lower precision can cause underflow or overflow in some values.
- This is not accounted for by the numerical analysis shown in the previous slides
- Two possible solutions:
 1. **Check the condition number** of the low precision block to verify there were no catastrophic overflows or underflows.

Overflow and underflow


- Storing the block in lower precision can cause underflow or overflow in some values.
- This is not accounted for by the numerical analysis shown in the previous slides
- Two possible solutions:
 1. **Check the condition number** of the low precision block to verify there were no catastrophic overflows or underflows.
 2. **Use a custom storage format** that preserves the number of exponent bits:



Preconditioner application performance



Time-to-solution, adaptive- vs full-precision BJ



Summary

- Employ inherently parallel preconditioning schemes on accelerator hardware
- Trade more computation for better access patterns
- Take advantage of large register field of NVIDIA GPUs
- Avoid data exchange by “renaming” CUDA threads
- Take advantage of low precision to accelerate preconditioner application
- Independently select precision for each subproblem to preserve preconditioner quality


KIT, Karlsruhe, 12 June 2019



Linear systems, Ginkgo


Sources of linear systems

- Real-world problem transformed into a linear system via:
 - PDE discretizations, graph representations
 - Large number of unknowns (1M+, full matrix 8TB)
 - Most matrix elements are zero
- Possible approach: iterative methods
 - Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver
 - Preconditioners
 - reduce the number of iterations



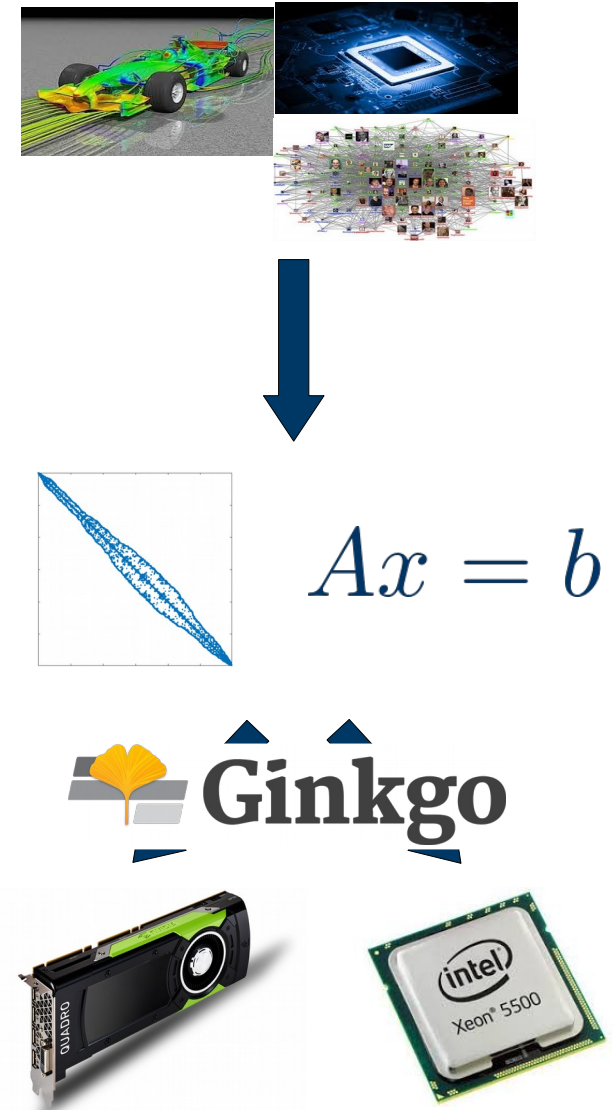
Sources of linear systems

- Real-world problem transformed into a linear system via:
 - PDE discretizations, graph representations
 - Large number of unknowns (1M+, full matrix 8TB)
 - Most matrix elements are zero
- Possible approach: iterative methods
 - Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver
 - Preconditioners
 - reduce the number of iterations
 - Special hardware (e.g. GPUs)
 - Probably not a good idea to implement everything from scratch...



Sources of linear systems

- Real-world problem transformed into a linear system via:
 - PDE discretizations, graph representations
 - Large number of unknowns (1M+, full matrix 8TB)
 - Most matrix elements are zero
- Possible approach: iterative methods
 - Krylov-subspace based linear solvers
 - SpMV
 - BLAS-1 operations
 - Sparse matrix formats & SpMV
 - accelerate each iteration of the solver
 - Preconditioners
 - reduce the number of iterations
 - Special hardware (e.g. GPUs)
 - Probably not a good idea to implement everything from scratch...
 - Use a library instead:



KIT, Karlsruhe, 12 June 2019



Preconditioning

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n}$$

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an
equivalent preconditioned system

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an
equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an
equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

~~$$M^{-1}A$$~~

**Do not compute the preconditioned
system matrix explicitly!**

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an
equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

~~$$M^{-1}A$$~~

$$y := (M^{-1}A)x$$

**Do not compute the preconditioned
system matrix explicitly!**

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

$$\cancel{M^{-1}A}$$

Do not compute the preconditioned system matrix explicitly!

$$y := (M^{-1}A)x$$



$$z := Ax$$

$$y := M^{-1}z$$

Preconditioner application

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

~~$$M^{-1}A$$~~

Do not compute the preconditioned system matrix explicitly!

$$y := (M^{-1}A)x$$



Generate the preconditioner matrix, and store it in a form suitable for application

$$A \rightsquigarrow M$$

Preconditioner setup



$$\begin{aligned} z &:= Ax \\ y &:= M^{-1}z \end{aligned}$$

Preconditioner application

Preconditioning

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \quad \longrightarrow \quad M^{-1}Ax = M^{-1}b$$

Replace the original system with an equivalent preconditioned system

$$M \approx A \quad M^{-1} \text{ easy to compute}$$

~~$$M^{-1}A$$~~

Do not compute the preconditioned system matrix explicitly!

$$y := (M^{-1}A)x$$



Generate the preconditioner matrix, and store it in a form suitable for application

$$A \rightsquigarrow M$$

Preconditioner setup



$$\begin{aligned} z &:= Ax \\ y &:= M^{-1}z \end{aligned}$$

Preconditioner application

Trade-off:
faster convergence,
but more work per iteration